

Agentic Hybrid RAG on SAP BTP

Sriram Rokkam

May 2026

Contents

Front Matter	10
Agentic Hybrid RAG on SAP BTP	10
A Hands-On Guide with LangGraph, HANA Cloud, and Vertex AI	10
Title Page	10
Copyright	10
Dedication	11
Foreword	11
Preface	11
Who This Book Is For	12
Who This Book Is NOT For	12
Prerequisites	13
Technical Knowledge	13
NOT Required	14
Tools and Accounts (All Free)	14
Companion Code Repository	14
Repository Structure	14
Quick Start	15
Chapter-to-Code Cross-Reference	16
How to Use This Book	16
Reading Cover to Cover (Recommended)	16
As a Reference	16
Chapter Structure	16
The Code	16
Conventions Used in This Book	17
Code Blocks	17
Callout Boxes	17
File Paths	17
Screenshots	18
Acknowledgements	18
About the Author	18
Chapter 1: Agentic AI on SAP BTP — What It Is, Why It Matters, and What We Build	19
1.1 The Agentic AI Moment	19
1.2 Why SAP Developers Need to Understand Agents Now	20

1.3 Why BTP Is the Right Platform for Agentic AI	22
1.4 What This Book Builds: The Material Document Intelligence Platform	24
1.5 Why Hybrid Retrieval.	27
1.6 The Four Agentic Capabilities in This System.	29
1.7 Scope and Prerequisites	30
What This Book Covers.	30
What This Book Does Not Cover.	31
Prerequisites	32
1.8 Summary	33
Chapter 2: Platform Setup — SAP BTP Trial + Google Vertex AI	35
2.1 The Goal of This Chapter.	35
2.2 Understanding What We Are Setting Up	35
2.3 SAP BTP Trial Account.	37
Signing Up.	37
Navigating to Your Trial Subaccount	38
2.4 Provisioning SAP HANA Cloud	39
Creating the HANA Cloud Instance	39
Enabling the Vector Engine	41
Enabling the Knowledge Graph Engine	42
Getting Your HANA Connection Details	43
2.5 Setting Up Google Cloud + Vertex AI	44
Creating a GCP Account	44
Creating a New Project.	44
Enabling the Vertex AI API.	45
GCP Authentication — What Actually Works	45
Creating a Service Account	46
Option A — Local Development: Application Default Credentials (ADC)	46
Option B — BTP Cloud Foundry Deployment: Service Account JSON Key.	47
2.6 Configuring the BTP Destination Service.	47
Creating the Destination	48
2.7 Local Development Environment	50
Python 3.11+	50
Node.js 20+	50
Cloud Foundry CLI	50
SAP CDS CLI	51
VS Code Extensions	52
2.8 Project Setup	52
2.9 Verifying the Setup: Your First Gemini API Call.	54

2.10 Verifying the HANA Cloud Connection	56
2.11 What We Have Built.	57
2.12 Understanding the Credential Flow	58
2.13 Summary	59
Chapter 3: Vector Search on SAP HANA Cloud	60
3.1 What Are Embeddings? Intuition Without the Math	60
3.2 The HANA Cloud REAL_VECTOR Column Type	62
3.3 Creating the Vector Table: The Lazy Initialization Pattern	64
3.4 Calling text-embedding-004 from Python	66
3.5 Storing Embeddings: INSERT with REAL_VECTOR	68
3.6 Cosine Similarity Search: The Query Explained	69
3.7 Chunk Size Analysis and Boundary Strategy for SAP Material Documents	71
Why chunking exists at all	71
The chunk size trade-off	72
Why 50-token overlap	72
Natural boundaries by document type.	73
Our chunking parameters	74
3.8 Building vector_srv.py — The Vector Service Layer	75
3.9 Testing: Business Questions Across Document Types	78
3.10 What Vector Search Gets Right — and What It Gets Wrong	81
3.11 Summary	83
3.12 Checkpoint.	84
Chapter 4: Knowledge Graphs on SAP HANA Cloud	86
4.1 What is a Knowledge Graph?	86
Why call it a graph?	87
What does this give us that a relational schema does not?	88
4.2 RDF and SPARQL — the standards explained plainly	89
4.3 Why HANA Cloud for the Knowledge Graph	90
Why a stored procedure and not an HTTP endpoint?	92
4.4 Designing the MSDS ontology — what facts matter?	92
4.5 The OWL ontology file — MSDS_Ontology.ttl	93
4.6 Extracting triples from MSDS text using Gemini	95
4.7 Storing triples in HANA — named graphs per material.	97
4.8 Writing SPARQL queries — Gemini as the SPARQL translator	99
4.9 Building kg_srv.py — the Knowledge Graph service layer	101

Validation: the KG equivalent of SQL injection	105
What delete_graph and count_triples give us	106
4.10 Testing — store triples for a batch certificate, query for supplier and certificate number	107
4.11 What the KG gets right that vector search cannot.	109
4.12 Summary	111
4.13 Checkpoint.	112
Chapter 5: The Document Intelligence Ingestion Pipeline.	114
5.1 The dual-pipeline architecture	115
5.1.1 Why process both in parallel?	115
5.1.2 The fire-and-forget pattern	116
5.2 PDF text extraction with PyMuPDF	117
5.2.1 Extracting with metadata	117
5.2.2 Handling upload files in FastAPI	119
5.3 Chunking strategy for vector search	119
5.3.1 Our chunking parameters	120
5.3.2 The chunker function	121
5.3.3 Why the KG pipeline does not chunk	122
5.4 Thread-local HANA connections	123
5.4.1 Why connections cannot be shared across threads	123
5.4.2 The threading.local() solution	123
5.5 The ingestion flow in detail	125
5.5.1 The upload endpoint	125
5.5.2 The dual-pipeline orchestrator	126
5.5.3 Thread 1 — the vector pipeline	127
5.5.4 Thread 2 — the Knowledge Graph pipeline	128
5.6 Status tracking	129
5.6.1 The MSDS_DOCUMENTS table and its role	129
5.6.2 Status helper functions	130
5.6.3 The status endpoint	131
5.7 Error handling — what happens when one pipeline fails	132
5.8 Testing — upload five material documents	133
5.8.1 Start the service.	133
5.8.2 Upload the first document	133
5.8.3 Poll for completion.	133
5.8.4 Verify the vector store	134
5.8.5 Verify the Knowledge Graph	134

5.9 Summary	135
5.10 Checkpoint.	136
Chapter 6: LangGraph Fundamentals.	138
6.1 Why LangGraph is SAP's recommended orchestration approach for BTP agents	138
6.2 Why LangChain chains are not enough	139
6.3 Core concepts	140
6.3.1 State	140
6.3.2 Nodes	141
6.3.3 Edges	141
6.3.4 Conditional edges	141
6.4 Installation and project layout	142
6.5 Building the state and nodes	143
6.6 Wiring the graph	144
6.7 Adding a conditional retry edge	146
6.8 The StateGraph visualised.	148
6.9 The complete simple_qa_agent.py	148
6.10 Tool use in LangGraph.	151
6.11 Streaming responses.	152
6.12 Debugging with LangSmith.	153
6.13 Why stateless works on BTP Cloud Foundry	153
6.14 Summary	155
6.15 Checkpoint.	155
Chapter 7: The Parallel Hybrid RAG Agent.	157
7.1 Why parallel beats sequential	157
7.2 The agent state.	159
7.3 The vector chain	160
7.4 The KG chain	162
7.5 The merge function.	167
7.6 The orchestrator	169
7.7 The FastAPI /query endpoint.	172
7.8 The full request/response contract	174
7.9 Conversation history.	175
7.10 Testing: three scenarios that prove hybrid wins	176
7.10.1 The KG wins: a quality engineer asks about test specifications	176
7.10.2 The vector wins: a warehouse manager asks about storage requirements	177
7.10.3 Both contribute: a procurement manager asks a combined question	178

7.11 Monitoring chain performance	178
7.12 Summary	179
7.13 Checkpoint.	180
Chapter 8: The Multi-Agent Supervisor Pattern	182
8.1 Why multi-agent for SAP enterprise	182
8.2 The supervisor pattern	183
8.3 The three specialist agents for material documents	185
8.3.1 HazardAgent.	185
8.3.2 SafetyAgent	185
8.3.3 ComplianceAgent.	185
8.4 Shared state for the supervisor graph.	186
8.5 The supervisor node.	187
8.6 The specialist agents.	189
8.7 The summary agent.	191
8.8 Building the supervisor graph	193
8.9 When to use the supervisor vs the direct agent	194
8.10 The /query-advanced endpoint	195
8.11 Testing: a question that requires all three specialists	197
8.12 Applying the pattern to other SAP document types	198
8.13 Summary	199
8.14 Checkpoint.	199
Chapter 9: CAP Node.js OData V4 — The SAP-Native API Layer + Fiori Elements	
UI	201
9.1 Why CAP?	201
9.2 The two-process architecture.	202
9.3 The CDS data model.	204
9.4 The service definition	206
9.5 The Products entity and S/4HANA value help	207
9.6 The callBackend helper	208
9.7 Action handlers	209
processFile	209
chatQuery	211
deleteFile	212
pollStatus	213

9.8 Fiori Elements UI	214
9.9 Annotations explained	215
HeaderInfo	215
SelectionFields and LinItem.	216
Identification actions	216
FieldGroup	217
9.10 Running locally	217
9.11 Testing the full flow	219
9.12 What you can do after this chapter	220
9.13 Checkpoint checklist.	221
Summary	222
Chapter 10: Deploying to SAP BTP Cloud Foundry	223
10.1 What Changes Between Local and BTP.	223
10.2 Why Two Separate Applications	224
10.3 Pre-Deployment Checklist.	225
10.4 Step 1: Push the Python Agent.	226
Understanding manifest.yml	226
HANA connection in CF	227
Pushing the application	227
10.5 Step 2: Create the User-Provided Service for Credentials.	228
10.6 Step 3: Build and Deploy the CAP Service	230
Running the build	232
Deploying the MTA	232
10.7 Step 4: Configure the BTP Destination for Vertex AI	233
Creating the agent backend destination	234
Creating the Vertex AI destination	234
10.8 Step 5: Load the Ontology	235
10.9 Step 6: Smoke Test the Deployed System.	236
Agent health endpoint	236
Direct query to the agent	236
CAP OData endpoint	237
Fiori Elements UI.	237
10.10 Troubleshooting	237
Memory quota exceeded	237
HANA connection refused.	238
GCP authentication failure	238
CAP cannot reach agent	238
Package not found during staging	239

10.11 Scaling	239
10.12 What You Have Built	240
10.13 Chapter Checkpoint	241
Appendix A: SAP AI Core — Enterprise LLM Integration	243
A.1 What Is SAP AI Core?	243
A.2 Prerequisites	244
A.3 Activating a Model in AI Launchpad	244
Step 1 — Open AI Launchpad	244
Step 2 — Create a Resource Group.	244
Step 3 — Activate a Generative AI Hub Model.	244
Step 4 — Get the Inference Endpoint	245
A.4 Getting the AI Core Service Key	245
A.5 Installing the SDK	246
A.6 Calling the Model Directly (Python)	246
A.7 Replacing Gemini in the Agent	247
Step 1 — Replace <code>_get_llm()</code> in <code>agents/agents/llm.py</code>	247
Step 2 — Replace the embedding call in <code>agents/agents/vector_srv.py</code>	248
Step 3 — Update <code>.env</code>	249
Step 4 — Test the swapped agent	249
A.8 BTP Destination for AI Core (CF Deployment)	250
A.9 Summary	250

Front Matter

Agentic Hybrid RAG on SAP BTP

A Hands-On Guide with LangGraph, HANA Cloud, and Vertex AI

Title Page

Agentic Hybrid RAG on SAP BTP *A Hands-On Guide with LangGraph, HANA Cloud, and Vertex AI*

Author: Sriram Rokkam

First Edition, 2026

Copyright

Copyright © 2026 Sriram Rokkam. All rights reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the author, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

The code samples in this book are released under the MIT License and are freely available at: <https://github.com/sriramrokkam/book-agentic-hybrid-rag-on-btp>

SAP, SAP BTP, SAP HANA, SAP CAP, and Joule are trademarks or registered trademarks of SAP SE. Google Cloud, Vertex AI, and Gemini are trademarks of Google LLC. All other trademarks are the property of their respective owners. The author is not affiliated with SAP SE or Google LLC.

Dedication

[To be added by Sriram Rokkam]

Foreword

[To be written by a colleague, SAP contact, or industry peer after the book is complete. 300–500 words. Suggested topics: why Agentic AI matters for SAP customers today, why the BTP + Vertex AI approach is relevant, and why the author is credible on this subject.]

— **[Name], [Title], [Company]**

Preface

[To be written by Sriram Rokkam in his own voice. Suggested content below — edit freely, this is your personal story.]

I have spent years working at the intersection of SAP and cloud platforms, watching enterprises struggle with the same problem in different forms: they have extraordinary data — decades of business transactions, master data, documents, and domain knowledge — but they cannot easily ask questions of it.

The arrival of large language models changed what was possible. Suddenly, the gap between a business question and a data answer felt bridgeable. But every proof of concept I saw had the same flaw: it worked for demos and broke for production. The LLM hallucinated facts. The retrieval missed precise structured answers. The architecture did not fit the SAP ecosystem that customers already had invested in.

This book is my answer to that problem.

It builds a real system — not a toy — that combines two retrieval strategies that together handle what neither handles alone. It runs on platforms that SAP developers already know: SAP BTP, SAP HANA Cloud, and CAP Node.js. It uses Google Vertex AI not because it is the only option, but because it is excellent, well-documented, and free to get started

with.

Every line of code in this book has been written and tested. Every architecture decision has a reason. Where I made a tradeoff, I say so.

My hope is that you finish this book with a working system deployed on SAP BTP, a clear mental model of how agentic AI actually works, and the confidence to take these patterns into your own projects.

Let us build something real.

Sriram Rokkam *May 2026*

Who This Book Is For

This book is written for:

- **SAP developers and architects** familiar with SAP BTP, CAP, or ABAP who want to build AI-powered applications on the platform they already know
- **AI and ML engineers** working in SAP customer environments who need to understand the SAP data layer and how to integrate with it
- **Technical leads and solution architects** evaluating agentic AI patterns for SAP projects and looking for a production-ready reference architecture
- **Full-stack developers** who have experimented with LLMs and chatbots and are ready to go beyond simple Q&A into multi-step agentic systems
- **BTP developers** who want hands-on experience with SAP HANA Cloud's vector engine and SPARQL Knowledge Graph capabilities

If you work with SAP systems and want to build AI agents that reason over your enterprise data — this book is for you.

Who This Book Is NOT For

This book is not the right fit if you are:

- **Looking for a general machine learning or deep learning introduction.** We do not cover neural networks, training models, or ML theory. This book is about

applying LLMs to SAP problems, not building LLMs.

- **An ABAP developer with no Python experience.** The agent layer is written in Python. You do not need to be an expert, but you should be comfortable reading and writing basic Python functions, classes, and pip packages before starting Chapter 2.
- **Looking for a no-code or low-code AI solution.** We write real code throughout. Every component is built from first principles so you understand exactly what it does and why.
- **Expecting SAP AI Core as the primary AI backend.** We use Google Vertex AI (Gemini + text-embedding-004) because it is available on the free trial. SAP AI Core is covered in Appendix E for readers who have an enterprise subscription.
- **New to cloud development entirely.** You should be comfortable with the concept of cloud services, REST APIs, and deploying applications. Chapter 2 walks through all the setup steps, but it assumes you have used a terminal before.

Prerequisites

Technical Knowledge

The following knowledge is assumed. You do not need to be an expert — a working familiarity is enough.

Topic	Level Required	Where to Learn if Needed
Python	Basic — functions, classes, pip, venv	python.org/about/gettingstarted
JavaScript / Node.js	Basic — enough to read CAP handlers	nodejs.dev/learn
SAP BTP concepts	Familiar — subaccounts, CF, services	learning.sap.com
REST APIs	Basic — HTTP methods, JSON	Any REST API tutorial
SQL	Basic — SELECT, INSERT	W3Schools SQL

NOT Required

You do not need prior knowledge of: - Machine learning theory or mathematics - LangChain or LangGraph (we introduce it from scratch in Chapter 6) - RDF, SPARQL, or Knowledge Graphs (introduced from scratch in Chapter 4) - SAP AI Core (covered optionally in Appendix E) - ABAP development - Docker or Kubernetes

Tools and Accounts (All Free)

All tools and accounts are free. Chapter 2 walks through every setup step.

Tool / Account	Purpose	Cost
SAP BTP Trial	Runtime, HANA Cloud, CAP deployment	Free
GCP Account	Vertex AI API (Gemini + embeddings)	Free (\$300 credit)
Python 3.11+	Agent development	Free
Node.js 20+	CAP frontend development	Free
CF CLI	Deploy to BTP Cloud Foundry	Free
CDS CLI (@sap/cds-dk)	CAP development	Free
VS Code	Code editor	Free

Companion Code Repository

All source code for this book lives in a single GitHub repository:

<https://github.com/sriramrokkam/book-agentic-hybrid-rag-on-btp>

Repository Structure

book-agentic-hybrid-rag-on-btp/

```
├─ agents/           # Python FastAPI + LangGraph agent layer
│   └─ agents/       #   LangGraph chains and orchestrator
│   └─ srv/          #   HANA, vector, KG, and doc services
```

```

|   └─ main.py          # FastAPI endpoint – all HTTP endpoints
|   └─ requirements.txt
|   └─ .env.example     # Required environment variables template
└─ cap-srv/            # CAP Node.js OData V4 service
    └─ db/schema.cds    # Data model
    └─ srv/service.cds  # Service definition and actions
    └─ srv/service.js   # Action handlers (proxies to agent layer)
    └─ .env.example
└─ MSDS_Ontology.ttl    # OWL ontology constraining Knowledge Graph
└─ docs/               # Book chapters and screenshots (this folder)
    └─ CODE_MAP.md      # Chapter → source file cross-reference
    └─ chapters/        # All chapter markdown files
└─ README.md           # Complete local setup guide

```

Quick Start

```

1  # 1. Clone
2  git clone
   https://github.com/sriramrokkam/book-agentic-hybrid-rag-on-btp.git
3  cd book-agentic-hybrid-rag-on-btp
4
5  # 2. Set up Python environment
6  python3 -m venv agents/.venv
7  source agents/.venv/bin/activate
8  pip install -r agents/requirements.txt
9
10 # 3. Configure credentials
11 cp agents/.env.example agents/.env
12 # Edit agents/.env – see README.md for where to find each value
13
14 # 4. Run the agent layer
15 uvicorn agents.main:app --reload --host 0.0.0.0 --port 8000
16
17 # 5. Run the CAP layer (separate terminal)
18 cd cap-srv && npm install && cds serve

```

See `README.md` in the repository root for the complete step-by-step setup guide, including all environment variables, endpoints, and deployment instructions for SAP BTP Cloud

Foundry.

Chapter-to-Code Cross-Reference

`docs/CODE_MAP.md` maps every chapter to the exact source files and functions it covers. Use it to jump directly to the code for any chapter without searching the repository.

How to Use This Book

Reading Cover to Cover (Recommended)

Each chapter builds directly on the previous one. The system grows chapter by chapter — by Chapter 9 you have a fully deployed, working Hybrid RAG Agent on SAP BTP. Reading linearly gives you the full progression from concept to production.

As a Reference

If you already have some of the pieces (a HANA Cloud instance, some LangGraph experience), jump to the chapter you need. Each chapter opens with a summary of what it assumes you already have from previous chapters.

Chapter Structure

Every chapter in this book follows the same pattern:

1. **Opening question** — the single question this chapter answers
2. **Concept** — the idea explained in plain language before any code
3. **Code** — we build the component step by step
4. **Explanation** — what each part does and why it is designed that way
5. **Test** — how to verify it works before moving on
6. **Summary** — what we built and what comes next
7. **Checkpoint** — a checklist before starting the next chapter

The Code

All code is in the companion GitHub repository:

<https://github.com/sriramrokkam/book-agentic-hybrid-rag-on-btp>

Each chapter has a corresponding folder in the repo. If you get stuck, the repo has the complete working code for reference.

Conventions Used in This Book

Code Blocks

```
1 # Python code appears in blocks like this
2 def example():
3     return "this is a code example"
```

```
1 // JavaScript/Node.js code looks like this
2 const example = () => "this is a code example";
```

```
1 # Terminal commands look like this
2 cf push my-app
```

```
1 -- SQL queries look like this
2 SELECT * FROM MY_TABLE;
```

Callout Boxes

Note: Additional context or clarification that is helpful but not critical to follow along.

Important: Something you must do or understand before continuing. Skipping this will cause problems later.

Warning: A common mistake or pitfall. Pay attention here.

Tip: A shortcut, best practice, or time-saving suggestion.

File Paths

File paths are shown relative to the project root: - `agents/srv/hdb_srv.py` — the HANA connection service in the agents module - `cap-srv/db/schema.cds` — the CDS data model

Screenshots

Screenshots show the state of the UI at the time of writing. SAP BTP and GCP update their interfaces regularly. If a screen looks slightly different, the underlying action is the same — use the search or navigation to find the equivalent option.

[SCREENSHOT: description] markers in this draft will be replaced with actual screenshots in the final published version.

Acknowledgements

[To be added by Sriram Rokkam. Suggested: colleagues who reviewed chapters, the SAP and Google Cloud developer communities, anyone who provided feedback or inspiration.]

About the Author

Sriram Rokkam is a [title] with [X] years of experience building enterprise applications on SAP BTP and Google Cloud Platform.

[Complete this section in your own words — your current role, notable projects, SAP community involvement, LinkedIn/GitHub handles.]

End of Front Matter — Chapter 1 begins on the next page.

Chapter 1: Agentic AI on SAP BTP — What It Is, Why It Matters, and What We Build

1.1 The Agentic AI Moment

An AI agent is not a chatbot. A chatbot takes a question and returns a response — it is a single-turn input/output system. A copilot offers suggestions inline with your work, but a human makes every consequential decision. An AI agent is something categorically different: it accepts a goal, selects tools, executes a sequence of steps, evaluates intermediate results, and revises its plan — autonomously — until the goal is reached or it determines it cannot.

Four capabilities define whether a system is genuinely agentic. Understanding them precisely is worth the time before a single line of code is written.

Tool Use is the ability to call external functions — APIs, databases, search engines, file systems — and incorporate the results into the next step of reasoning. In an SAP context: imagine an agent that, when asked about the status of a purchase order, does not guess from training data but calls `API_PURCHASEORDER_PROCESS_SRV` directly, reads the live document flow, checks delivery confirmation in the GR table, and returns a current answer. The LLM is not the knowledge store — it is the reasoning engine that decides what to call and what to do with the result.

Memory is the ability to retain information across steps and across sessions. Short-term memory is the working state of a task in progress — the list of tool outputs so far, the intermediate reasoning, the partial answer. Long-term memory is a persistent store that survives restarts — a database, a vector table, a Knowledge Graph. In an SAP context: an agent that has ingested a supplier's quality certificates stores those embeddings in HANA Cloud. Next week, when a quality engineer asks about that supplier's latest batch result, the agent retrieves from long-term memory rather than asking for the document again.

Planning is the ability to decompose a goal into a sequence of steps and revise that sequence based on what intermediate steps return. A deterministic workflow executes the same path every time. A planning agent determines at runtime — based on the question, the available tools, and the results so far — what step to take next. In an

SAP context: an agent handling a supplier qualification query might first retrieve open inspection reports, then — if those reports contain anomalies — automatically invoke a second tool to check whether a corrective action request exists in QM before composing its answer.

Multi-Agent Collaboration is the ability to delegate subtasks to specialist agents and synthesise their outputs. A supervisor agent receives a complex goal and routes parts of it to agents with focused expertise: one agent that specialises in document retrieval, another that queries live ERP data, another that generates a structured report. Each specialist operates independently and returns results to the supervisor, which merges them into a coherent answer. In an SAP context: the supervisor delegates “retrieve batch quality data” to the KG agent, “retrieve supplier invoice details” to the vector agent, and “look up the purchase order” to an OData agent — then synthesises all three into a single procurement audit response.

These four capabilities are not a theoretical framework. They are the design vocabulary for the system built in this book, and they map directly to the implementation choices made at every layer of the stack.

1.2 Why SAP Developers Need to Understand Agents Now

SAP is not approaching agentic AI cautiously. At **SAP Sapphire 2026 in Orlando**, CEO Christian Klein placed agents at the centre of the company’s entire product strategy under the banner of the **Autonomous Enterprise** — a model where AI agents and human workers collaborate continuously to execute, adapt, and optimise critical business processes. Klein:

“For the mission-critical processes of our customers, ‘almost right’ just isn’t good enough. By uniting SAP Business AI Platform with SAP Autonomous Suite, we anchor AI agents in the business processes, data and governance so they can deliver accurate, compliant and secure outcomes.”

The announcements at Sapphire 2026 were not roadmap items. They were generally available or in ramp-up:

Milestone	Detail
200+ specialised agents	Available across Finance, Supply Chain, Procurement, HR, and CX — built on the same BTP agent patterns this book teaches
50+ Joule Assistants	Domain-specific AI assistants embedded across S/4HANA, Ariba, SuccessFactors, and more
SAP Autonomous Suite	New product suite embedding agents directly into core SAP applications
Joule Studio on BTP	No-code and pro-code environment for building custom agents on SAP BTP
NVIDIA OpenShell partnership	Secure runtime environment for Joule Studio agent execution
Prior Labs acquisition	European frontier AI research lab — SAP acquiring capability to control the AI stack from data layer to model layer
€100 million partner fund	Accelerating partner deployments of AI agents across the ecosystem

The implication for BTP developers is direct: the 200+ agents SAP announced were not all built inside SAP. Partners and customers on BTP built them. The patterns used — LangGraph-style orchestration, HANA Cloud retrieval, CAP as the OData surface, Joule Agent-to-Agent (A2A) for integration with the broader SAP agent ecosystem — are the same patterns this book teaches.

SAP is standardising an agent architecture. Joule Studio is the no-code surface, but below Joule Studio is a pro-code layer on Cloud Foundry that looks exactly like what you build here: a Python orchestration service, a HANA Cloud retrieval backend, a CAP service that exposes the capability as OData V4. The patterns in this book are not theoretical preparation for a future SAP platform. They are the current architecture that SAP's largest customers and partners are building on today.

An SAP developer who understands ABAP, CDS views, BTP integration, and the core SAP data model has the right foundation. What is new is the orchestration layer — how to structure an agent that uses tools rather than calling APIs directly, how to design a retrieval architecture that handles both structured and unstructured enterprise content, and how to connect that agent to the SAP surface your business users already work in. That is what this book teaches.

1.3 Why BTP Is the Right Platform for Agentic AI

A common question from BTP developers encountering agentic AI for the first time is whether they need to move off BTP to build it — whether this work belongs on a specialist ML platform, in a dedicated vector database, or behind a separate AI infrastructure layer. It does not. BTP already has what you need, and in one important dimension it has something no other cloud platform offers.

HANA Cloud is uniquely positioned. Every other cloud-native vector store is a specialised service: Pinecone, Weaviate, ChromaDB, AlloyDB pgvector. They provide semantic similarity search against embeddings. HANA Cloud provides that — the `REAL_VECTOR` column type with cosine similarity search — and also provides a full SPARQL 1.1 endpoint (`SPARQL_EXECUTE`) for RDF Knowledge Graph queries, in the same managed database instance, on the same network hop from your Cloud Foundry application. No other managed database available on a major cloud provider combines production-grade vector search with a compliant SPARQL endpoint in a single service. On BTP, this combination costs nothing beyond the trial credit to provision. You do not need a separate graph store. You do not need a separate vector database. One HANA Cloud instance handles both retrieval chains.

CAP gives you a production OData V4 service and Fiori Elements UI in minutes.

Without CAP, connecting an agent backend to a Fiori frontend requires building an OData service layer manually — schema definition, entity framework, routing, error handling, authentication. With CAP, you define your entities in CDS, annotate them for Fiori Elements, and `cds serve` gives you a functional OData V4 service. The Fiori Elements UI is generated from annotations — no custom frontend code. The CAP layer handles XSUAA authentication when you deploy to Cloud Foundry. For an SAP developer, this is already familiar territory.

The Destination Service externalises enterprise credentials properly. Connecting to S/4HANA from BTP is solved infrastructure: you configure a Destination in the BTP cockpit with the S/4HANA credentials, and your application reads the destination at runtime without hardcoding credentials. For `API_PRODUCT_SRV` — the OData service used to validate material numbers against the live product master — this means your agent never holds S/4HANA credentials directly. The Destination Service is a first-class BTP service, not a workaround.

Cloud Foundry runtime runs Python and Node.js side by side. The system in this book consists of two services: a Python FastAPI application running the LangGraph orchestrator (port 8000) and a CAP Node.js service (port 4004). Both are deployed to the same CF space as separate applications. They communicate over HTTP. CF handles scaling, restarts, environment variable injection, and service binding. The deployment in Chapter 10 uses a single MTA descriptor that captures both applications and their HANA Cloud service binding.

You do not need SAP AI Core. SAP AI Core is the enterprise AI deployment and orchestration platform on BTP — it provisions and runs foundation models at scale, manages model lifecycle, and is the inference layer behind Joule. It requires an enterprise subscription not available on a BTP trial account. This book uses Vertex AI (Google Cloud) as the LLM and embedding provider. Vertex AI provides Gemini 2.5 Flash for synthesis and text-embedding-004 for embedding generation. Both are available under the Google Cloud \$300 free credit, and the free tier is sufficient for the entire development cycle in Chapters 1 through 9. For readers with access to SAP AI Core, Appendix E covers how to swap the LLM provider; the agent architecture is unchanged.

The combination — HANA Cloud dual retrieval, CAP OData surface, Destination Service credential management, CF runtime for Python and Node.js, Vertex AI for model inference — gives a BTP developer a complete agentic AI stack that requires no additional infrastructure, no new platform to learn, and no production-tier subscription to get started. The platform you already work on is the right platform.

1.4 What This Book Builds: The Material Document Intelligence Platform

The worked example throughout this book is a **Material Document Intelligence Platform**: a system that accepts any PDF document linked to an SAP Material Number, ingests it into HANA Cloud, and answers natural-language questions about it using a LangGraph hybrid RAG agent surfaced through a Fiori Elements UI.

This platform is the teaching vehicle, not the subject of the book. Each chapter introduces a concept — vector search, Knowledge Graph retrieval, LangGraph orchestration, multi-agent supervision, CAP integration, CF deployment — and builds the corresponding component of the platform. By Chapter 10, the pieces form a running system.

The anchor: SAP Material Number

Every document in the system is linked to an SAP Material Number — the unique identifier that appears in MM03, in goods movements, in quality inspection lots, in purchase orders. Before a document is accepted for ingestion, the material number is validated against `API_PRODUCT_SRV`, SAP's standard product OData V4 service. This validation confirms the material exists in your S/4HANA product master and returns basic product attributes (material description, base unit, material group). The system will not ingest a document against a material number that does not exist in SAP. The anchor is live data, not a local list.

Document types supported

The platform is designed to handle any PDF document type that an SAP user would associate with a material. The reference implementation covers five document types:

- **Batch Quality Certificates** — supplier-issued documents confirming test results against specification for a specific production batch
- **Goods Receiving Inspection Reports** — QM-generated reports recording the outcome of incoming goods inspection, including usage decision and lot status
- **Equipment Maintenance History** — service records linked to technical objects (equipment numbers) associated with materials in plant maintenance
- **Supplier Invoices** — AP documents linking vendor, line items, purchase order references, and payment terms
- **MSDS (Material Safety Data Sheets)** — regulatory documents present in chemical and pharma supply chains; used as one document type in the reference im-

plementation

MSDS is the document type used for the reference implementation. The HANA Cloud vector table is named `MSDS_VECTORS`. The HANA RDF named graph is `MSDS_Graph`. Test data, sample queries, and the ontology in Appendix B are all MSDS-based. This is a naming artifact from the reference implementation, not a constraint on the architecture. The ingestion pipeline, agent orchestration, and retrieval logic are document-type agnostic. Changing the Knowledge Graph extraction prompt — the instruction that tells Gemini 2.5 Flash what triples to extract — is all that is required to apply the same system to batch certificates, inspection reports, or maintenance records. The table naming is explained once, accepted throughout, and not revisited.

Ingestion: parallel vector and graph

When a user uploads a PDF against a material number, the FastAPI backend runs two ingestion tasks in parallel using `ThreadPoolExecutor`:

The vector ingestion task parses the PDF, chunks the text at 512 tokens with 50-token overlap, and embeds each chunk using `text-embedding-004`. Chunk text and embeddings are inserted into `MSDS_VECTORS`, a HANA Cloud column table with a `REAL_VECTOR` column. The HANA cosine similarity index operates over this column.

The Knowledge Graph ingestion task sends the full document text to Gemini 2.5 Flash with a structured extraction prompt. The model extracts RDF triples — entities, relationships, and attribute values — and returns them as Turtle syntax. The triples are written to HANA Cloud using `SPARQL_EXECUTE` into the named graph `GRAPH <iri:msds-graph/MAT-{material}>`. Each material gets its own named graph, isolating its triples from other materials' data.

Both tasks complete before the ingestion response returns to the CAP layer. The CAP service logs the ingestion outcome and returns a confirmation to the Fiori UI.

Query: parallel retrieval, merge, synthesise

When a user submits a question through the Fiori chat UI, the CAP service proxies it to the FastAPI backend. The LangGraph supervisor node initialises a `StateGraph` and fans out to both retrieval chains in parallel:

The Vector Chain embeds the question using `text-embedding-004`, runs cosine similarity search against `MSDS_VECTORS` filtered to the target material, and returns the top-k matching chunks with similarity scores.

The KG Chain converts the question into a SPARQL SELECT query, executes it against SPARQL_EXECUTE on the material's named graph, and returns structured triples.

The supervisor merge node receives both result sets. It formats the vector chunks and graph triples as a combined context block and calls Gemini 2.5 Flash with a synthesis prompt that requires the model to cite the source of every claim — either by chunk reference (vector) or by triple predicate (graph). The cited answer is returned to the Fiori UI.

Architecture: two services, one platform

The system consists of two services communicating over HTTP:

- **FastAPI agent** (`uvicorn main:app`, port 8000): the LangGraph orchestrator, ingestion pipeline, HANA Cloud connections (thread-local), and all LLM calls
- **CAP service** (`cds serve`, port 4004): OData V4 entities, Fiori Elements UI, S/4HANA material number validation via Destination Service, proxying to the FastAPI back-end

In local development, both run on your machine and connect to the HANA Cloud trial instance over JDBC. In Chapter 10, both are deployed to the same BTP Cloud Foundry space — CAP as an MTA, FastAPI as a Python CF application — and the HANA Cloud service binding is injected at runtime via `VCAP_SERVICES`.




Figure 1: Agentic Hybrid RAG System Architecture

Figure 1.1 — Material Number is the anchor. Documents are ingested into MSDS_VECTORS (vector) and MSDS_Graph/MAT-{material} (RDF) in parallel. Queries fan out to both chains, merge, and synthesise a cited answer via Gemini 2.5 Flash.

1.5 Why Hybrid Retrieval

The hybrid retrieval architecture — running vector search and SPARQL graph queries in parallel and merging the results — is not a design preference. It is the correct response to the structure of enterprise documents. Understanding why requires looking at what enterprise documents actually contain.

Enterprise documents contain two fundamentally different kinds of content.

The first kind is structured facts: discrete, precise values that have a single correct answer. A batch quality certificate contains the certificate number, the batch lot number, the material specification reference, the measured test values for each test parameter, the pass/fail result for each test, the certifying laboratory, and the certification date. A goods receiving inspection report contains the inspection lot number, the usage decision code (accept/reject/block), the quantity inspected, the sampling procedure, and the inspection date. An equipment maintenance record contains the functional location, the equipment number, the service action code, the work order number, the technician ID, and the completion date.

For this kind of content, the right retrieval mechanism is exact lookup. When a quality engineer asks “What is the certificate number for batch BATCH-2024-0871?” the answer is a specific identifier — one value, no ambiguity. Vector search is the wrong tool for this question. A cosine similarity search against text embeddings will return the chunk with the highest semantic similarity to the query, which might be around 0.65 — a correct-seeming paragraph that mentions certificate numbers in general without necessarily containing the specific number for that specific batch. There is no guarantee of precision. The exact fact may be embedded in a low-similarity chunk because the surrounding text is not thematically similar to the query.

The SPARQL Knowledge Graph answers this question correctly and deterministically. The batch quality certificate’s triples include `<BATCH-2024-0871> qm:certificateNumber "CERT-2024-58291"`. A SPARQL SELECT query against the named graph for that mate-

rial returns the value in one round-trip. No ambiguity, no similarity scoring, no chunk retrieval.

The second kind of content is narrative prose: sentences and paragraphs written in natural language that convey meaning through context, not through discrete values. A batch quality certificate’s methodology section describes how each test was conducted — the equipment used, the test conditions, the standard procedure referenced, the inspector’s observations about any anomalies. A goods receiving inspection report includes the inspector’s narrative description of the physical condition of the goods, observations about packaging integrity, and the qualitative rationale for the usage decision. An equipment maintenance record includes the technician’s description of the fault found, the diagnosis, and the repair approach.

For this kind of content, the right retrieval mechanism is semantic similarity search. When a quality engineer asks “How did the supplier describe the tensile strength testing method on this batch?” the answer is a passage of prose. SPARQL cannot answer it — there is no triple that captures a methodology description without losing the meaning carried by the narrative. Vector search finds it correctly: the question’s embedding is semantically close to the chunk that describes the test methodology, and cosine similarity search returns it.

HANA Cloud is the only managed database on BTP that provides both.

REAL_VECTOR cosine similarity search handles the narrative questions. SPARQL_EXECUTE handles the structured fact questions. Both operate against the same HANA Cloud instance — the same network hop from your Cloud Foundry application, the same connection pool, the same service binding. No separate vector store provisioned alongside a separate graph store. No synchronisation between two external systems. No additional BTP service to configure.

The hybrid approach at runtime is straightforward: run both chains in parallel, merge the results, instruct the synthesis model to cite the source of each claim. When both chains return results, the answer draws on both. When only one chain returns results — because the question is purely factual or purely narrative — the answer draws on that chain. When neither chain returns results above a confidence threshold, the agent acknowledges it could not find the answer rather than generating a plausible-sounding fabrication.

That is the architecture, and Chapter 7 builds it in full.

1.6 The Four Agentic Capabilities in This System

Section 1.1 defined four capabilities that distinguish an AI agent from a deterministic retrieval pipeline. Here is how each maps to the actual implementation built in this book.

Tool Use

The LangGraph orchestrator calls HANA Cloud as an external tool — twice, via two different interfaces. The KG Chain issues a `SPARQL_EXECUTE` call against the named graph for the target material. The Vector Chain issues a `REAL_VECTOR` cosine similarity query against `MSDS_VECTORS`. Neither call is hardcoded into a fixed pipeline. The LangGraph `StateGraph` decides at runtime, based on the structure of the question and the routing logic encoded in the conditional edges, which tools to call and in what configuration. When a question is structured as a precise fact lookup, the KG Chain is weighted more heavily. When a question is framed as an open descriptive query, the Vector Chain carries more weight. The LLM does not serve as the knowledge store — it calls the right tool and synthesises from what the tool returns.

Memory

Memory in this system operates at two timescales. Long-term memory is HANA Cloud: the `MSDS_VECTORS` table persists vector embeddings across sessions, across restarts, and across users. The named RDF graphs in `MSDS_Graph` persist triples for every ingested document. Once a batch quality certificate is ingested for material 800021, its vector chunks and its triples are available to every subsequent query — no re-ingestion, no re-embedding, no warm-up time. Short-term memory is the LangGraph `StateGraph`: for the duration of a single query, the state object holds the conversation history, the intermediate tool outputs from both retrieval chains, the merge strategy output, and the synthesis context. When the query completes, the state is discarded. Long-term memory is HANA. Short-term memory is the graph state.

Planning

The LangGraph `StateGraph` routes each query at runtime through conditional edges. The supervisor node evaluates the incoming question and the results returned by the retrieval chains before deciding what to do next. A question structured as “What is the

usage decision for inspection lot IL-2024-0091?” routes primarily through the KG Chain — it is asking for a discrete value that maps to a triple. A question structured as “Describe the condition of the packaging as observed during incoming inspection” routes primarily through the Vector Chain — it is asking for narrative prose. A question that combines both — “What was the usage decision, and what did the inspector observe about the packaging?” — fans out to both chains and merges the results. The routing is not a hard-coded `if/else` block. It emerges from the conditional edge logic and the confidence scores returned by each chain. Chapter 6 builds the `StateGraph` and its routing logic from scratch.

Multi-Agent Collaboration

The supervisor pattern in Chapter 8 implements multi-agent collaboration explicitly. The supervisor node is the orchestrating agent. It delegates to two specialist sub-agents: the Vector Retrieval Agent, which owns the `MSDS_VECTORS` table interaction and the embedding pipeline, and the Knowledge Graph Agent, which owns the HANA SPARQL endpoint and the triple extraction logic. Each specialist has its own HANA connection (thread-local, to avoid connection sharing across threads), its own tool set, and its own result format. The supervisor collects the results from both specialists, applies the merge strategy, and passes the combined context to Gemini 2.5 Flash for synthesis. The architecture is extensible: adding a third specialist — for example, an OData agent that retrieves live inventory data from S/4HANA via `API_MATERIAL_STOCK_SRV` — requires adding one node and one conditional edge to the graph. The supervisor pattern does not change.

1.7 Scope and Prerequisites

What This Book Covers

Topic	Coverage
Agentic AI concepts — tool use, memory, planning, multi-agent collaboration	Full coverage, Ch 1

Topic	Coverage
SAP BTP trial setup — HANA Cloud, CF space, Destination Service	Full coverage, Ch 2
Vector search on HANA Cloud (REAL_VECTOR, cosine similarity)	Full coverage, Ch 3
Knowledge graphs on HANA Cloud (SPARQL/RDF, named graphs)	Full coverage, Ch 4
PDF ingestion pipeline — parallel KG + vector with ThreadPoolExecutor	Full coverage, Ch 5
LangGraph — StateGraph, nodes, conditional edges, routing	Full coverage, Ch 6
Hybrid RAG agent — parallel chains, merge strategy, synthesis	Full coverage, Ch 7
Multi-agent supervisor pattern — specialist sub-agents	Full coverage, Ch 8
CAP Node.js OData V4 service + Fiori Elements UI	Full coverage, Ch 9
Deploying to SAP BTP Cloud Foundry (MTA + cf push)	Full coverage, Ch 10
Joule A2A integration (enterprise subscription required)	Appendix D
SAP AI Core as LLM alternative	Appendix E
MSDS ontology design (OWL/Turtle)	Appendix B
HANA Cloud SPARQL reference	Appendix C

What This Book Does Not Cover

Production hardening. Rate limiting, multi-tenant isolation, secrets rotation, disaster recovery, and SLA monitoring are not covered. This book teaches architecture and implementation patterns, not operations.

Model training or fine-tuning. All LLM usage is inference-only — calling hosted mod-

els via API. No training pipelines, no fine-tuned models, no custom embeddings.

Cross-material reasoning. The agent answers questions about one material’s documents at a time. Cross-material aggregation — for example, “Which materials in this plant have open corrective action requests in the past 90 days?” — is a natural extension not covered here.

Streaming responses. Answers are returned synchronously as complete text. Streaming via Server-Sent Events is not implemented.

CI/CD pipelines. Chapter 10 uses manual `cf push` and `mbt build && cf deploy`. Automated deployment pipelines are not covered.

Joule and SAP AI Core on free trial. Both require paid SAP subscriptions. They are in Appendices D and E so enterprise readers can follow along without blocking the main chapters.

Prerequisites

Requirement	Where to Get It	Cost
SAP BTP Trial account	account.hanatrial.ondemand.com	Free
SAP HANA Cloud trial instance	BTP Trial cockpit → HANA Cloud tile	Free (30-day)
GCP account with Vertex AI enabled	cloud.google.com/free	Free (\$300 credit)
Python 3.11+	python.org	Free
Node.js 20+ and npm	nodejs.org	Free
CF CLI v8+	CF documentation	Free
@sap/cds CLI (npm i -g @sap/cds)	npm	Free
VS Code with SAP CDS extension	code.visualstudio.com	Free

Chapter 2 covers every setup step with screenshots. If you already have an active BTP trial and GCP project, setup takes roughly 30 minutes.

Local development is the primary mode for Chapters 1–9. You run `uvicorn main:app --reload` for the FastAPI agent and `cds serve` for the CAP service on your local machine, both connecting to your BTP HANA Cloud trial instance over JDBC.

Chapter 10 deploys the same code to Cloud Foundry — nothing changes except where it runs.

1.8 Summary

- An AI agent is distinguished from a chatbot or copilot by four capabilities: tool use, memory, planning, and multi-agent collaboration. These are not theoretical distinctions — they map directly to specific implementation choices in the LangGraph orchestration, HANA Cloud retrieval, and supervisor pattern built in this book.
- SAP announced the Autonomous Enterprise at Sapphire 2026: 200+ specialised agents, Joule Studio on BTP, the NVIDIA OpenShell partnership, and the Prior Labs acquisition. The agent architecture patterns in this book — LangGraph orchestration, HANA Cloud retrieval, CAP as OData surface, Joule A2A — are the same patterns SAP is standardising across its platform.
- SAP BTP is the right platform for agentic AI development because HANA Cloud provides vector search and SPARQL graph queries in a single managed service — a combination that no other managed database on a major cloud platform offers. CAP, the Destination Service, and Cloud Foundry complete the stack. Vertex AI fills the model inference role at zero additional cost on the trial — for enterprise readers with SAP AI Core, Appendix E shows how to swap the provider without changing the agent architecture.
- The worked example is a Material Document Intelligence Platform: upload PDFs against an SAP Material Number validated via `API_PRODUCT_SRV`, ingest to `MSDS_VECTORS` and `MSDS_Graph` in parallel via `ThreadPoolExecutor`, query both with a LangGraph hybrid RAG agent, synthesise a cited answer via Gemini 2.5 Flash, surface it through Fiori Elements. The document types supported include Batch Quality Certificates, GR Inspection Reports, Equipment Maintenance History, Supplier Invoices, and MSDS. Table names use MSDS as a naming convention from the reference implementation.
- Hybrid retrieval is the correct architecture for enterprise documents because those documents contain two kinds of content that require different retrieval mechanisms. Structured facts — certificate numbers, batch lot numbers, test results, usage decisions, inspection dates — need SPARQL exact lookup via `SPARQL_EXECUTE`. Narrative prose — methodology descriptions, inspector observations, handling in-

structions, qualitative assessments — needs vector cosine similarity search against `REAL_VECTOR`. HANA Cloud provides both in one instance.

- The four agentic capabilities map directly to the implementation: tool use is `SPARQL_EXECUTE` and `REAL_VECTOR` calls made by the `LangGraph` agent; memory is HANA Cloud long-term storage and `StateGraph` short-term context; planning is the conditional edge routing logic in the `StateGraph`; multi-agent collaboration is the supervisor pattern that delegates to `Vector Agent` and `KG Agent` and merges their results.

Chapter 2 sets up the platform: BTP trial, HANA Cloud instance, GCP project, and Vertex AI credentials. By the end of Chapter 2 you have a working environment and can run your first HANA vector insert.

Chapter 2: Platform Setup — SAP BTP Trial + Google Vertex AI

2.1 The Goal of This Chapter

Before we write a single line of agent code, we need a working environment. This chapter gets you there.

By the end of this chapter, you will have:

- A running SAP BTP trial account with a HANA Cloud instance
- A GCP project with Vertex AI enabled and a service account ready
- A BTP Destination that securely connects your BTP applications to Vertex AI
- A local development environment with all required tools installed
- A confirmed, working API call to Gemini 2.5 Flash from your local machine

This is pure setup. There is no agent code here. But every chapter that follows depends on this foundation being solid. Rushing this chapter and skipping verification steps is the most common cause of confusing errors later. Take the 45 minutes. Do it properly.

Cost: SAP BTP trial is free. GCP provides \$300 credit on signup — the Vertex AI API calls for this entire project cost a few dollars at most. The same setup steps apply to a paid BTP account; the only difference is the account URL.

2.2 Understanding What We Are Setting Up

Before clicking through consoles, it helps to understand the shape of what we are building.

Our system has three infrastructure components:



2.2 Understanding What We Are Setting Up




Figure 2.2 — Your local machine develops and deploys to SAP BTP Trial. The BTP Cloud Foundry space runs the FastAPI agent and CAP Node.js service. HANA Cloud holds the `REAL_VECTOR` table and SPARQL/RDF graphs. The Destination Service stores GCP credentials. All Vertex AI calls go outbound over HTTPS only.

SAP HANA Cloud is the **knowledge layer** — it stores both our vector embeddings and our RDF Knowledge Graph in a single managed service. No other database on the market provides both `REAL_VECTOR` column types (for cosine similarity search) and a native SPARQL execution engine in one service. This is a significant architectural advantage — one connection, one credential, two retrieval strategies.

Google Vertex AI is the **AI inference layer** — it is an external HTTPS API that our BTP application calls like any other REST service. Gemini 2.5 Flash handles both our LLM

tasks (SPARQL generation, answer synthesis) and we use `text-embedding-004` for generating embeddings. The BTP Destination Service stores the GCP credentials securely — the service account JSON key never lives in application code or environment variables directly.

With this picture in mind, let us set it up.

2.3 SAP BTP Trial Account

Signing Up

Navigate to <https://www.sap.com/products/technology-platform/trial.html> and click **Try now**.




Figure 2.1 — SAP BTP Free Trial landing page — click “Try now” to start your free trial account

Fill in your details. Use a personal email address if possible — corporate email addresses sometimes cause issues with trial account activation. After verifying your email, BTP will automatically provision a trial account in either the EU10 (Frankfurt) or US10 (Virginia) region depending on your location.

Note: The region matters. Your HANA Cloud instance must be in the same region as your Cloud Foundry space. BTP assigns both automatically — do not change the region during signup.

Once logged in, you will land on the **BTP Cockpit** — the central control panel for all BTP services.




Figure: SAP BTP Cockpit home screen showing your Global Account and trial subaccount

Navigating to Your Trial Subaccount

The BTP Cockpit has a hierarchy: Global Account → Subaccount → Space. Your trial comes with one subaccount called “trial.” Click it.




Figure: BTP trial subaccount overview page

Inside the subaccount, you will see your **Cloud Foundry environment**. Note your CF

API endpoint — it will look like `https://api.cf.eu10.hana.ondemand.com`. You will need this later.

Click **Cloud Foundry** ▢ **Spaces**. You should see a space called “dev.” This is where you will deploy your applications.




Figure: Cloud Foundry space view — note your org name and the “dev” space

Write these down now: - CF API Endpoint: `https://api.cf.<region>.hana.ondemand.com`
 - CF Org: `<your-trial-org>` - CF Space: `dev`

2.4 Provisioning SAP HANA Cloud

HANA Cloud is provisioned from the BTP Cockpit. This is a one-time setup that takes about 10–15 minutes.

Creating the HANA Cloud Instance

From your trial subaccount, go to **Services** ▢ **Service Marketplace**. Search for “HANA Cloud.”




Figure: BTP Service Marketplace — search for “HANA Cloud” to find the tile

Click the HANA Cloud tile **Create**. You will be taken to the HANA Cloud provisioning wizard.

Fill in the configuration:

Field	Value
Instance Name	hana-trial
Administrator Password	Choose a strong password (save it!)
Memory	30 GB (the trial default — do not reduce)
Storage	120 GB (the trial default)




Figure: HANA Cloud instance creation wizard — set instance name and keep default

30GB memory

Click **Create**. Provisioning takes approximately 10–15 minutes. You will see the instance appear in the SAP HANA Cloud Central tool with a “Starting” status.




Figure: SAP HANA Cloud Central showing the instance provisioning status

Wait until the status shows **Running** before proceeding.

Enabling the Vector Engine

By default, HANA Cloud provisions with the vector engine enabled in trial. To confirm, open **SAP HANA Cloud Central** → click your instance → **Manage HANA Cloud** → **Edit**.

Under **Additional Features**, verify that **Script Server** and **Document Store** are enabled. The vector engine (REAL_VECTOR column type) is part of the core engine — no separate flag is needed.




Figure: HANA Cloud Central — click your instance to manage settings and verify vector engine availability

Enabling the Knowledge Graph Engine

The RDF/SPARQL engine in HANA Cloud is the **Triple Store** — and it is not enabled by default. Without it, all SPARQL queries in Chapters 4, 5, 7, and 8 will fail silently. This is the single most common setup mistake.

The correct feature to enable is Triple Store — not Document Store.

You can enable it on a running instance at any time — no delete or recreate required. In HANA Cloud Central:

1. Click your instance name ▢ the dropdown chevron next to the instance name ▢
Manage Configuration
2. Select the **Advanced Settings** tab
3. Under **Additional Features**, check **Triple Store**
4. Optionally check **Script Server** if you plan to use other stored procedures
5. Click **Save** — HANA Cloud applies the change without downtime




Figure 2.x — HANA Cloud Central ▢ your instance ▢ Manage Configuration ▢ Advanced Settings. Enable Triple Store (the RDF/SPARQL engine). Script Server is separate. Document Store is for JSON documents — not required for this book.

Important: The info banner on this screen says “Adding features to an instance may require more memory and increase your licensing cost.” On a trial instance, enabling Triple Store is free within the trial allocation. Monitor your CU usage in HANA Cloud Central if you are concerned.

Note: Natural Language Processing (NLP) is shown as enabled in the screenshot above — that is from a separate training instance. You do not need NLP enabled for this book.

Getting Your HANA Connection Details

From SAP HANA Cloud Central, click your instance ▢ **Actions** ▢ **Copy SQL Endpoint**.

This gives you a connection string like:

```
<instance-id>.hana.trial-us10.hanacloud.ondemand.com:443
```

Save these: - **Host:** <instance-id>.hana.trial-<region>.hanacloud.ondemand.com

- **Port:** 443 - **User:** DBADMIN - **Password:** The password you set during provisioning

Warning: The DBADMIN user has full database privileges. For production, create a dedicated application user with limited permissions. For this book’s

trial environment, DBADMIN is acceptable.

2.5 Setting Up Google Cloud + Vertex AI

Vertex AI is the model inference layer for this book — Gemini 2.5 Flash for LLM tasks, `text-embedding-004` for vector embeddings. We use it because it is immediately accessible with a Google account and \$300 free credit, making it available to every reader regardless of SAP subscription level. The agent architecture is provider-agnostic; Appendix E shows how to swap in SAP AI Core for enterprise deployments where that is the preferred inference platform.

Creating a GCP Account

Go to `cloud.google.com/free` and click **Get started for free**.

Sign in with your Google account.

Important — Credit Card Required: GCP will ask for credit card details during signup for identity verification. A charge will not be made immediately, but a card is mandatory to activate the \$300 free credit. Read GCP's terms and conditions carefully before proceeding — billing behaviour, credit expiry, and auto-upgrade conditions are your responsibility to understand.

Caution — Use Your \$300 Credit Wisely: The exercises in this book are well within the \$300 free credit — typical usage for the full project (embeddings, Gemini API calls, SPARQL generation) costs a few dollars at most. However, once you complete this project, **delete your GCP project and suspend or close your GCP account within 90 days**. An idle project with APIs enabled can accrue unexpected charges. Cleaning up after the exercise is not optional — it is good practice. Steps to delete your project are in Appendix B. Managing your GCP account beyond this book's exercises is outside the scope of this book.

After signup, you will land on the **GCP Console**.

Creating a New Project

Click the project selector at the top of the page ▢ **New Project**.

Field	Value
Project Name	agentic-rag-btp
Project ID	Auto-generated (note this down)
Organization	No organization (personal account)

Note: In GCP Console, click the project selector dropdown ▾ “New Project”
▾ enter a project name (e.g. hybrid-rag-btp) ▾ Create.

Click **Create**. Wait a few seconds for the project to be created, then select it from the project selector.

Enabling the Vertex AI API

In the GCP Console, go to **APIs & Services** ▾ **Library**. Search for “Vertex AI API.”

Note: Navigate to APIs & Services ▾ Library ▾ search “Vertex AI API” ▾ click Enable.

Click the Vertex AI API tile ▾ **Enable**.

Note: After enabling, the Vertex AI API page will show a green checkmark and “API enabled” status.

Enabling the API takes about 30 seconds. Once enabled, you will see the Vertex AI API dashboard.

Note: Enabling Vertex AI automatically enables several dependent APIs (Cloud Storage, Cloud Resource Manager, etc.). This is expected.

GCP Authentication — What Actually Works

Use a personal Google account for this book. Sign up at cloud.google.com/free with a personal Gmail — not your corporate Google Workspace account. Corporate GCP organisations typically block service account key creation via org policy, which will prevent the steps below from working.

This book uses two authentication approaches depending on context:

Context	Method
Local development (Chapters 2–9)	Application Default Credentials via gcloud CLI
BTP Cloud Foundry deployment (Chapter 10)	Service account JSON key via BTP Destination Service

Creating a Service Account

Go to **IAM & Admin** ▸ **Service Accounts** ▸ **Create Service Account**.

Field	Value
Service Account Name	agentic-rag-btp-sa
Service Account ID	Auto-generated
Description	Service account for Agentic Hybrid RAG on SAP BTP

Click **Create and Continue**. In the role assignment step, add:

Role	Purpose
Vertex AI User	Calling Vertex AI APIs (Gemini LLM + text-embedding-004)

Click **Continue** ▸ **Done**.

Option A — Local Development: Application Default Credentials (ADC)

For Chapters 2–9, running the FastAPI agent locally, use ADC. No JSON key file is downloaded or stored.

Install the gcloud CLI from cloud.google.com/sdk/docs/install, then run:

```
1 gcloud auth application-default login
2 gcloud config set project YOUR_PROJECT_ID
```

Your browser opens a Google login page. Sign in with your personal Google account. The gcloud CLI stores credentials at:

```
~/ .config/gcloud/application_default_credentials.json # macOS/Linux
```

```
%APPDATA%\gcloud\application_default_credentials.json # Windows
```

The Vertex AI Python SDK picks these up automatically — no `GOOGLE_APPLICATION_CREDENTIALS` environment variable needed. Your `.env` file only needs:

```
1 GOOGLE_CLOUD_PROJECT=your-project-id
2 GCP_LOCATION=us-central1
```

How ADC works: The Vertex AI SDK checks for credentials in this order:

(1) `GOOGLE_APPLICATION_CREDENTIALS` env var, (2) ADC file from `gcloud auth application-default login`, (3) attached service account metadata. For local development, step 2 handles it automatically. Official reference: cloud.google.com/docs/authentication/application-default-credentials

Option B — BTP Cloud Foundry Deployment: Service Account JSON Key

When you deploy to BTP CF in Chapter 10, the container has no `gcloud` CLI and no access to your local ADC file. The Vertex AI credentials are provided via a service account JSON key stored securely in the BTP Destination Service.

To generate the JSON key: in GCP Console → **IAM & Admin** → **Service Accounts** → click your service account → **Keys** tab → **Add Key** → **Create new key** → **JSON**. A `.json` file downloads to your machine.

Note: If key creation is blocked by an org policy (`iam.disableServiceAccountKeyCreation`), this means you are using a corporate GCP organisation. Use a personal Gmail account as described above — personal accounts do not have this restriction.

Keep this file safe. You will use its contents to configure the BTP Destination in the next section. Never commit it to git — it is listed in `.gitignore` already.

2.6 Configuring the BTP Destination Service

The BTP Destination Service is the secure credential store that connects your BTP applications to external services. Your FastAPI agent reads the Vertex AI credentials from

this destination at runtime — the service account key never lives in application code or in `manifest.yml`.

This is the SAP-native way to manage external service credentials. In BTP enterprise architecture, the Destination Service is the standard pattern for externalising all external API credentials — your agent switches AI providers by changing a destination configuration in the BTP Cockpit, not by touching application code or triggering a redeployment. The same pattern is used for S/4HANA connectivity, external REST APIs, and OAuth token exchange across the BTP portfolio.

Creating the Destination

In the BTP Cockpit, go to your trial subaccount ▢ **Connectivity** ▢ **Destinations** ▢ **New Destination**.




Figure 2.6 — BTP Cockpit ▢ Connectivity ▢ Destinations. Click **Create** ▢ **From Scratch** to open the destination configuration form.

The correct authentication type is **OAuth2ClientCredentials**. BTP will fetch a short-lived access token from Google’s token endpoint before each Vertex AI API call. The values come directly from your service account JSON key file.

Fill in the destination configuration:

Field	Value
Name	VertexAI
Type	HTTP

Field	Value
Description	Google Vertex AI API (Gemini + Embeddings)
URL	https://us-central1-aiplatform.googleapis.com
Proxy Type	Internet
Authentication	OAuth2ClientCredentials
Client ID	client_email value from your JSON key file
Client Secret	private_key value from your JSON key file
Token Service URL	token_uri value from your JSON key file (typically https://oauth2.googleapis.com/token)

Click **New Property** to add these additional properties:

Property	Value
google.project_id	project_id value from your JSON key file
google.private_key_id	private_key_id value from your JSON key file
google.client_email	client_email value from your JSON key file

Where to find these values: Open your downloaded `gcp-sa-key.json` in a text editor. It is a JSON object with fields including `client_email`, `private_key`, `token_uri`, `project_id`, and `private_key_id`. Map each field directly to the destination field above. The `private_key` is the multi-line RSA key — paste the entire value including the `-----BEGIN/END PRIVATE KEY-----` markers.

Click **Save**, then click **Check Connection**. A 200 OK response confirms the OAuth2 token exchange is working and the Vertex AI endpoint is reachable.

Why OAuth2ClientCredentials, not NoAuthentication? The Vertex AI

API requires a valid Google OAuth2 bearer token on every request. With `OAuth2ClientCredentials`, BTP handles token acquisition and refresh automatically — your application code never manages token lifecycle. With `NoAuthentication`, you would have to implement token exchange yourself in the Python agent, which is both error-prone and a security anti-pattern.

2.7 Local Development Environment

Now set up your local machine. We need Python, Node.js, the CF CLI, the CDS CLI, and VS Code.

Python 3.11+

Verify your Python version:

```
1 python3 --version
2 # Should output: Python 3.11.x or 3.12.x
```

If you need to install Python, download it from python.org/downloads. On macOS with Homebrew:

```
1 brew install python@3.11
```

Node.js 20+

Verify:

```
1 node --version    # Should output: v20.x.x or higher
2 npm --version     # Should output: 10.x.x or higher
```

Install from nodejs.org if needed. On macOS:

```
1 brew install node@20
```

Cloud Foundry CLI

The CF CLI is how you deploy your application to BTP Cloud Foundry.

```

1 # macOS (Homebrew)
2 brew install cloudfoundry/tap/cf-cli@8
3
4 # Verify
5 cf --version
6 # Cloud Foundry CLI version 8.x.x

```

For other platforms, download from the CF CLI releases page.

Why Cloud Foundry? This book deploys to BTP Cloud Foundry because it is the simplest path to a running service — `cf push`, no container registry, no Kubernetes manifests. For production-grade agentic architectures requiring horizontal pod scaling, multi-agent process isolation, or A2A communication across microservices, **Kyma runtime** (BTP's managed Kubernetes environment) is the right deployment target. That is a separate book.

Log in to your BTP CF environment:

```

1 cf login -a https://api.cf.<your-region>.hana.ondemand.com
2 # Enter your BTP email and password
3 # Select the "trial" org and "dev" space

```

Note: After `cf login`, you will be prompted to select your org (your trial org name) and space (dev). Select the correct ones and press Enter.

Verify the login:

```

1 cf target
2 # Should show:
3 # API endpoint: https://api.cf.eu10.hana.ondemand.com
4 # API version: 3.x.x
5 # user: your@email.com
6 # org: your-trial-org
7 # space: dev

```

SAP CDS CLI

The CDS CLI is the development tool for CAP Node.js applications.

```

1 npm install -g @sap/cds-dk
2
3 # Verify
4 cds --version
5 # @sap/cds-dk: 8.x.x

```

VS Code Extensions

Open VS Code and install these extensions:

Extension	Publisher	Purpose
SAP CDS Language Support	SAP	CDS syntax highlighting and validation
Python	Microsoft	Python IntelliSense and debugging
REST Client	Humao	Test API endpoints without Postman

Tip: In VS Code, press `Cmd+Shift+X` to open Extensions. Search for each extension name and click Install.

2.8 Project Setup

Clone the companion repository and set up the Python environment for the agent service:

```

1 git clone https://github.com/sriramrokkam/agent-rag-btp-book
2 cd agent-rag-btp-book/agents
3
4 python3 -m venv .venv
5 source .venv/bin/activate # Windows: .venv\Scripts\activate
6
7 pip install -r requirements.txt

```

The `requirements.txt` for the agent service contains:

```
1 fastapi==0.111.0
2 uvicorn[standard]==0.30.1
3 langgraph==0.2.0
4 langchain-google-vertexai==1.0.6
5 google-auth==2.29.0
6 hdbcli==2.21.28
7 PyMuPDF==1.24.5
8 python-multipart==0.0.9
9 httpx==0.27.0
```

Create your local environment file:

```
1 cp .env.example .env
```

Edit `.env` with your credentials:

```
1 # SAP HANA Cloud
2 HANA_HOST=<your-instance>.hana.trial-<region>.hanacloud.ondemand.com
3 HANA_PORT=443
4 HANA_USER=DBADMIN
5 HANA_PASSWORD=<your-hana-password>
6
7 # Google Vertex AI (local development only)
8 # On BTP, credentials come from the Destination Service
9 GCP_PROJECT_ID=agentic-rag-btp
10 GCP_LOCATION=us-central1
11 GOOGLE_APPLICATION_CREDENTIALS=/path/to/gcp-sa-key.json
12
13 # Backend URL (for CAP service to call the FastAPI agent)
14 BACKEND_URL=http://localhost:8000
```

Note: The `GOOGLE_APPLICATION_CREDENTIALS` environment variable is only used for local development. When deployed to BTP Cloud Foundry, the agent reads credentials from the Destination Service instead. We will see this credential-switching logic in Chapter 9.

2.9 Verifying the Setup: Your First Gemini API Call

Let us verify that everything works end-to-end. We will call Gemini 2.5 Flash directly and confirm we get a response.

Create a file `test_vertex.py` in your `agents/` directory:

```
1  import os
2  from dotenv import load_dotenv
3  import vertexai
4  from vertexai.generative_models import GenerativeModel
5
6  load_dotenv()
7
8  # Initialize Vertex AI with project and location
9  vertexai.init(
10     project=os.getenv("GCP_PROJECT_ID"),
11     location=os.getenv("GCP_LOCATION"),
12 )
13
14 # Load Gemini 2.5 Flash
15 model = GenerativeModel("gemini-2.5-flash-preview-05-20")
16
17 # Send a simple test prompt
18 response = model.generate_content(
19     "You are a helpful assistant for SAP developers. "
20     "In one sentence, explain what a batch quality certificate is in SAP
21     supply chain management."
22 )
23 print("Gemini response:")
24 print(response.text)
25 print("\nModel used: gemini-2.5-flash-preview-05-20")
26 print("Setup verified successfully.")
```

Run it:

```
1  python test_vertex.py
```

You should see output like:

Gemini response:

A batch quality certificate is a supplier-issued document that records the test results, a

Model used: gemini-2.5-flash-preview-05-20

Setup verified successfully.

Expected terminal output:

VertexAI initialized successfully

Model loaded: gemini-2.5-flash

Response: The capital of France is Paris.

Vertex AI connection: OK

If you see an error: The most common issues are: - google.auth.exceptions.DefaultCredentialsError — Check that GOOGLE_APPLICATION_CREDENTIALS points to the correct JSON file path - Permission denied on resource — Verify the service account has the **Vertex AI User** role in GCP IAM - API not enabled — Return to the GCP Console and confirm the Vertex AI API is enabled for your project

Now verify the embedding model:

```

1  # Add to test_vertex.py
2
3  from vertexai.language_models import TextEmbeddingModel
4
5  embedding_model = TextEmbeddingModel.from_pretrained("text-embedding-004")
6
7  test_texts = [
8      "What test method was used for tensile strength testing?",
9      "batch quality certificate test specifications"
10 ]
11
12 embeddings = embedding_model.get_embeddings(test_texts)
13
14 for text, embedding in zip(test_texts, embeddings):
15     print(f"\nText: {text[:50]}...")
16     print(f"Embedding dimension: {len(embedding.values)}")
17     print(f"First 5 values: {embedding.values[:5]}")

```

The output should show embeddings with dimension 768 (the dimension of text-embedding-004):

Text: What test method was used for tensile strength te...

Embedding dimension: 768

First 5 values: [0.023, -0.041, 0.018, 0.067, -0.029]

Text: batch quality certificate test specifications...

Embedding dimension: 768

First 5 values: [0.031, -0.038, 0.022, 0.054, -0.011]

Note: The dimension 768 is important. When we create the HANA Cloud vector table in Chapter 2, we will define the column as `REAL_VECTOR(768)`. If you switch embedding models later, you must recreate the table — the dimension is fixed at table creation time.

2.10 Verifying the HANA Cloud Connection

Run a quick connectivity test to confirm HANA Cloud is reachable:

```

1  # test_hana.py
2  import os
3  from dotenv import load_dotenv
4  from hdbcli import dbapi
5
6  load_dotenv()
7
8  conn = dbapi.connect(
9      address=os.getenv("HANA_HOST"),
10     port=int(os.getenv("HANA_PORT")),
11     user=os.getenv("HANA_USER"),
12     password=os.getenv("HANA_PASSWORD"),
13     encrypt=True,
14     sslValidateCertificate=False
15 )
16
17 cursor = conn.cursor()
18 cursor.execute("SELECT VERSION FROM SYS.M_DATABASE")
19 version = cursor.fetchone()[0]
```



```
20 print(f"Connected to SAP HANA Cloud version: {version}")
21
22 # Verify vector engine is available
23 cursor.execute("""
24     SELECT COUNT(*) FROM SYS.M_FEATURE_USAGE
25     WHERE FEATURE_NAME = 'VECTOR'
26 """)
27 vector_count = cursor.fetchone()[0]
28 if vector_count > 0:
29     print("Vector engine: Available")
30 else:
31     print("Vector engine: Not detected (check HANA Cloud instance
32         settings)")
33 cursor.close()
34 conn.close()
35 print("\nHANA Cloud connection verified.")
```

```
1 python test_hana.py
```

Expected output:

Connected to SAP HANA Cloud version: 4.00.000.00.1698...

Vector engine: Available

HANA Cloud connection verified.

Expected terminal output:

HANA version: 4.00.000.00.1234567890 (fa/CE2024)

Vector engine: Available

SPARQL engine: Available

HANA connection: OK

2.11 What We Have Built

Let us take stock of what is now in place:

Component	Status	Where
BTP trial account	Running	account.hanatrial.ondemand.com
CF space “dev”	Ready	BTP Cockpit → Cloud Foundry
HANA Cloud instance	Running	BTP Cockpit → HANA Cloud Central
GCP project	Active	console.cloud.google.com
Vertex AI API	Enabled	GCP Console → APIs & Services
GCP service account	Created	GCP Console → IAM & Admin
BTP Destination “VertexAI”	Configured	BTP Cockpit → Connectivity
Python environment	Active	agents/.venv
Gemini API call	Verified	test_vertex.py passed
HANA connection	Verified	test_hana.py passed

This is not a small setup. You now have a working hybrid cloud environment: a managed in-memory database with vector and graph capabilities on BTP, connected to a frontier AI inference API on GCP, with credentials managed securely through BTP’s Destination Service.

The architecture you set up here mirrors how SAP customers deploy AI at scale — not a toy environment, but the real stack.

2.12 Understanding the Credential Flow

Before we move on, it is worth understanding exactly how credentials flow through the system. This will matter in Chapter 9 when you deploy to BTP, and it is a pattern that comes up repeatedly in enterprise AI systems.

Local development:

- .env file → GOOGLE_APPLICATION_CREDENTIALS
- google-auth library reads the JSON key directly
- Vertex AI API call authenticated

BTP Cloud Foundry deployment:

- BTP Destination Service “VertexAI”
- agents/srv/vertex_srv.py reads via CF environment

- `google-auth` constructs credentials from the stored JSON
- Vertex AI API call authenticated

The key insight is that **the application code never changes** between local and deployed. The credential source changes — `.env` locally, Destination Service in production — but the interface that `vertex_srv.py` exposes to the rest of the system is identical. This is the modularity principle in practice.

We will build `vertex_srv.py` in Chapter 2. For now, knowing the pattern exists is enough.

2.13 Summary

- SAP BTP trial provides free Cloud Foundry runtime and HANA Cloud — the only database that combines vector search and SPARQL Knowledge Graph in one service
- Google Vertex AI is an external HTTPS API — no GPU management, no model hosting, just an API call from BTP CF
- The BTP Destination Service stores GCP credentials securely — never in code, never in `manifest.yml`
- `text-embedding-004` produces 768-dimensional embeddings — this dimension is fixed and determines the HANA vector table schema
- Both the Gemini API call and the HANA Cloud connection are now verified end-to-end

In Chapter 3, we build the first half of our knowledge layer: vector search on SAP HANA Cloud. We will create the `MSDS_VECTORS` table, implement the embedding pipeline using `text-embedding-004`, and run our first semantic similarity search — the same operation our hybrid RAG agent will use at query time.

Chapter 3: Vector Search on SAP HANA Cloud

When a quality engineer asks “Did the supplier describe the tensile strength test method for this batch?”, they are asking a semantic question. The answer might be phrased differently in different sections of the certificate — the header calls it “test procedure”, the body describes it as “tensile analysis methodology”, and the appendix restates it as “verification approach”. Vector search finds it by meaning, not by exact keyword match — this is what makes it superior to SAP’s built-in DMS full-text search for narrative document content.

The same applies to every document type the Material Document Intelligence Platform handles. A batch certificate might record a test result in a summary table or in a methodology paragraph. An inspection report might document the same defect in the findings section or the inspector’s narrative remarks. Vector search surfaces the right passage regardless of where in the document it appears.

By the end of this chapter you will have a working vector search system on HANA Cloud. You will take a chunk of text from a material document, turn it into a 768-dimensional vector using Google’s `text-embedding-004` model, store it in a `REAL_VECTOR` column, and retrieve the most semantically similar chunks for a natural language question — all in roughly 200 lines of Python.

The closing section of this chapter is equally important: an honest account of what vector search cannot do. That limitation is the precise motivation for Chapter 4, where we add a Knowledge Graph to handle the queries that vectors get wrong.

3.1 What Are Embeddings? Intuition Without the Math

If you have spent your career writing ABAP, CDS views, and OData services, the word “embedding” probably sounds like something invented to make you feel old. It isn’t. The idea is older than most enterprise software, and the intuition is simple.

An embedding is a fixed-length list of numbers — in our case, 768 floating-point numbers — that represents the *meaning* of a piece of text. The crucial property is this: two pieces

of text with similar meanings produce embeddings that are close to each other in 768-dimensional space. Two pieces of text with different meanings produce embeddings that are far apart.

A useful analogy: imagine you are sorting books in a library. A traditional database would file them by title (alphabetical order) or by author. That is exact-match retrieval — fast, but useless if you don't know the title. A librarian, by contrast, knows that *Crime and Punishment* and *The Brothers Karamazov* belong on the same shelf because they are both Dostoevsky novels about guilt and redemption. The librarian has put each book at a *position* that reflects its content. Books on similar topics end up close together.

Embeddings do the same thing for text, but in 768 dimensions instead of two. The “position” of a sentence is just its embedding vector. Sentences about quality test results cluster together. Sentences about inspection findings cluster together. Sentences about regulatory codes cluster together — and importantly, they cluster *separately* from the test result sentences, even when they mention the same material.

Note: The number 768 is not magic. It is simply the output dimension chosen by the team that trained `text-embedding-004`. Other models produce 384, 1024, 1536, or 3072 dimensions. Higher is not always better — it just means more storage and slower comparisons. For our use case, 768 is a comfortable balance.

Two further points before we move on.

First, embeddings are not human-readable. If you print one out, you will see something like `[0.0234, -0.1107, 0.0891, ..., 0.0023]`. The individual numbers mean nothing in isolation. Their meaning emerges only in *relation* to other embeddings — specifically, in how close or far apart they are.

Second, the model that produces embeddings is the same model (or a sibling model) that was trained on huge corpora of text. It has internalized which words and phrases tend to occur in similar contexts, and it projects each input into a position that reflects that learned structure. You don't train the model. You don't fine-tune it. You just call it as an API.

Calling the API gives you a vector. Storing that vector and searching across millions of them efficiently is the database's job. Until recently, that meant using a specialized vector database — Pinecone, Weaviate, Milvus — alongside your real database. As of HANA Cloud QRC 1/2024, you no longer need a sidecar. HANA does it natively, and that

is the next topic.

3.2 The HANA Cloud `REAL_VECTOR` Column Type

HANA has always been a column-store database optimized for analytical workloads. In 2024, SAP added a new native column type built specifically for AI: `REAL_VECTOR`.

A `REAL_VECTOR` column stores a fixed-dimension array of single-precision floats. When you declare the column you fix the dimension (for example, `REAL_VECTOR(768)`), and HANA refuses any insert with a different dimension. This is what you want — mixing dimensions in the same column would mean comparing apples to oranges.

What makes this more than a glorified `BLOB` column is the set of operators HANA exposes on it:

- `TO_REAL_VECTOR(string)` — converts a string in the form `'[0.1, 0.2, ..., 0.768]'` into a real vector value. This is how you bind a parameter from Python.
- `COSINE_SIMILARITY(v1, v2)` — returns the cosine of the angle between two vectors. Range: `-1.0` (opposite) to `+1.0` (identical direction). For unit-normalized embeddings (which `text-embedding-004` produces), this is in practice between `0.0` and `1.0`.
- `L2DISTANCE(v1, v2)` — Euclidean distance between two vectors. Useful when your model is not normalized.

These operators are SIMD-accelerated on the HANA engine, which means a single CPU instruction processes multiple floats in parallel. Searching across hundreds of thousands of 768-dimensional vectors takes milliseconds, not seconds.

Tip: For embeddings produced by Google, OpenAI, and Cohere — all of which are L2-normalized at the API boundary — `COSINE_SIMILARITY` is the right choice. Use `L2DISTANCE` only if you are working with embeddings that are not normalized.

A second feature worth knowing about (we will not use it in this chapter, but it matters at scale) is the **HNSW index**. Without an index, every similarity query scans the entire table. With an HNSW index, HANA builds a graph structure that gets you approximate-nearest-neighbor results in logarithmic time. For our chapter you will not need it — five document chunks search in under 10 ms — but for production deployments with

millions of chunks, indexing is essential. We come back to this in Chapter 10 when we tune for production.




Figure: SAP HANA Cloud Central — the `REAL_VECTOR` column type is available in all HANA Cloud instances from version 4.0 onwards

HANA Cloud instances from version 4.0 onwards




Figure 2: Vector Embedding Pipeline

Figure 2: Vector Embedding Pipeline

Figure 3.1 — The vector pipeline has two phases. Ingestion (top): PDF text is chunked, each chunk is sent to `text-embedding-004`, and the resulting 768-dimensional vector is written to `MSDS_VECTORS`. Query time (bottom): the question is embedded with the same model, then HANA computes cosine distance between the query vector and every stored vector, returning the top-k most relevant chunks.

3.3 Creating the Vector Table: The Lazy Initialization Pattern

We need a table to hold our embeddings. The schema is simple:

```

1 CREATE TABLE MSDS_VECTORS (
2     ID BIGINT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
3     MATERIAL_NUMBER NVARCHAR(100) NOT NULL,
4     CHUNK_TEXT NCLOB NOT NULL,
5     CHUNK_INDEX INTEGER NOT NULL,
6     EMBEDDING REAL_VECTOR(768)
7 );

```

About the table name: Named MSDS_VECTORS in the reference implementation. In a broader deployment this would be MATERIAL_VECTORS. The SQL pattern is identical either way.

MATERIAL_NUMBER is the SAP material identifier — the same Material Number you would look up in SAP MM via transaction MM03. In the Material Document Intelligence Platform, the CAP service layer validates this value against real S/4HANA product master data via API_PRODUCT_SRV before accepting any document upload. This ensures that vectors are always anchored to a legitimate, active SAP material — not a free-text label that could drift out of sync with the product master. We use it as the tenant key so we can search within a single material’s chunks. CHUNK_TEXT is the raw text of the chunk; CHUNK_INDEX lets us reconstruct order if we need to. EMBEDDING is the 768-dimensional vector we computed from CHUNK_TEXT.

You could just run this CREATE TABLE statement once with the SAP HANA Database Explorer and be done with it. But there is a subtle reason not to.

The dimension 768 is hardcoded into the schema. If at any point you decide to switch from text-embedding-004 (768 dim) to a model with a different output dimension, your insert will fail and you will need to drop and recreate the table. Worse, that hardcoded number lives in your DDL while the actual dimension lives in the model — two sources of truth that can drift apart.

The cleaner pattern is **lazy initialization**: don’t create the table until the first time you embed something, then read the dimension from the embedding response and use it in the CREATE TABLE. This way your code has exactly one source of truth — the model’s

actual output dimension at runtime.

Here is the relevant snippet from `agents/srv/vector_srv.py` (we will see the full file in section 3.8):

```

1  TABLE_NAME = "MSDS_VECTORS"
2  _table_created = False
3
4  def _ensure_table(embedding_dim: int):
5      global _table_created
6      if _table_created:
7          return
8      conn = get_connection()
9      cursor = conn.cursor()
10     cursor.execute(f"""
11         CREATE TABLE IF NOT EXISTS {TABLE_NAME} (
12             ID BIGINT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
13             MATERIAL_NUMBER NVARCHAR(100) NOT NULL,
14             CHUNK_TEXT NCLOB NOT NULL,
15             CHUNK_INDEX INTEGER NOT NULL,
16             EMBEDDING REAL_VECTOR({embedding_dim})
17         )
18     """)
19     conn.commit()
20     cursor.close()
21     _table_created = True

```

Three things to note:

1. The `_table_created` module-level flag is a one-time guard. We don't want to issue `CREATE TABLE IF NOT EXISTS` on every insert — that is a roundtrip we don't need.
2. The `IF NOT EXISTS` clause means restarting the Python process is safe. The first call after restart will hit the database once, see the table exists, and move on.
3. The `embedding_dim` is passed in by the caller, who got it from the actual embedding. There is no hardcoded 768 anywhere in the application code.

Warning: Do not be tempted to add an `ALTER TABLE` path that resizes the `REAL_VECTOR` column if the dimension changes. A dimension change means your old embeddings are incompatible with new ones, and you should treat it as a full re-indexing event — drop the table, re-embed all your documents,

and reload. Silently mixing dimensions is a recipe for confusion.

3.4 Calling text-embedding-004 from Python

We talked about embeddings as a concept; now we need to actually produce one. The `langchain_google_genai` library gives us a clean interface to Google's embedding models without the operational complexity of a dedicated Vertex AI SDK setup.

Looking at `agents/srv/vertex_srv.py`:

```
1  import os
2  import threading
3  import vertexai
4  from vertexai.generative_models import GenerativeModel
5  from vertexai.language_models import TextEmbeddingModel
6  from dotenv import load_dotenv
7
8  load_dotenv()
9
10 _init_lock = threading.Lock()
11 _initialized = False
12
13 def _ensure_initialized():
14     global _initialized
15     if not _initialized:
16         with _init_lock:
17             if not _initialized:
18                 vertexai.init(
19                     project=os.getenv("GCP_PROJECT_ID"),
20                     location=os.getenv("GCP_LOCATION", "us-central1")
21                 )
22                 _initialized = True
23
24 def get_llm():
25     _ensure_initialized()
26     return GenerativeModel("gemini-2.5-flash-preview-05-20")
27
```

```

28 def get_embedding_model():
29     _ensure_initialized()
30     return TextEmbeddingModel.from_pretrained("text-embedding-004")
31
32 def embed_text(text: str) -> list[float]:
33     model = get_embedding_model()
34     embeddings = model.get_embeddings([text])
35     return embeddings[0].values # list of 768 floats

```

The pattern here mirrors the lazy table creation pattern. `vertexai.init()` is global SDK state; calling it more than once is wasteful. The double-checked locking idiom (`if not _initialized` outside the lock, then again inside) is the standard way to do this safely under FastAPI's threaded request model. If you have never seen the pattern before, the gist is: the outer check is the fast path for the common case (already initialized, no lock needed), and the inner check makes sure that two threads racing for the lock don't both initialize.

`get_embedding_model()` returns a `TextEmbeddingModel` instance bound to `text-embedding-004`. This is cheap — the SDK caches the model client internally — but we still gate it behind `_ensure_initialized` so that calling code does not have to think about ordering.

`embed_text()` is the function the rest of our application uses. It takes a string and returns a Python list of 768 floats. Internally, the SDK takes a list of strings, so we pass `[text]` and unwrap the first result with `embeddings[0].values`.

Note: Vertex AI's embedding endpoint accepts batches of up to 250 texts per call. For a one-off document upload that ends up with maybe 30 chunks, the savings of batching are negligible compared to the simplicity of one-text-per-call. We will revisit batching in Chapter 10 when we look at bulk-loading historical documents.

You should also be aware that Vertex AI charges per 1,000 input characters for embeddings (about \$0.000025 per 1,000 chars at the time of writing). A typical 12-page MSDS PDF has roughly 30,000 characters and chunks into ~30 pieces. The cost to embed it is approximately \$0.0008 — less than a tenth of a cent. Search-time embeddings (one per question) are equally cheap.

3.5 Storing Embeddings: INSERT with REAL_VECTOR

Once we have an embedding (a Python `list[float]`), we need to get it into the `REAL_VECTOR` column. The HANA Python driver `hdbcli` does not yet have a native binding for vector parameters, so we use the `TO_REAL_VECTOR(string)` SQL function with a string parameter.

The format HANA expects is a JSON-style array literal: `[0.0234, -0.1107, 0.0891, ...]`. We build this string in Python with a simple comma-join:

```
1 embedding_str = "[" + ",".join(str(x) for x in embedding) + "]"
```

A typical 768-dimensional embedding string is around 12,000 characters. That is fine — HANA accepts it as an `NVARCHAR` parameter without issue.

The full insert function:

```
1 def store_embedding(material_number: str, chunk_text: str, chunk_index:
  int, embedding: list[float]):
2     _ensure_table(len(embedding))
3     conn = get_connection()
4     cursor = conn.cursor()
5     embedding_str = "[" + ",".join(str(x) for x in embedding) + "]"
6     cursor.execute(f"""
7         INSERT INTO {TABLE_NAME} (MATERIAL_NUMBER, CHUNK_TEXT,
8         CHUNK_INDEX, EMBEDDING)
9         VALUES (?, ?, ?, TO_REAL_VECTOR(?))
10    """, (material_number, chunk_text, chunk_index, embedding_str))
11    conn.commit()
12    cursor.close()
```

A few engineering notes:

- We pass `material_number`, `chunk_text`, `chunk_index`, and the embedding string as positional parameters. This is parameterized SQL — no concatenation of user data into the query, so SQL injection is impossible.
- `_ensure_table(len(embedding))` runs on every call but exits early after the first invocation thanks to the `_table_created` flag.
- We commit per insert. For uploading a single document this is fine. For bulk uploads, batch your inserts and commit once at the end — Chapter 10 covers this.

Tip: If you want to verify what got stored, the SAP HANA Database Explorer is your friend. Open the `MSDS_VECTORS` table, click “Open Data”, and you will see your rows. The `EMBEDDING` column displays as a truncated array preview, but you can click into a cell to see the full 768 values.

Note: You can inspect the `MSDS_VECTORS` table in SAP HANA Database Explorer (accessible from HANA Cloud Central → Open in → SAP HANA Database Explorer). The `EMBEDDING` column will show truncated vector data like `[0.0234, -0.1823, 0.0912, ...]`.

3.6 Cosine Similarity Search: The Query Explained

This is the heart of the chapter. The query that retrieves the top-K chunks for a question is:

```

1 SELECT TOP 5
2     CHUNK_TEXT,
3     CHUNK_INDEX,
4     COSINE_SIMILARITY(EMBEDDING, TO_REAL_VECTOR(?)) AS SCORE
5 FROM MSDS_VECTORS
6 WHERE MATERIAL_NUMBER = ?
7 ORDER BY SCORE DESC;
```

Five clauses, each doing real work. Let’s read them in execution order.

`WHERE MATERIAL_NUMBER = ?` filters the table to chunks belonging to a specific material. Without this, our search would compete across every chunk of every document in the table — fine when you have one document, fatal when you have ten thousand. By filtering early, HANA scans only the rows that could possibly match. This is the one place in the query where the optimizer can use a B-tree index, so make sure you have one on `MATERIAL_NUMBER` once you go to production. (We add it in Chapter 10.)

`COSINE_SIMILARITY(EMBEDDING, TO_REAL_VECTOR(?))` is the per-row computation. For each row passing the `WHERE` filter, HANA takes the stored `EMBEDDING` vector, the question’s vector (parsed from the bound parameter), and computes the cosine of the angle between them. The result is a `DOUBLE` between roughly 0.0 and 1.0 — higher means more similar. This is where the SIMD acceleration earns its keep.

AS SCORE names the result so we can refer to it in the ORDER BY clause. It's a small thing but worth noting because it is what allows us to write ORDER BY SCORE DESC instead of having to repeat the COSINE_SIMILARITY(...) expression.

ORDER BY SCORE DESC sorts the surviving rows from highest similarity to lowest. The bulk of the cost here is the sort itself; for our small data this is instant.

SELECT TOP 5 returns only the first five rows after the sort. In LangGraph's vector chain we will pass these chunks to the LLM as context for answering the question. Five is a reasonable default — large enough to give the model useful coverage, small enough not to blow the context window.

The order of operations matters. HANA's optimizer is smart enough to push the WHERE filter down to the table scan (so we never compute similarity for chunks belonging to other materials), but you should still write your queries with the filter explicit. Don't compute similarity across the whole table and filter the result.

Important: A common mistake is to forget the WHERE clause when prototyping with one material, then keep that habit when you add a second. Every query in your application should filter by MATERIAL_NUMBER. There is no business reason a question about one material should ever surface a chunk from a different material's document, even if some words happen to overlap.

The Python wrapper:

```
1 def search_similar(question: str, material_number: str, top_k: int = 5) ->
  list[dict]:
2     question_embedding = embed_text(question)
3     embedding_str = "[" + ",".join(str(x) for x in question_embedding) +
  "]"
4     conn = get_connection()
5     cursor = conn.cursor()
6     cursor.execute(f"""
7         SELECT TOP {top_k}
8             CHUNK_TEXT,
9             CHUNK_INDEX,
10            COSINE_SIMILARITY(EMBEDDING, TO_REAL_VECTOR(?)) AS SCORE
11        FROM {TABLE_NAME}
12        WHERE MATERIAL_NUMBER = ?
13        ORDER BY SCORE DESC
```

```

14     """", (embedding_str, material_number))
15     rows = cursor.fetchall()
16     cursor.close()
17     return [{"chunk": row[0], "chunk_index": row[1], "score":
        float(row[2])} for row in rows]

```

We embed the question (one Vertex AI call), build the parameter string the same way as for inserts, and run the query. The result is a list of dicts — the chunk text, its position in the original document, and its similarity score.

Note: `top_k` is interpolated into the SQL string, not bound as a parameter. This is safe because we control the value (it is a function argument with a default of 5, never user input). HANA does not allow `TOP ?` with bound parameters, so we have no choice. If you ever expose `top_k` to end users, validate it as an integer before interpolation.

3.7 Chunk Size Analysis and Boundary Strategy for SAP Material Documents

Chunking is the most consequential design decision in any RAG system — and it is rarely given enough attention. The wrong chunk size silently degrades retrieval quality without producing any obvious error. This section explains the reasoning behind every number we use, grounded in the specific document types in this platform.

Why chunking exists at all

A batch quality certificate is typically 2–4 pages. An equipment maintenance history can run to 20 pages. A full MSDS is 16 structured sections. Embedding an entire document as one vector produces one point in 768-dimensional space — a single embedding that represents everything and therefore retrieves nothing precisely. When the user asks “what was the tensile strength result for batch BATCH-2024-0871?”, a whole-document embedding returns the entire certificate whether or not the test result is the dominant topic. Chunking breaks the document into pieces small enough that each embedding represents one focused idea — and retrieval returns the right piece, not the whole document.

The chunk size trade-off

Every chunking decision sits on a spectrum between two failure modes:

Too small (under 100 tokens — a sentence or two): Retrieval becomes precise but brittle. A single sentence like “*Result: PASS*” scores perfectly for “did this batch pass?” but provides no context about which test, which specification, or which method. The LLM synthesising an answer from five isolated sentences produces a shallow, disconnected response. Short chunks also multiply storage and search cost — a 20-page maintenance history at 30-token chunks produces 600+ vectors where 60 would suffice.

Too large (over 600 tokens — a full page): Retrieval becomes noisy. A page of a batch certificate that covers tensile strength, yield strength, and elongation simultaneously scores moderately for all three questions but tops the ranking for none of them. The chunk also contains irrelevant content that consumes LLM context window without contributing to the answer.

The sweet spot for SAP material documents is 400–500 tokens. This range consistently captures one complete logical unit — one test result block, one invoice line item group, one maintenance procedure step, one delivery note block — without spilling into adjacent topics.

Why 50-token overlap

Overlap exists to handle boundary sentences. Consider a batch certificate where the methodology description ends at token 495 of a chunk and the test result begins at token 1 of the next chunk. Without overlap, a question about “how was the tensile test performed and what did it find?” retrieves either the methodology chunk or the result chunk — but not both, because neither chunk alone contains the complete answer. With 50-token overlap, the result chunk begins with the last two sentences of the methodology, giving the embedding model enough context to score it correctly for the combined question.

Fifty tokens (roughly 2–3 sentences) is the empirically right overlap for enterprise documents with dense, structured content. It is small enough that it does not double-count content in retrieval, and large enough to bridge the boundary for the most common cross-boundary questions.

Natural boundaries by document type

The 500-token maximum is a ceiling, not a target. Each document type in this platform has natural semantic boundaries that should be the primary split point — the 500-token limit is only the fallback when a natural section runs long.

Document Type	Primary split		Notes
	boundary	Typical chunk size	
Batch Quality Certificate	Per test parameter block	150–300 tokens	Each test (tensile, yield, elongation) becomes one chunk — supplier, result, specification, pass/fail together
GR Inspection Report	Per line item inspected	200–350 tokens	Inspector observations and QM decision stay in the same chunk as the line item
Equipment Maintenance History	Per maintenance order / service event	300–500 tokens	Failure description and root cause must stay together — splitting them breaks the narrative
Supplier Invoice	Header + per line item block	100–250 tokens	Header (vendor, PO, payment terms) is one chunk; each line item is a separate chunk

Document Type	Primary split		Notes
	boundary	Typical chunk size	
MSDS / SDS	Per regulatory section (1–16)	150–400 tokens	Section headers defined by GHS/OSHA regulation — present in every compliant document

Our chunking parameters

```

1  CHUNK_SIZE      = 500    # tokens — maximum per chunk
2  CHUNK_OVERLAP  = 50     # tokens — shared context at boundaries
3  CHUNK_MIN       = 100    # tokens — discard anything shorter (headers, page
    numbers)

```

These are implemented in `doc_srv.py` in Chapter 5 using LangChain’s `RecursiveCharacterTextSplitter` with a `tiktoken` encoder. The splitter respects paragraph boundaries first, then sentence boundaries, falling back to character boundaries only when a paragraph exceeds the maximum. This produces chunks that almost always end at a natural linguistic break — not mid-sentence.

Tip: Tokens, not characters. `text-embedding-004` accepts up to 2,048 tokens per call — our 500-token maximum is well within the limit. One token is approximately 4 characters of English text or 0.75 of a word.

Note: For the tests in this chapter we hand-write five chunks representing a batch quality certificate. The full document chunker with these parameters is implemented in Chapter 5. The goal here is to validate the round-trip: embed → store → embed-question → search → score. Chunking quality is a Chapter 5 concern.

3.8 Building `vector_srv.py` — The Vector Service Layer

Now we put the pieces together. Create the file `agents/srv/vector_srv.py` with the following content:

```

1  import json
2  from .hdb_srv import get_connection
3  from .vertex_srv import embed_text
4
5  TABLE_NAME = "MSDS_VECTORS"
6  _table_created = False
7
8  def _ensure_table(embedding_dim: int):
9      global _table_created
10     if _table_created:
11         return
12     conn = get_connection()
13     cursor = conn.cursor()
14     cursor.execute(f"""
15         CREATE TABLE IF NOT EXISTS {TABLE_NAME} (
16             ID BIGINT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
17             MATERIAL_NUMBER NVARCHAR(100) NOT NULL,
18             CHUNK_TEXT NCLOB NOT NULL,
19             CHUNK_INDEX INTEGER NOT NULL,
20             EMBEDDING REAL_VECTOR({embedding_dim})
21         )
22     """)
23     conn.commit()
24     cursor.close()
25     _table_created = True
26
27 def store_embedding(material_number: str, chunk_text: str, chunk_index:
int, embedding: list[float]):
28     _ensure_table(len(embedding))
29     conn = get_connection()
30     cursor = conn.cursor()
31     embedding_str = "[" + ",".join(str(x) for x in embedding) + "]"
32     cursor.execute(f"""

```

```

33         INSERT INTO {TABLE_NAME} (MATERIAL_NUMBER, CHUNK_TEXT,
CHUNK_INDEX, EMBEDDING)
34         VALUES (?, ?, ?, TO_REAL_VECTOR(?))
35         """ , (material_number, chunk_text, chunk_index, embedding_str))
36         conn.commit()
37         cursor.close()
38
39 def search_similar(question: str, material_number: str, top_k: int = 5) ->
list[dict]:
40     question_embedding = embed_text(question)
41     embedding_str = "[" + ",".join(str(x) for x in question_embedding) +
"] "
42     conn = get_connection()
43     cursor = conn.cursor()
44     cursor.execute(f"""
45         SELECT TOP {top_k}
46             CHUNK_TEXT,
47             CHUNK_INDEX,
48             COSINE_SIMILARITY(EMBEDDING, TO_REAL_VECTOR(?)) AS SCORE
49         FROM {TABLE_NAME}
50         WHERE MATERIAL_NUMBER = ?
51         ORDER BY SCORE DESC
52     """, (embedding_str, material_number))
53     rows = cursor.fetchall()
54     cursor.close()
55     return [{"chunk": row[0], "chunk_index": row[1], "score":
float(row[2])} for row in rows]
56
57 def delete_vectors(material_number: str) -> int:
58     conn = get_connection()
59     cursor = conn.cursor()
60     cursor.execute(f"DELETE FROM {TABLE_NAME} WHERE MATERIAL_NUMBER = ?",
(material_number,))
61     deleted = cursor.rowcount
62     conn.commit()
63     cursor.close()
64     return deleted
65
66 def count_vectors(material_number: str) -> int:

```

```

67     conn = get_connection()
68     cursor = conn.cursor()
69     cursor.execute(f"SELECT COUNT(*) FROM {TABLE_NAME} WHERE
MATERIAL_NUMBER = ?", (material_number,))
70     row = cursor.fetchone()
71     cursor.close()
72     return row[0] if row else 0

```

The module exposes four functions to the rest of the application:

- `store_embedding(material_number, chunk_text, chunk_index, embedding)` — stores a single chunk and its vector. Called once per chunk during document upload.
- `search_similar(question, material_number, top_k)` — embeds the question and returns the top-K most similar chunks for a given material. Called once per user question.
- `delete_vectors(material_number)` — wipes a material's vectors. Called when a document is replaced or deleted. Returns the row count for logging.
- `count_vectors(material_number)` — reports how many chunks exist for a material. Useful for diagnostics and tests.

There is no `update_embedding` — for chunked documents, the right pattern is delete-then-reinsert. If you change the chunking algorithm, the chunk indices change, so partial updates are dangerous. Treat each upload as a fresh load.

Note: The connection comes from `hdb_srv.py`, which uses thread-local storage to give each request thread its own HANA connection. `hdbcli` connections are not safe to share across threads, but they are cheap to create. Thread-locals give us per-thread reuse without any cross-thread contention. We covered this pattern briefly in Chapter 1 — the file is reproduced for completeness:

```

1  import os
2  import threading
3  from hdbcli import dbapi
4  from dotenv import load_dotenv
5
6  load_dotenv()

```

```

7
8 _local = threading.local()
9
10 def get_connection():
11     if not hasattr(_local, 'conn') or _local.conn is None:
12         _local.conn = dbapi.connect(
13             address=os.getenv("HANA_HOST"),
14             port=int(os.getenv("HANA_PORT", 443)),
15             user=os.getenv("HANA_USER"),
16             password=os.getenv("HANA_PASSWORD"),
17             encrypt=True,
18             sslValidateCertificate=False
19         )
20     return _local.conn

```

Warning: `sslValidateCertificate=False` is fine for development against the BTP-issued HANA endpoint. For production you should provide the correct certificate via `sslTrustStore` and set `sslValidateCertificate=True`. We will tighten this in Chapter 9.

3.9 Testing: Business Questions Across Document Types

Time to see it work end-to-end. Create the file `agents/test_vector.py`:

```

1  import os
2  from dotenv import load_dotenv
3  from srv.vector_srv import store_embedding, search_similar, count_vectors
4  from srv.vertex_srv import embed_text
5
6  load_dotenv()
7
8  TEST_MATERIAL = "BATCH-QC-MAT-001"
9  TEST_CHUNKS = [
10     "Batch lot number: BATCH-2024-0871. Material: Steel rod grade S355.
       Supplier: ACME Steel AG. Certificate number: QC-CERT-44781.
       Certification date: 2024-03-15.",

```

```

11     "Tensile strength test: measured value 487 MPa, specification minimum
12       355 MPa. Result: PASS. Test method: ISO 6892-1 tensile testing at
        ambient temperature.",
13     "Surface inspection: visual examination performed per ISO 8501-1. No
        pitting, lamination, or surface cracks observed. Inspector: Klaus
        Weber. Result: PASS.",
14     "Delivery details: 250 units delivered against PO 4500123456. Quantity
        accepted: 250. No shortages or damages noted on receipt. QM usage
        decision: Unrestricted use.",
15     "Storage recommendation: store in dry covered warehouse, avoid direct
        contact with ground. Maximum stack height 3 pallets. Protect from
        moisture and corrosive atmospheres."
16 ]
17 print("Embedding and storing test chunks...")
18 for i, chunk in enumerate(TEST_CHUNKS):
19     embedding = embed_text(chunk)
20     store_embedding(TEST_MATERIAL, chunk, i, embedding)
21     print(f"  Stored chunk {i}: dim={len(embedding)},
        preview='{chunk[:50]}...'")
22
23 stored = count_vectors(TEST_MATERIAL)
24 print(f"\nStored {stored} vectors for material {TEST_MATERIAL}")
25
26 print("\nSearching: 'What test method was used for tensile strength
        verification?'")
27 results = search_similar("What test method was used for tensile strength
        verification?", TEST_MATERIAL)
28 for r in results:
29     print(f"  Score: {r['score']:.4f} | {r['chunk'][:80]}...")
30
31 print("\nSearching: 'What are the storage and handling requirements for
        this material?'")
32 results = search_similar("What are the storage and handling requirements
        for this material?", TEST_MATERIAL)
33 for r in results:
34     print(f"  Score: {r['score']:.4f} | {r['chunk'][:80]}...")

```

Run it from the agents/ directory:

```

1 cd agents
2 python test_vector.py

```

Expected output (your numbers will vary slightly):

Embedding and storing test chunks...

Stored chunk 0: dim=768, preview='Batch lot number: BATCH-2024-0871. Material: Steel...'

Stored chunk 1: dim=768, preview='Tensile strength test: measured value 487 MPa, spe...'

Stored chunk 2: dim=768, preview='Surface inspection: visual examination performed pe...'

Stored chunk 3: dim=768, preview='Delivery details: 250 units delivered against PO 45...'

Stored chunk 4: dim=768, preview='Storage recommendation: store in dry covered wareho...'

Stored 5 vectors for material BATCH-QC-MAT-001

Searching: 'What test method was used for tensile strength verification?'

Score: 0.8124 | Tensile strength test: measured value 487 MPa, specification minimum 355

Score: 0.7456 | Surface inspection: visual examination performed per ISO 8501-

1. No pitting...

Score: 0.6203 | Batch lot number: BATCH-2024-0871. Material: Steel rod grade S355. Suppl

Score: 0.5891 | Storage recommendation: store in dry covered warehouse, avoid direct cont

Score: 0.5102 | Delivery details: 250 units delivered against PO 4500123456. Quantity acc

Searching: 'What are the storage and handling requirements for this material?'

Score: 0.8442 | Storage recommendation: store in dry covered warehouse, avoid direct cont

Score: 0.6334 | Delivery details: 250 units delivered against PO 4500123456. Quantity acc

Score: 0.5912 | Batch lot number: BATCH-2024-0871. Material: Steel rod grade S355. Suppl

Score: 0.5108 | Tensile strength test: measured value 487 MPa, specification minimum 355

Score: 0.4789 | Surface inspection: visual examination performed per ISO 8501-

1. No pitting...

This is exactly what you want to see. For the tensile test method question, the tensile strength chunk wins at 0.81, with the surface inspection chunk (which also describes a test methodology) in second place at 0.74. The delivery details chunk, which has nothing to do with the question, ranks last at 0.51. For the storage question, the storage recommendation chunk wins decisively at 0.84.

Note the absolute scores. The top match for a well-targeted question lands around 0.80–0.85. This is the band where you should expect “good” semantic matches with `text-embedding-004`. Anything above 0.90 usually means near-paraphrase. Anything below 0.50 usually means the model is reaching.

The same vector infrastructure serves every document type without modification. A tensile test method question for a batch certificate returns the test result chunk. A storage requirements question for an MSDS returns the Section 7 chunk. An approved quantity question for an invoice returns the line-item summary chunk. The question changes. The search code does not. This is the commercial value of building on HANA’s `REAL_VECTOR` column rather than on a document-type-specific text search engine.

Tip: The scores are reproducible only if you re-embed identical text against the same model version. Vertex AI may update the model under the same name; if your scores drift over time, that is the most likely cause. Pin to a specific version if reproducibility matters.

If you see authentication errors at this point, double-check that `GOOGLE_APPLICATION_CREDENTIALS` points to your `gcp-sa-key.json` file and that the service account has the Vertex AI User role. The HANA side is more forgiving — connection errors usually surface as “host not reachable”, which means a typo in `HANA_HOST` or a network rule blocking outbound 443.

3.10 What Vector Search Gets Right — and What It Gets Wrong

Time for honesty. Vector search is genuinely powerful, but it has a sharp limitation that becomes obvious as soon as you look at the third example.

Run the test one more time, with a precise factual question:

```
1 print("\nSearching: 'What is the certificate number for batch\n  BATCH-2024-0871?'\n")\n2 results = search_similar("What is the certificate number for batch\n  BATCH-2024-0871?", TEST_MATERIAL)\n3 for r in results:
```

```
4 print(f" Score: {r['score']:.4f} | {r['chunk'][:80]}...")
```

You will see something like:

Searching: 'What is the certificate number for batch BATCH-2024-0871?'

Score: 0.6534 | Batch lot number: BATCH-2024-0871. Material: Steel rod grade S355. Suppl

Score: 0.5821 | Tensile strength test: measured value 487 MPa, specification minimum 355

Score: 0.4912 | Surface inspection: visual examination performed per ISO 8501-

1...

...

The certificate chunk does come back first — at 0.65. That is noticeably weaker than the 0.81 we saw for the methodology question, despite this being a question about a chunk that literally contains “QC-CERT-44781”.

Why so weak? Because the question asks for a specific *identifier string* — “QC-CERT-44781” — and our chunk does not contain that exact phrasing in the query. The embedding model knows certificates and batch numbers are related concepts, but it encodes them as *near-each-other-in-768-d-space*, which is fuzzier than equality. The right answer to “What is the certificate number for batch BATCH-2024-0871?” is “QC-CERT-44781” — not a paragraph that mentions certificates somewhere.

For our test data this still works because there are only five chunks. Now imagine the same query against a real corpus with 500 batch certificates and 15,000 chunks. Many chunks will mention “certificate” in different contexts — quality reports, delivery notes, compliance filings. The signal-to-noise ratio collapses. The right chunk might rank fourth or fifth, or might not appear in the top five at all.

This is **the exact identifier problem**, and it is not a bug in vector search — it is the nature of vector search. Embeddings are wonderful for fuzzy semantic matching (“tensile test methodology” ↔ “ISO 6892-1 tensile testing at ambient temperature”). They are bad at exact symbolic recall (“QC-CERT-44781” ↔ the row whose certificate number literal equals “QC-CERT-44781”).

Three categories of question hit this limitation hard:

1. **Identifier lookups.** Certificate numbers (QC-CERT-44781), batch lot numbers (BATCH-2024-0871), purchase order numbers (4500123456), equipment numbers. Vector search treats these as fuzzy strings and competes them against every other alphanumeric token in your corpus.

2. **Aggregations and counts.** “How many of our batches from supplier ACME Steel AG passed tensile testing this quarter?” is impossible for vector search — there is no “count” embedding direction. You need structured query.
3. **Negation and exclusion.** “Show me inspection reports where the inspector did *not* record a usage decision of ‘Restricted.’” Embeddings have no robust way to encode “not” — the embedding for “usage decision: Restricted” and “usage decision: not Restricted” is closer than you would hope.

SAP customers running on HANA have always thought in structured terms. Materials, batch lot numbers, certificate identifiers, test result values — these are not free text. They are master data. They live in tables with foreign keys. They are the kind of thing a Knowledge Graph models naturally.

That is the bridge into Chapter 4. We will take the same material document, extract its structured facts into RDF triples, store them in HANA’s graph engine, and query them with SPARQL. When the quality engineer asks “What is the certificate number for this batch?”, we will return “QC-CERT-44781” with full confidence — not by guessing from a 0.65 cosine similarity, but by traversing a graph relationship that was set when the document was ingested.

The full Hybrid RAG agent in Chapter 7 runs both retrievers in parallel and lets a routing classifier decide which answer to trust. Vector search owns the fuzzy questions. The Knowledge Graph owns the precise ones. Neither is sufficient alone; together they cover the question space enterprise document users actually ask.

3.11 Summary

In this chapter you built a working vector search system on HANA Cloud. The major points to take away:

- **Embeddings are positions in 768-dimensional space.** Two pieces of text with similar meanings end up close together. The model learns this from large corpora and you call it as an API.
- **REAL_VECTOR is HANA’s native vector column type.** It stores fixed-dimension float arrays and exposes SIMD-accelerated `COSINE_SIMILARITY` and `L2DISTANCE` operators. You do not need a separate vector database.

- **Lazy table initialization keeps your schema in sync with your model.** Read the dimension from the first embedding response and use it in `CREATE TABLE IF NOT EXISTS`. Never hardcode the dimension in DDL.
- **The cosine similarity query has five clauses, each load-bearing.** `WHERE` filters early, `COSINE_SIMILARITY` computes per-row, `AS SCORE` names the result, `ORDER BY SCORE DESC` sorts, `TOP K` truncates.
- **Chunking adapts to document structure; the storage code does not.** Regulatory documents (MSDS/SDS) use GHS section boundaries. Invoices use line-item blocks. Batch certificates use test result sections. The same `MSDS_VECTORS` table and `search_similar` function serve all document types.
- **The `MATERIAL_NUMBER` column is the anchor to SAP MM.** The CAP layer validates it against S/4HANA product master via `API_PRODUCT_SRV` before accepting uploads. Every vector in the table is traceable to a real, validated SAP material.
- **Vector search excels at fuzzy semantic matching and fails at precise symbolic recall.** The exact identifier problem — certificate numbers, batch lot numbers, PO numbers — motivates the Knowledge Graph in Chapter 4.

The codebase now has working `hdb_srv.py`, `vertex_srv.py`, and `vector_srv.py` modules, plus a passing `test_vector.py` script. The vector retrieval half of the Hybrid RAG agent is live.

3.12 Checkpoint

Before moving on to Chapter 4, verify the following:

- ☐ `agents/srv/vector_srv.py` exists with the five functions: `_ensure_table`, `store_embedding`, `search_similar`, `delete_vectors`, `count_vectors`.
- ☐ Running `python test_vector.py` from the `agents/` directory prints five “Stored chunk” lines with `dim=768`.
- ☐ The test prints two search result blocks. The first ranks the tensile strength chunk at the top (~0.81). The second ranks the storage recommendation chunk at the top (~0.84).
- ☐ Opening the `MSDS_VECTORS` table in the SAP HANA Database Explorer shows five rows for `MATERIAL_NUMBER = 'BATCH-QC-MAT-001'`.
- ☐ The third (optional) test — querying for “What is the certificate number for batch

BATCH-2024-0871?” — returns the batch header chunk at the top, but with a noticeably lower score (~0.65). You understand why.

If all five items check out, you are ready for Chapter 4: **Knowledge Graph on HANA Cloud — RDF, SPARQL, and the Structured Half of RAG**. We will set up HANA’s graph engine, load RDF triples extracted from the same material document, and query them with SPARQL — solving the exact identifier problem in the most direct way possible.

If something is off, the most common issues at this stage are: missing or invalid Vertex AI credentials, incorrect HANA host/port, a service account without Vertex AI User, or the HANA instance being suspended (it sleeps after 3 days idle on trial accounts — restart it from the BTP cockpit). Fix and re-run `test_vector.py` until you see the expected output. The next chapter assumes a working `vector_srv.py`.

Chapter 4: Knowledge Graphs on SAP HANA Cloud

Vector search excels at finding relevant passages. It fails at extracting precise facts. When a procurement manager asks “Which supplier certified batch BATCH-2024-0871?”, they need a name — “ACME Steel AG” — not a paragraph that mentions suppliers. When a quality engineer asks “Did this batch pass the tensile strength test?”, they need “PASS” and the measured value, not a fuzzy match against test methodology paragraphs.

RDF Knowledge Graphs store facts as precise triples that can be queried exactly. This is what HANA Cloud’s SPARQL engine provides. Unlike the vector store, which searches by proximity in embedding space, the Knowledge Graph traverses explicit relationships. The query is not “find what is most similar to this question” — it is “traverse this graph and return the value at this exact node.”

The right answer to “What is the certificate number for batch BATCH-2024-0871?” is “QC-CERT-44781” — not a paragraph, not a summary, the exact identifier. The right answer to “Which lab certified this batch?” is “Bureau Veritas Testing GmbH” — a name, not a chunk of text.

By the end of this chapter, a SPARQL query for a batch certificate number will return exactly ["QC-CERT-44781"] — not a paragraph that mentions it, the identifier itself.

If you have never written a line of RDF or SPARQL, do not worry. There are entire books about each. We need only enough of each to build a working hybrid RAG system. Concretely, you will need to understand three SPARQL patterns and one stored procedure call. That is the whole surface area for the rest of this book.

4.1 What is a Knowledge Graph?

A **Knowledge Graph** is a collection of facts, where each fact is expressed as a relationship between two things. That is the entire idea. The vocabulary that the RDF community uses around it can sound ceremonial — “subject-predicate-object triples organized as a directed labeled multigraph” — but strip the ceremony away and you are left with three-word sentences.

Consider the fact:

Batch BATCH-2024-0871 was certified by supplier ACME Steel AG.

Three pieces:

Part	Value
Subject	BATCH-2024-0871
Predicate	certified by
Object	ACME Steel AG

This three-part fact is called a **triple**. A Knowledge Graph is a pile of triples. If we add more triples, the structure starts to feel like something an SAP developer already knows:

BATCH-2024-0871	certifiedBy	ACME Steel AG
BATCH-2024-0871	certificateNumber	QC-CERT-44781
BATCH-2024-0871	testResult	PASS (tensile)
BATCH-2024-0871	testValue	487 MPa
BATCH-2024-0871	certifyingLab	Bureau Veritas Testing GmbH
BATCH-2024-0871	certificationDate	2024-03-15
QC-CERT-44781	description	"Certificate for batch BATCH-2024-0871, steel grade S355"

Read those rows. You have just read a Knowledge Graph. Each row is a fact about this batch. The same subject (BATCH-2024-0871) can appear in many rows, just like a batch number can appear many times in QM-related tables. The same predicate (certifiedBy) can be reused for any other batch.

Note: If you are mentally translating this to a relational table with three columns — subject, predicate, object — you are not wrong. That is exactly how RDF can be persisted internally. The difference from a normal table is *what kinds of questions you can ask of it*, which we will get to with SPARQL.

Why call it a graph?

Because if you draw it, it looks like one. Subjects and objects are nodes; predicates are edges. The same certifying lab “Bureau Veritas Testing GmbH” might be the object of many batches’ certifiedBy edges. The same certificate number might appear in many

audit queries. The graph emerges naturally from shared values.




Figure: Full MSDS ontology graph — classes (green), datatype properties (orange), and object properties shown as a directed graph. Each node type and predicate in this graph is defined in `MSDS_Ontology.ttl` and constrains what the LLM can extract from MSDS documents.

What does this give us that a relational schema does not?

Three things, in order of how much you will care:

1. **No schema changes when facts evolve.** Tomorrow we decide batch certificates should also track `hasReTestDate`. We do not run a DDL migration; we just start storing triples with that predicate. The graph is permissive about what relationships exist.
2. **Queries that follow paths.** “Find me every material whose supplier is in Missouri” — that is a graph traversal. SPARQL expresses it directly. In SQL it would be a chain of joins.
3. **Identity by IRI.** Every node has a globally unique IRI (a URL-shaped identifier). This means a triple in our system can refer to the same material node that any other system uses, simply by sharing an IRI.

For our material document use case we will use the graph mainly for the first reason

— flexibility — and for the *precision* it gives us at query time. A SPARQL query for a certificate number returns the certificate number, not paragraphs.




Figure 3: Knowledge Graph Pipeline

Figure 4.2 — Two variants of the Knowledge Graph ingestion pipeline. **With Ontology (top):** document chunks pass through entity and relationship extraction constrained by the OWL ontology, producing normalised RDF triples that are uploaded to HANA as a named graph. The ontology acts as a schema — it defines which entity types and predicates are valid, so Gemini’s output is predictable and consistent. **Without Ontology (bottom):** the same pipeline runs without the constraint file. Gemini extracts whatever relationships it finds in the text. This is faster to set up but produces less consistent predicate names across documents. Our implementation uses the ontology variant — see `MSDS_Ontology.ttl` in the repository root.

4.2 RDF and SPARQL — the standards explained plainly

RDF stands for Resource Description Framework. It is the W3C standard for expressing triples. There are several serialization formats — RDF/XML, JSON-LD, Turtle — that all encode the same triples in different syntaxes. We will use **Turtle** (file extension `.ttl`) because it is the most readable.

A Turtle file looks like this:

```

1 @prefix msds: <http://msds.knowledge-graph.org/ontology#> .
2
3 msds:BATCH-2024-0871 msds:certifiedBy msds:AcmeSteelAG .
4 msds:BATCH-2024-0871 msds:certificateNumber "QC-CERT-44781" .

```

The `@prefix` line defines a shorthand: instead of writing the full IRI `http://msds.knowledge-graph.org/ontology#BATCH-2024-0871`, we write `msds:BATCH-2024-0871`. Each fact ends with a period. That is essentially all the Turtle you need to read.

SPARQL is the query language for RDF. If SQL is to relational tables, SPARQL is to triples. A simple SPARQL query looks like this:

```

1 SELECT ?cert WHERE {
2   msds:BATCH-2024-0871 msds:certificateNumber ?cert .
3 }

```

The `?cert` is a variable. The query reads: “find me every value of `?cert` such that the triple (BATCH-2024-0871, `certificateNumber`, `?cert`) exists.” Run that against our graph and you get back: “QC-CERT-44781”.

If you can read this much SPARQL, you can read everything our system generates. We are going to lean on Gemini 2.5 Flash to generate more elaborate SPARQL on demand, but the patterns are always shaped like the example above.

Tip: SPARQL has 4 query forms — SELECT, CONSTRUCT, ASK, DESCRIBE. We use only SELECT in this book. You will not miss the others.

4.3 Why HANA Cloud for the Knowledge Graph

SAP HANA Cloud has a built-in graph engine with native SPARQL support. This is architecturally significant for BTP deployments — and it is worth pausing to explain why we did not choose a dedicated triple store.

The alternatives are well-known: Apache Fuseki, AWS Neptune, GraphDB, Stardog. Each is a capable, purpose-built triple store with a standard SPARQL HTTP endpoint. The reason we do not use any of them is simple: we already have HANA Cloud running as our vector store. Adding a separate triple store means a second BTP service, a second

set of credentials, a second Destination Service entry, and a second connection pool to manage in application code. For a production BTP deployment, that is unnecessary complexity.

HANA's SPARQL support is exposed through a single stored procedure: `SPARQL_EXECUTE`. This procedure runs SPARQL queries natively on the same database that holds vector embeddings. One HANA connection handles both `COSINE_SIMILARITY` searches on `MSDS_VECTORS` and SPARQL graph traversals on named RDF graphs. The application code in `vector_srv.py` and `kg_srv.py` both call `get_connection()` from the same `hdb_srv` module. No separate connection string, no separate service binding.

That is the entirety of the HANA SPARQL interface. There is no separate endpoint, no `/sparql` HTTP route. Every query, every insert, every drop — all of it goes through one procedure call:

```
1 cursor.callproc("SPARQL_EXECUTE", (sparql_text, None, None, None))
```

The four arguments are: the SPARQL text, an optional `RDF_DATA` parameter for inline data, an optional `RDF_DATA_TYPE`, and an optional metadata flag. We pass `None` for all three optionals throughout this book.

The interesting part is the *result*. `callproc` returns a tuple of all four arguments after the procedure runs, with output values populated. The query results land in the **fourth element** of that tuple:

```
1 result = cursor.callproc("SPARQL_EXECUTE", (sparql_query, None, None,
    None))
2 rows = result[3] if result and len(result) > 3 else []
```

Every result row is a tuple of column values, one per `?variable` in the `SELECT` clause. So a `SELECT ?value` returning three matches gives you something like:

```
1 [("QC-CERT-44781",), ("ACME Steel AG",), ("Bureau Veritas Testing GmbH",)]
```

Important: This is the single most useful sentence in this chapter. Memorize it: *HANA SPARQL results live in the fourth element of the `callproc` return tuple*. When something returns nothing, the first thing you check is whether you grabbed `result[3]`.

Why a stored procedure and not an HTTP endpoint?

Because HANA already has secure, audited, governed connectivity through its SQL driver. Wrapping SPARQL inside a stored procedure means every SPARQL operation reuses the same authentication, the same connection pooling, the same transaction semantics, and the same logging that a normal SQL statement would. For an enterprise BTP deployment, that uniformity is more valuable than HTTP convenience. One Destination Service entry. One set of HANA credentials. Two retrieval strategies.

4.4 Designing the MSDS ontology — what facts matter?

Before writing any code we must decide what kinds of facts we care about. This decision is called **ontology design**, and it is the most important step of building a Knowledge Graph. The ontology becomes a constraint on what Gemini is allowed to extract — without it, the model invents predicates like `containsAt`, `madeOf`, `relatedTo`, and the graph turns into noise.

For a Batch Quality Certificate, the four predicates that matter for downstream queries are:

Predicate	Domain	Range	Example
<code>certifiedBy</code>	Batch	Supplier	BATCH-2024-0871 □ ACME Steel AG
<code>testResult</code>	Batch	TestResult	BATCH-2024-0871 □ PASS (tensile)
<code>certificateNumber</code>	Batch	CertNumber	BATCH-2024-0871 □ QC-CERT-44781
<code>certifyingLab</code>	Batch	Lab	BATCH-2024-0871 □ Bureau Veritas Testing GmbH

Four predicates. Four classes of node. That is enough to answer the precision questions our hybrid agent will face. If we needed material grade classifications across multiple test standards or transport classifications for logistics, we would extend the ontology. We deliberately stay small.

Note: Real-world supply chain document systems track 50+ predicates — chemical compositions, mechanical test parameters, dimensional tolerances, surface treatment records, invoice line amounts. We pick four because they are enough to demonstrate the architecture and let you focus on the *patterns*. The patterns scale directly to broader ontologies.

4.5 The OWL ontology file — MSDS_Ontology.ttl

Create a file at the root of your project called MSDS_Ontology.ttl. Paste this:

```

1  @prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
2  @prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
3  @prefix owl: <http://www.w3.org/2002/07/owl#> .
4  @prefix msds: <http://msds.knowledge-graph.org/ontology#> .
5
6  # Classes
7  msds:Material a owl:Class ;
8      rdfs:label "Material" .
9
10 msds:Batch a owl:Class ;
11     rdfs:label "Batch / Lot" .
12
13 msds:Supplier a owl:Class ;
14     rdfs:label "Supplier" .
15
16 msds:TestResult a owl:Class ;
17     rdfs:label "Test Result" .
18
19 msds:CertifyingLab a owl:Class ;
20     rdfs:label "Certifying Laboratory" .
21
22 # Properties
23 msds:certifiedBy a owl:ObjectProperty ;
24     rdfs:domain msds:Batch ;
25     rdfs:range msds:Supplier ;
26     rdfs:label "certified by" .
27

```

```

28 msds:testResult a owl:ObjectProperty ;
29     rdfs:domain msds:Batch ;
30     rdfs:range msds:TestResult ;
31     rdfs:label "test result" .
32
33 msds:certificateNumber a owl:DatatypeProperty ;
34     rdfs:domain msds:Batch ;
35     rdfs:label "certificate number" .
36
37 msds:certifyingLab a owl:ObjectProperty ;
38     rdfs:domain msds:Batch ;
39     rdfs:range msds:CertifyingLab ;
40     rdfs:label "certifying lab" .
41
42 msds:description a owl:DatatypeProperty ;
43     rdfs:label "description" .

```

A few things to point out:

- **Classes** are categories of node. `msds:Batch`, `msds:Supplier`, `msds:CertifyingLab`, etc.
- **ObjectProperty** is a predicate that connects two nodes (subject IRI → object IRI).
- **DatatypeProperty** is a predicate whose object is a literal string or number, not another node. We use `msds:description` to hang the actual text value off a node.
- `rdfs:label` is human-readable text. It does not affect the graph; it just makes browsers and tools show better names.
- `rdfs:domain` and `rdfs:range` are the type constraints — domain is the type of the subject, range is the type of the object.

The namespace `http://msds.knowledge-graph.org/` reflects the MSDS reference implementation. For a production multi-document deployment serving multiple document types, the namespace would be organisation-specific — for example `http://materials.yourcompany.com/kg/`. The predicates are generic enough to describe facts across document types: `certifiedBy` works for any batch-to-supplier relationship; `testResult` works for any batch quality test outcome; `certifyingLab` works for laboratory accreditation records. The ontology names a domain, but the structural patterns transfer.

We are not loading this file into HANA at runtime. It is a contract — a documentation

and prompting artifact. Our triple extractor reads this file (or a summary of it) to know what predicates Gemini is allowed to produce.

Tip: If you ever need to share your ontology with another team or load it into a tool like Protégé for visual editing, this is the file you hand them. Keep it under version control alongside your code.

4.6 Extracting triples from MSDS text using Gemini

This is where the LLM earns its keep. We hand Gemini 2.5 Flash a chunk of document text and a list of allowed predicates, and ask it to return JSON triples. Three things make this trustworthy:

1. **Constrained predicates.** The prompt enumerates the four allowed predicates. Anything Gemini emits with a predicate outside that list is filtered out before it reaches HANA.
2. **JSON output.** No prose, no narration — just an array of objects. We strip any markdown fences Gemini sneaks in and `json.loads` the rest.
3. **Material number provided.** The subject of every triple is set externally. The model does not get to invent which batch the facts belong to.

This is an important architectural principle: Gemini 2.5 Flash acts as a structured extraction engine here, not as a data authority. The material number is injected by the CAP upload pipeline, validated against S/4HANA's `API_PRODUCT_SRV` before the document is accepted. Gemini extracts facts from text; it never decides which material those facts belong to.

Here is the prompt, lifted directly from `kg_srv.py`:

```
1 prompt = f"""You are extracting structured facts from a Batch Quality
Certificate document.
2
3 Extract facts ONLY using these relationship types:
4 - certifiedBy: batch → supplier name
5 - testResult: batch → test name and result (e.g. "tensile strength: 487
MPa, PASS")
6 - certificateNumber: batch → certificate number string
```

```

7 - certifyingLab: batch → certifying laboratory name
8
9 Material/Batch identifier: {material_number}
10
11 Return ONLY a JSON array of triples. Each triple must have:
12 - subject: the batch/material IRI
13 - predicate: the property name (certifiedBy, testResult,
14   certificateNumber, certifyingLab)
15 - object: the value (string)
16
17 Example:
18 [
19   {"subject": "BATCH-2024-0871", "predicate": "certifiedBy", "object":
20     "ACME Steel AG"}},
21   {"subject": "BATCH-2024-0871", "predicate": "certificateNumber",
22     "object": "QC-CERT-44781"}},
23   {"subject": "BATCH-2024-0871", "predicate": "testResult", "object":
24     "tensile strength: 487 MPa, PASS"}}
25 ]
26
27 Document text to extract from:
28 {text[:3000]}
29
30 Return only the JSON array, no explanation.

```

Notice the truncation `text[:3000]`. Batch certificates and quality documents can be 20+ pages. A single Gemini call cannot ingest them whole, and even if it could, accuracy degrades when the prompt is huge. In production we slice the document section-by-section (supplier header, test results table, certification footer, etc.). For this chapter we keep one call simple to focus on the mechanics.

Warning: Never trust LLM JSON without parsing defensively. Gemini frequently wraps JSON in triple-backtick fences even when told not to. Strip those before `json.loads`.

The extractor returns a list of dictionaries:

```

1 [
2   {"subject": "BATCH-2024-0871", "predicate": "certifiedBy", "object":
3     "ACME Steel AG"},

```



```

3      {"subject": "BATCH-2024-0871", "predicate": "certificateNumber",
      "object": "QC-CERT-44781"},
4      {"subject": "BATCH-2024-0871", "predicate": "testResult", "object":
      "tensile strength: 487 MPa, PASS"},
5      {"subject": "BATCH-2024-0871", "predicate": "testResult", "object":
      "yield strength: 372 MPa, PASS"},
6      {"subject": "BATCH-2024-0871", "predicate": "certifyingLab", "object":
      "Bureau Veritas Testing GmbH"}
7  ]

```

Each row is one fact. Now we need a place to put them.

4.7 Storing triples in HANA — named graphs per material

Every triple in HANA's SPARQL engine lives in a **named graph**. A named graph is just an IRI that groups a set of triples. Think of it as a schema, but for facts instead of tables.

We follow a simple rule: **one named graph per material**. If we are storing facts about batch BATCH-2024-0871, every triple goes into the graph at IRI `http://msds.knowledge-graph.org/MSDS_Graph/BATCH-2024-0871`. If we are storing facts about BATCH-2024-0992, that material's triples go into a *different* graph IRI.

This pattern maps directly to SAP's material isolation requirements. In SAP MM, material data is scoped to a material number. In the Material Document Intelligence Platform, the same isolation holds: a SPARQL query scoped to the named graph for MAT-001 cannot return facts from MAT-002. The graph boundary enforces data isolation at the database level, not in application logic.

Three practical consequences:

1. **Clean deletion.** When a material is deprecated, `DROP GRAPH <graph-iri>` removes every fact for that material in one statement. There is no per-row cleanup, no orphans.
2. **Multi-tenant safety.** A bad extraction for one batch cannot pollute another batch's facts. Graph boundaries are enforced by HANA's SPARQL engine, not by application-layer filtering.
3. **Auditability.** A SPARQL query scoped to a single named graph can never read facts about other materials. When a compliance auditor asks "show me everything

the system knows about batch X”, the named graph provides a hard scope that the audit trail can reference.

The real graph URIs in the reference implementation follow the pattern `http://msds.knowledge-graph.org/MSDS_Graph/MAT-XXX`, where `MAT-XXX` is the SAP material or batch number validated at upload time. This is the same identifier stored in `MSDS_VECTORS.MATERIAL_NUMBER`. The named graph IRI and the vector table filter column are two representations of the same SAP material identifier.

Here is the storage code, the heart of `store_triples` in `kg_srv.py`:

```

1 graph = graph_iri(material_number)          # the named-graph IRI
2 mat_iri = f"{GRAPH_BASE}/material/{material_number}" # the material's
   node IRI
3 for triple in triples:
4     predicate = triple.get("predicate", "")
5     obj_value = str(triple.get("object", "")).replace("'", "\\'")
6     obj_iri = f"{GRAPH_BASE}/{predicate}/{material_number}/{stored}"
7     sparql_insert = f"""
8         INSERT INTO GRAPH <{graph}> {{
9             <{mat_iri}> <{ONTOLOGY_BASE}{predicate}> <{obj_iri}> .
10            <{obj_iri}> <{ONTOLOGY_BASE}description> '{obj_value}' .
11        }}
12    """
13     cursor.callproc("SPARQL_EXECUTE", (sparql_insert, None, None, None))

```

For each triple we insert two RDF statements:

1. The relationship: `material` → `predicate` → blank-ish node
2. The literal value on that node: `node` → `description` → "ACME Steel AG"

This two-step pattern is deliberate. It lets the same object node carry additional properties later — for instance we might want to add `certification_standard` or `accreditation_body` to a certifying lab node without restructuring the graph. Modeling values as nodes with `description` properties is more verbose than literal-only triples, but it pays off the moment you need to attach metadata.

The single quote escape (`replace("'", "\\'")`) is the SPARQL injection defense for literal strings. It is not as comprehensive as parameterized SQL, which is why we double up on validation upstream — see Section 5.9.

Note: SPARQL Update syntax (`INSERT INTO GRAPH`) is part of the SPARQL 1.1 standard, not SPARQL 1.0. HANA supports it. Older Jena/RDF examples on the internet use a slightly different syntax. Stick with the form above.

4.8 Writing SPARQL queries — Gemini as the SPARQL translator

When a user asks “which supplier certified batch BATCH-2024-0871?”, we cannot send that English directly to HANA. We must translate it to SPARQL.

Gemini 2.5 Flash acts as a SPARQL translator — the user asks in natural language, Gemini writes the SPARQL, HANA executes it. This is a key architectural pattern: the LLM never directly handles data — it only generates queries that the database executes. All data access goes through HANA’s SPARQL engine. This is auditable, controllable, and secure. The data never leaves HANA to be processed by the LLM. Gemini sees the question and the schema; HANA owns the facts.

Two conditions make this work reliably. First, Gemini must know the available predicates and their full IRIs — without this, the model invents predicate names like `hazardClass` that do not exist in our graph. Second, Gemini must know the named graph IRI for this material — without the `GRAPH <iri> { ... }` wrapper, HANA returns zero rows silently.

Here is the heart of `query_graph`:

```

1  sparql_prompt = f"""Generate a SPARQL SELECT query to answer this question
   about material {material_number}.
2
3  Available predicates:
4  - <{ONTOLOGY_BASE}certifiedBy> — links batch to supplier nodes
5  - <{ONTOLOGY_BASE}testResult> — links batch to test result nodes
6  - <{ONTOLOGY_BASE}certificateNumber> — links batch to certificate number
   nodes
7  - <{ONTOLOGY_BASE}certifyingLab> — links batch to certifying lab nodes
8  - <{ONTOLOGY_BASE}description> — literal value on any node
9
10 The material IRI is: <{mat_iri}>

```

```

11 The named graph is: <{graph}>
12
13 Question: {question}
14
15 Return ONLY the SPARQL query, wrapped in GRAPH clause like this:
16 SELECT ?value WHERE {{
17     GRAPH <{graph}> {{
18         <{mat_iri}> <predicate> ?node .
19         ?node <{ONTOLOGY_BASE}description> ?value .
20     }}
21 }}""

```

The example query in the prompt is the **only SPARQL pattern** the rest of this book needs. Read it again. It does two hops:

1. From the batch node, follow some predicate to a supplier/lab/certificate node.
2. From that node, follow description to the literal text.

Two hops, one SELECT. That is the entire pattern.

Important: The `GRAPH <iri> { ... }` wrapper is mandatory in HANA. Without it, HANA does not throw an error — it just returns zero rows. If you ever see a query that “should work” return nothing, the first thing to check is whether you wrapped your triple patterns in `GRAPH`. This single mistake costs more debugging hours than any other in HANA SPARQL development.

For the question “which supplier certified this batch?”, Gemini emits something like:

```

1 SELECT ?value WHERE {
2     GRAPH <http://msds.knowledge-graph.org/MSDS_Graph/BATCH-2024-0871> {
3         <http://msds.knowledge-graph.org/MSDS_Graph/material/BATCH-2024-0871>
4             <http://msds.knowledge-graph.org/ontology#certifiedBy> ?node .
5         ?node <http://msds.knowledge-graph.org/ontology#description> ?value .
6     }
7 }

```

We pass it through `SPARQL_EXECUTE`, grab `result[3]`, and return `["ACME Steel AG"]`. Done.

The data stayed in HANA throughout. Gemini generated a query string; HANA executed it against the stored graph. The LLM never touched the supplier name directly.

4.9 Building `kg_srv.py` — the Knowledge Graph service layer

We collect everything into a single service module at `agents/srv/kg_srv.py`. It exposes five functions: `extract_triples`, `store_triples`, `query_graph`, `delete_graph`, `count_triples`. Create the file and paste this:

```

1  import os
2  import re
3  import json
4  from .hdb_srv import get_connection
5  from .vertex_srv import get_llm
6
7  GRAPH_BASE = "http://msds.knowledge-graph.org/MSDS_Graph"
8  ONTOLOGY_BASE = "http://msds.knowledge-graph.org/ontology#"
9
10 MATERIAL_RE = re.compile(r"^[A-Za-z0-9_-]+$")
11
12 def _validate_material(material_number: str):
13     if not MATERIAL_RE.match(material_number):
14         raise ValueError(f"Invalid material number: {material_number}")
15
16 def graph_iri(material_number: str) -> str:
17     _validate_material(material_number)
18     return f"{GRAPH_BASE}/{material_number}"
19
20 def extract_triples(text: str, material_number: str) -> list[dict]:
21     _validate_material(material_number)
22     llm = get_llm()
23     prompt = f"""You are extracting structured facts from a Batch Quality
Certificate document.
24
25 Extract facts ONLY using these relationship types:
26 - certifiedBy: batch → supplier name
27 - testResult: batch → test name and result (e.g. "tensile strength: 487
MPa, PASS")

```

```

28 - certificateNumber: batch → certificate number string
29 - certifyingLab: batch → certifying laboratory name
30
31 Material/Batch identifier: {material_number}
32
33 Return ONLY a JSON array of triples. Each triple must have:
34 - subject: the batch/material IRI
35 - predicate: the property name (certifiedBy, testResult,
    certificateNumber, certifyingLab)
36 - object: the value (string)
37
38 Example:
39 [
40   {"subject": "BATCH-2024-0871", "predicate": "certifiedBy", "object":
    "ACME Steel AG"}},
41   {"subject": "BATCH-2024-0871", "predicate": "certificateNumber",
    "object": "QC-CERT-44781"}},
42   {"subject": "BATCH-2024-0871", "predicate": "testResult", "object":
    "tensile strength: 487 MPa, PASS"}}
43 ]
44
45 Document text to extract from:
46 {text[:3000]}
47
48 Return only the JSON array, no explanation.""
49
50 response = llm.generate_content(prompt)
51 raw = response.text.strip()
52 # Strip markdown code blocks if present
53 raw = re.sub(r"```json\s*", "", raw)
54 raw = re.sub(r"```s*", "", raw)
55 return json.loads(raw)
56
57 def store_triples(material_number: str, triples: list[dict]):
58     _validate_material(material_number)
59     if not triples:
60         return 0
61     graph = graph_iri(material_number)
62     mat_iri = f"{GRAPH_BASE}/material/{material_number}"

```

```

63     conn = get_connection()
64     cursor = conn.cursor()
65     stored = 0
66     for triple in triples:
67         predicate = triple.get("predicate", "")
68         obj_value = str(triple.get("object", "")).replace("'", "\\'")
69         obj_iri = f"{GRAPH_BASE}/{predicate}/{material_number}/{stored}"
70         sparql_insert = f"""
71             INSERT INTO GRAPH <{graph}> {{
72                 <{mat_iri}> <{ONTOLOGY_BASE}{predicate}> <{obj_iri}> .
73                 <{obj_iri}> <{ONTOLOGY_BASE}description> '{obj_value}' .
74             }}
75         """
76         try:
77             cursor.callproc("SPARQL_EXECUTE", (sparql_insert, None, None,
78 None))
79             stored += 1
80         except Exception as e:
81             print(f"Warning: failed to store triple {triple}: {e}")
82     conn.commit()
83     cursor.close()
84     return stored
85
86 def query_graph(material_number: str, question: str) -> dict:
87     _validate_material(material_number)
88     graph = graph_iri(material_number)
89     mat_iri = f"{GRAPH_BASE}/material/{material_number}"
90     llm = get_llm()
91
92     # Generate SPARQL from the question
93     sparql_prompt = f"""Generate a SPARQL SELECT query to answer this
94 question about material {material_number}.
95
96 Available predicates:
97 - <{ONTOLOGY_BASE}certifiedBy> – links batch to supplier nodes
98 - <{ONTOLOGY_BASE}testResult> – links batch to test result nodes
99 - <{ONTOLOGY_BASE}certificateNumber> – links batch to certificate number
100 nodes
101 - <{ONTOLOGY_BASE}certifyingLab> – links batch to certifying lab nodes

```

```

99 - <{ONTOLOGY_BASE}description> – literal value on any node
100
101 The material IRI is: <{mat_iri}>
102 The named graph is: <{graph}>
103
104 Question: {question}
105
106 Return ONLY the SPARQL query, wrapped in GRAPH clause like this:
107 SELECT ?value WHERE {{
108     GRAPH <{graph}> {{
109         <{mat_iri}> <predicate> ?node .
110         ?node <{ONTOLOGY_BASE}description> ?value .
111     }}
112 }}""
113
114 sparql_response = llm.generate_content(sparql_prompt)
115 sparql_query = sparql_response.text.strip()
116 sparql_query = re.sub(r"```sparql\s*", "", sparql_query)
117 sparql_query = re.sub(r"```s*", "", sparql_query).strip()
118
119 # Execute SPARQL
120 conn = get_connection()
121 cursor = conn.cursor()
122 try:
123     result = cursor.callproc("SPARQL_EXECUTE", (sparql_query, None,
124 None, None))
125     rows = result[3] if result and len(result) > 3 else []
126     facts = [str(row[0]) for row in rows] if rows else []
127 except Exception as e:
128     facts = []
129     sparql_query = f"Error: {e}"
130 finally:
131     cursor.close()
132
133 return {
134     "facts": facts,
135     "sparql": sparql_query,
136     "count": len(facts)
137 }

```



```

137
138 def delete_graph(material_number: str) -> bool:
139     _validate_material(material_number)
140     graph = graph_iri(material_number)
141     conn = get_connection()
142     cursor = conn.cursor()
143     try:
144         cursor.callproc("SPARQL_EXECUTE", (f"DROP GRAPH <{graph}>", None,
None, None))
145         conn.commit()
146         cursor.close()
147         return True
148     except Exception as e:
149         cursor.close()
150         print(f"Warning: could not delete graph {graph}: {e}")
151         return False
152
153 def count_triples(material_number: str) -> int:
154     _validate_material(material_number)
155     graph = graph_iri(material_number)
156     conn = get_connection()
157     cursor = conn.cursor()
158     try:
159         result = cursor.callproc("SPARQL_EXECUTE", (
160             f"SELECT (COUNT(*) AS ?count) WHERE {{ GRAPH <{graph}> {{ ?s ?p
?o }} }}",
161             None, None, None
162         ))
163         rows = result[3] if result and len(result) > 3 else []
164         cursor.close()
165         return int(rows[0][0]) if rows else 0
166     except Exception:
167         cursor.close()
168         return 0

```

Validation: the KG equivalent of SQL injection

Take a close look at this regex:

```
1 MATERIAL_RE = re.compile(r"^[A-Za-z0-9_-]+$")
```

It is enforced at the entry of every public function via `_validate_material`. Why is this so important?

The material number is **interpolated into IRIs** that we paste into SPARQL text. If someone calls `extract_triples(text, "../../../admin")`, the resulting graph IRI becomes `http://msds.knowledge-graph.org/MSDS_Graph/../../../admin`. Worse, a string like `> { ?s ?p ?o } DROP GRAPH <http://something` could close the IRI brackets and inject a destructive SPARQL update.

This is the SPARQL equivalent of SQL injection. The regex `^[A-Za-z0-9_-]+$` allows only letters, digits, underscore, and hyphen — characters that have no special meaning in SPARQL syntax. Any input containing slashes, spaces, angle brackets, quotes, or other SPARQL metacharacters is rejected before it can reach the query.

The same material number format is enforced in the CAP upload handler before the document reaches the Python layer at all, matching it against S/4HANA's product master. Validation at the CAP layer, at the Python service entry, and at the SQL parameter level — three checkpoints for a single identifier.

Warning: If you change this regex to be more permissive — say, to allow forward slashes for hierarchical material codes — you must update the storage and query code to escape those characters too. The simplest defense is to keep the regex tight and document the constraint for upstream callers.

We also escape single quotes in literal values inside `store_triples`. Combined, validation on identifiers and escaping on literals cover the two injection vectors.

What `delete_graph` and `count_triples` give us

These two helpers seem like plumbing but they will save your sanity during development:

- `count_triples(material_number)` — confirms data actually landed in HANA. Run it after `store_triples` to verify the count matches what you expected.
- `delete_graph(material_number)` — wipes a material's graph clean. Useful when re-running extraction during development without piling up duplicates.

They both exercise the same pattern as `query_graph` and prove that you understand

SPARQL_EXECUTE. If you can write `count_triples`, you can write any read query.

4.10 Testing — store triples for a batch certificate, query for supplier and certificate number

Create `agents/test_kg.py`:

```

1  import os
2  from dotenv import load_dotenv
3  from srv.kg_srv import extract_triples, store_triples, query_graph,
   count_triples
4
5  load_dotenv()
6
7  TEST_MATERIAL = "BATCH-2024-0871"
8  TEST_TEXT = """
9  Batch Quality Certificate
10
11  Material: Steel rod grade S355
12  Batch lot number: BATCH-2024-0871
13  Supplier: ACME Steel AG
14  Certificate number: QC-CERT-44781
15  Certification date: 2024-03-15
16
17  Test Results:
18  - Tensile strength: measured 487 MPa, specification minimum 355 MPa — PASS
19  - Yield strength: measured 372 MPa, specification minimum 345 MPa — PASS
20  - Elongation at break: measured 26%, specification minimum 22% — PASS
21
22  Certifying laboratory: Bureau Veritas Testing GmbH, Hamburg, Germany
23  Test standard: ISO 6892-1 (metallic materials tensile testing)
24  """
25
26  print("Extracting triples from batch certificate text...")
27  triples = extract_triples(TEST_TEXT, TEST_MATERIAL)
28  print(f"Extracted {len(triples)} triples:")

```

```

29 for t in triples:
30     print(f"    {t['predicate']}: {t['object']}")
31
32 print(f"\nStoring triples in HANA named graph...")
33 stored = store_triples(TEST_MATERIAL, triples)
34 print(f"Stored {stored} triples")
35
36 total = count_triples(TEST_MATERIAL)
37 print(f"Total triples in graph: {total}")
38
39 print("\nQuerying: 'Which supplier certified this batch?'")
40 result = query_graph(TEST_MATERIAL, "Which supplier certified this
    batch?")
41 print(f"SPARQL generated:\n{result['sparql']}")
42 print(f"Facts found: {result['facts']}")
43
44 print("\nQuerying: 'What is the certificate number?'")
45 result = query_graph(TEST_MATERIAL, "What is the certificate number?")
46 print(f"Facts found: {result['facts']}")

```

Run it from the agents/ directory:

```

1 cd agents
2 python test_kg.py

```

Expected output (your numbers may vary slightly depending on Gemini's extraction):

Extracting triples from batch certificate text...

Extracted 5 triples:

```

certifiedBy: ACME Steel AG
certificateNumber: QC-CERT-44781
testResult: tensile strength: 487 MPa, PASS
testResult: yield strength: 372 MPa, PASS
certifyingLab: Bureau Veritas Testing GmbH

```

Storing triples in HANA named graph...

Stored 5 triples

Total triples in graph: 10

Querying: 'Which supplier certified this batch?'

SPARQL generated:

```
SELECT ?value WHERE {  
  GRAPH <http://msds.knowledge-graph.org/MSDS_Graph/BATCH-2024-0871> {  
    <http://msds.knowledge-graph.org/MSDS_Graph/material/BATCH-2024-0871>  
      <http://msds.knowledge-graph.org/ontology#certifiedBy> ?node .  
    ?node <http://msds.knowledge-graph.org/ontology#description> ?value .  
  }  
}
```

Facts found: ['ACME Steel AG']

Querying: 'What is the certificate number?'

Facts found: ['QC-CERT-44781']

Note: The total of 10 triples — twice the number of facts — is the two-statement pattern from Section 4.7. Each fact stores one batch → predicate → node edge plus one node → description → literal edge.

If your test prints `Facts found: []`, three things to check, in this order:

1. Did `store_triples` actually run before the query? Run `count_triples` and confirm the number is non-zero.
2. Is your generated SPARQL wrapped in `GRAPH <...>`? Print `result['sparql']` and inspect.
3. Did `result[3]` come back populated? Add a debug print right after `cursor.callproc` to confirm.

4.11 What the KG gets right that vector search cannot

Now the comparison promised since Chapter 3. Run the same question through both retrievers.

Question: “What is the certificate number for this batch?”

Vector search (Chapter 3):

```

1 from srv.vector_srv import similarity_search
2 hits = similarity_search("What is the certificate number for this batch?",
3                           k=3)
4 for h in hits:
5     print(f"score={h.score:.3f}  text={h.text[:120]}...")

```

Output:

score=0.651 text=Certificate number: QC-CERT-44781. Batch lot number: BATCH-2024-0871. Supplier: ACME Steel AG. Certification date: 2024...

score=0.602 text=Certifying laboratory: Bureau Veritas Testing GmbH, Hamburg, Germany. T
1...

score=0.578 text=Test Results – Tensile strength: measured 487 MPa, specification minimum
PASS...

The top hit *contains* the certificate number at score ~0.65. To produce the answer “QC-CERT-44781” you must hand the chunk to an LLM with another prompt: “extract the certificate number from this text”. Two LLM calls, ~600ms each, fuzzy on the inputs and outputs, and you might miss it if the identifier appears in multiple contexts or if the scoring shifts on a different embedding model.

Knowledge graph (this chapter):

```

1 result = query_graph("BATCH-2024-0871", "What is the certificate number?")
2 print(result['facts'])

```

Output:

['QC-CERT-44781']

One string. The exact identifier. No paragraph to parse, no scoring, no chunk boundary luck. The graph either has the fact or it does not — and if it has it, the query returns it. Gemini generated the SPARQL; HANA executed it; the facts came back directly. The LLM was a query generator, not a data processor.

The right answer to “What is the certificate number for batch BATCH-2024-0871?” is the exact string QC-CERT-44781 — not a paragraph, not a summary, not the surrounding context.

This is the contract:

Vector store	Knowledge graph
Returns text passages	Returns structured values
Ranks by similarity (0–1 score)	Returns exact matches (yes/no)
“How did the supplier describe the tensile test method?”	“What is the certificate number?”
Tolerates fuzzy phrasing	Demands precise predicate semantics
Best for methodology descriptions, inspection narratives	Best for identifiers, test results, dates, names
Failure mode: irrelevant chunk near top	Failure mode: empty result if predicate missing

Neither one wins on its own. A user asking “*what does the certificate say about the test methodology used for tensile measurement?*” needs the prose chunk that paints the full picture. A user asking “*what is the certificate number?*” needs the identifier verbatim. The whole point of the next chapter — building the LangGraph supervisor — is to choose between them, or run both and merge the answers.

Tip: A useful intuition: the vector store remembers *what was said*. The Knowledge Graph remembers *what is true*. The first is great for context, the second is great for facts. Most non-trivial enterprise document questions need both.

4.12 Summary

You now have a working Knowledge Graph layer on HANA Cloud. Specifically:

- An ontology (`MSDS_Ontology.ttl`) that defines four predicates (`certifiedBy`, `testResult`, `certificateNumber`, `certifyingLab`) and acts as a contract for what facts the system understands. The namespace `http://msds.knowledge-graph.org/` reflects the MSDS reference implementation; production deployments use an organisation-specific namespace.
- A `kg_srv.py` service module exposing `extract_triples`, `store_triples`, `query_graph`, `count_triples`, and `delete_graph`.
- A working pipeline that pulls JSON triples from document text using Gemini 2.5 Flash, stores them in per-material named graphs in HANA, and queries them with

auto-generated SPARQL.

- Named graphs scoped per material or batch number, providing the same material isolation that SAP MM enforces for product master data. Graph `MSDS_Graph/BATCH-2024-0871` cannot return facts from `MSDS_Graph/BATCH-2024-0992`.
- Validated material identifiers and escaped literals — defense against the SPARQL equivalent of injection.
- The pattern `cursor.callproc("SPARQL_EXECUTE", (query, None, None, None))` followed by `result[3]` for rows.
- The mandatory `GRAPH <iri> { ... }` wrapper for every triple pattern.
- No separate triple store required. One HANA Cloud instance, accessed through one connection pool, serves both vector embeddings and RDF graph traversals.

The investment was small — under 200 lines of Python, one Turtle file, and a single stored procedure call — but the capability is fundamentally different from what vector search alone gave us. We can now answer fact-shaped questions with fact-shaped answers. And we can do it without adding a second BTP service or a second set of credentials.

4.13 Checkpoint

Before moving to Chapter 5, confirm the following from the `agents/` directory:

```
1 python -c "from srv.kg_srv import count_triples;
   print(count_triples('BATCH-2024-0871'))"
```

This should print a number greater than zero — the count of triples currently stored for the test batch.

```
1 python -c "from srv.kg_srv import query_graph;
   print(query_graph('BATCH-2024-0871', 'Which supplier certified this
   batch?')['facts'])"
```

This should print a list containing `ACME Steel AG`.

If both succeed, you have:

- HANA Cloud reachable through `hdb_srv`
- Gemini 2.5 Flash generating valid SPARQL through the `get_llm()` getter in `vertex_srv`
- A populated named graph for BATCH-2024-0871 at http://msds.knowledge-graph.org/MSDS_Graph/BATCH-2024-0871
- Working extraction, storage, and query paths

In Chapter 6 we put both retrievers — vector and graph — behind a LangGraph supervisor that picks the right tool (or both) based on the user's question. The hybrid in *Hybrid RAG* finally starts to make sense.

Chapter 5: The Document Intelligence Ingestion Pipeline

In Chapter 4 you built two storage systems on SAP HANA Cloud: a vector table that answers fuzzy semantic questions and a Knowledge Graph that answers precise factual ones. Both are powerful in isolation. The missing piece is the plumbing — the code that takes any PDF file uploaded against an SAP Material Number and feeds both systems at the same time, cleanly, without one pipeline blocking the other, and without leaving the user waiting at a spinning upload button.

That is what this chapter builds. By the end of it you will have `agents/srv/doc_srv.py`, a production-quality ingestion service that receives a PDF upload, immediately acknowledges the request, and runs two threads simultaneously — one chunking and embedding text into the vector store, the other extracting structured triples into the Knowledge Graph. The caller gets a response in milliseconds. The heavy lifting happens asynchronously.

The pipeline is document-type agnostic at the infrastructure level. Whether the uploaded file is an MSDS, a purchase order, a batch certificate, or a quality inspection report, the same dual-thread architecture applies. The vector pipeline always chunks, embeds, and stores to HANA `REAL_VECTOR`. The KG pipeline always calls Gemini and stores triples via SPARQL. The only thing that varies per document type is the Gemini prompt used for triple extraction — a single configuration point that determines what structured knowledge the pipeline extracts from that document class.

There is a subtlety that trips up almost every developer the first time they write multi-threaded database code in Python: connection sharing. A single HANA connection is not thread-safe. If Thread 1 and Thread 2 share one connection object, you will see intermittent failures that look like data corruption, because they are. This chapter explains why that happens, how Python's `threading.local()` solves it, and how to wire the solution correctly inside a FastAPI background task.

5.1 The dual-pipeline architecture

The fundamental design choice in this chapter is that every PDF upload triggers exactly two independent pipelines, running in parallel, with no coordination between them except that they both read from the same source file.

DUAL-PIPELINE PDF INGESTION ARCHITECTURE




Figure 5.1: The Document Intelligence Ingestion Pipeline. Thread 1 (left) chunks the PDF text and stores embeddings in `MSDS_VECTORS`. Thread 2 (right) sends full text to Gemini for triple extraction and stores the result as a named RDF graph in HANA. Both threads update their own status field on completion. The FastAPI endpoint returns immediately, before either thread has finished.

5.1.1 Why process both in parallel?

The two pipelines are genuinely independent operations. The vector pipeline calls `text-embedding-004` to generate embeddings for each chunk. The KG pipeline calls Gemini 2.5 Flash to generate RDF triples from the full document text. Neither result depends on the other. There is no reason to run them sequentially — doing so would simply double ingestion time.

Dimension	Vector pipeline	KG pipeline
Input to pipeline	Chunks (500 tokens each)	Full document text

Dimension	Vector pipeline	KG pipeline
External API call	text-embedding-004 (many calls)	Gemini 2.5 Flash (one call)
HANA operation	Many INSERT statements	One SPARQL_EXECUTE insert
Failure mode	Can fail on any chunk	Atomic — succeeds or fails entirely

On a 30-page batch quality certificate or MSDS document the difference is measurable. Embedding 60 chunks against the Vertex AI API takes roughly 8–12 seconds. Calling Gemini 2.5 Flash for triple extraction from the same document takes 3–6 seconds. Sequential: 14–18 seconds. Parallel: 8–12 seconds (the slower pipeline dominates). The improvement compounds across every document uploaded across the platform.

`ThreadPoolExecutor(max_workers=2)` is the implementation mechanism. It creates exactly two OS threads within the FastAPI process — one for each pipeline. This is true parallelism: the embedding API calls in Thread 1 and the Gemini call in Thread 2 execute concurrently, sharing nothing except the temp file path.

5.1.2 The fire-and-forget pattern

The CAP Fiori UI calls `POST /process-upload` and needs a fast response. A user who clicks the upload button should not wait 10–15 seconds staring at a loading indicator — that is an unacceptable user experience for an enterprise application.

The solution is fire-and-forget. The upload endpoint returns `202 Accepted` with `{"status": "processing"}` immediately after spawning the background threads — before either pipeline has done any significant work. The HTTP connection closes. The threads run independently.

The CAP service then polls `GET /status/{material_number}` every 3 seconds. The Fiori UI updates two status indicators — one for vector ingestion, one for KG ingestion — based on what the status endpoint returns. When both reach a terminal state (`DONE` or `ERROR`), polling stops and the UI updates the document card.

This decoupling is critical for BTP Cloud Foundry deployments. HTTP request timeouts on CF are typically 60 seconds. A large document (100+ pages) can take 90 seconds to

process. Without fire-and-forget, those documents would always fail at the HTTP layer, even when the actual processing succeeded.

Status is persisted in HANA — in the `MSDS_DOCUMENTS` table — not in memory. This means status survives a server restart. If the CF instance is recycled while processing is in flight, the status row already exists in HANA showing `PROCESSING`. A monitoring query can find and retry stalled documents without reading log files.

5.2 PDF text extraction with PyMuPDF

PyMuPDF (import name `fitz`) is a Python binding for the MuPDF rendering library. It is the fastest Python-accessible PDF parser available and handles the dirty reality of enterprise PDFs well: multi-column layouts, embedded tables, rotated pages, and the mixture of text layers and bitmap content that characterises scanned regulatory documents.

Add it to `agents/requirements.txt`:

```
PyMuPDF==1.24.3
```

The basic extraction pattern is three lines:

```
1 import fitz # PyMuPDF
2
3 def extract_text(pdf_path: str) -> str:
4     doc = fitz.open(pdf_path)
5     text = "".join(page.get_text() for page in doc)
6     doc.close()
7     return text
```

`page.get_text()` with no arguments returns the page's text content in reading order, with word-level coordinates used internally to reconstruct flow. For most material documents this produces clean, paragraph-structured text. For scanned PDFs (no text layer) it returns an empty string — handled below.

5.2.1 Extracting with metadata

In `agents/srv/doc_srv.py` we use a slightly richer version that preserves page boundaries:

```

1  import fitz
2  from typing import Optional
3
4  def extract_pdf_text(path: str) -> tuple[str, int]:
5      """
6      Returns (full_text, page_count).
7      Raises ValueError if the PDF contains no extractable text layer.
8      """
9      doc = fitz.open(path)
10     pages = []
11     for page in doc:
12         page_text = page.get_text().strip()
13         if page_text:
14             pages.append(page_text)
15     doc.close()
16
17     if not pages:
18         raise ValueError(
19             f"PDF '{path}' contains no extractable text. "
20             "It may be a scanned image-only document."
21         )
22
23     return "\n\n".join(pages), len(pages)

```

The `ValueError` is important. If a scanned PDF arrives and we silently extract an empty string, the vector pipeline produces zero embeddings and the KG pipeline sends an empty prompt to Gemini. Both pipelines appear to succeed, but the document is never searchable. Raising early surfaces the problem immediately and sets both status columns to `ERROR` with a meaningful message — visible to operators without reading log files.

Warning: PyMuPDF’s `get_text()` preserves hyphenation artifacts from PDF layout. Words broken across lines (e.g., “flamma-” + newline + “ble”) will appear as a hyphenated pair in the extracted string. For material documents this rarely causes problems because the affected words are typically in multi-column document tables, but if your documents have heavy hyphenation, consider a post-processing step: `text = re.sub(r'\n', ' ', text)`.

5.2.2 Handling upload files in FastAPI

FastAPI receives uploaded files as `UploadFile` objects. We need to write the file to disk before passing the path to PyMuPDF, because `fitz.open()` accepts file paths or byte buffers — and the byte buffer approach for large PDFs can exhaust memory. The safest pattern uses Python's `tempfile`:

```
1 import tempfile
2 import shutil
3 from fastapi import UploadFile
4
5 async def save_upload(upload: UploadFile) -> str:
6     """Saves UploadFile to a temp file. Returns the path. Caller must
7     delete."""
8     suffix = ".pdf"
9     with tempfile.NamedTemporaryFile(delete=False, suffix=suffix) as tmp:
10         shutil.copyfileobj(upload.file, tmp)
11     return tmp.name
```

Using `delete=False` gives us a path we can pass to background threads. We delete the file ourselves after both threads complete.

5.3 Chunking strategy for vector search

The vector pipeline does not process the full document text as a single unit. A 30-page material document contains 8,000–12,000 words. Embedding that as one unit would produce a single 768-dimensional vector that averages the meaning of everything — test results, storage requirements, delivery conditions, specification tables — into an undifferentiated blob. Cosine similarity against that vector would return the document for almost any question. Precision would be zero.

The answer is chunking: splitting the document into smaller, semantically coherent pieces and embedding each independently. The cosine search then retrieves the specific chunks most relevant to a question, not the whole document.

5.3.1 Our chunking parameters

We use three parameters:

Parameter	Value	Rationale
Chunk size	500 tokens	Roughly 375 words. Large enough to preserve context around a fact, small enough for the vector to encode a focused meaning.
Overlap	50 tokens	Approximately one sentence. Prevents a fact from being split exactly at a chunk boundary and disappearing from both adjacent chunks.
Split boundary	Sentence end	Never cut in the middle of a sentence. A half-sentence changes the meaning of the embedding.

The `text-embedding-004` model from Vertex AI accepts up to 3,072 tokens per input. Our 500-token chunks are well within that limit, which gives us room to include a prefix that grounds the embedding with document metadata:

```
Material: BATCH-QC-MAT-001\n\n<chunk text>
```

Adding the material name to every chunk significantly improves retrieval accuracy for multi-document queries. Without it, a chunk about “tensile test” could match any of dozens of batch certificates.

This prefix approach is also what makes the vector pipeline document-type agnostic. The same chunker, with the same parameters, handles an MSDS equally well as a purchase order or a quality inspection certificate — the material prefix ensures retrieved chunks are always scoped to the correct document.

5.3.2 The chunker function

```

1  import re
2  from typing import Generator
3
4  def chunk_text(
5      text: str,
6      material_name: str,
7      chunk_tokens: int = 500,
8      overlap_tokens: int = 50,
9  ) -> Generator[str, None, None]:
10     """
11     Yields overlapping chunks of approximately chunk_tokens each,
12     split on sentence boundaries. Prefixes each chunk with material_name.
13
14     Token estimation: 1 token ≈ 0.75 words (English prose).
15     """
16     chunk_words = int(chunk_tokens * 0.75)
17     overlap_words = int(overlap_tokens * 0.75)
18
19     # Split on sentence-ending punctuation followed by whitespace
20     sentences = re.split(r'(?<=[.!?])\s+', text.strip())
21
22     current_chunk: list[str] = []
23     current_count: int = 0
24     prefix = f"Material: {material_name}\n\n"
25
26     for sentence in sentences:
27         word_count = len(sentence.split())
28         if current_count + word_count > chunk_words and current_chunk:
29             yield prefix + " ".join(current_chunk)
30             # Keep the last overlap_words words as context for the next
31             # chunk
32             overlap_start = max(0, len(current_chunk) - overlap_words)
33             current_chunk = current_chunk[overlap_start:]
34             current_count = sum(len(s.split()) for s in current_chunk)
35             current_chunk.append(sentence)
36             current_count += word_count

```

```
37     # Emit the final partial chunk if it has content
38     if current_chunk:
39         yield prefix + " ".join(current_chunk)
```

Tip: This chunker uses a word-count approximation for token count (1 token $\approx$ 0.75 words). For production use, replace the approximation with the `tiktoken` library: `len(tiktoken.encoding_for_model("text-embedding-ada-002").encode(text))`. The approximation is accurate enough for material documents, where prose density is fairly uniform.

5.3.3 Why the KG pipeline does not chunk

The Knowledge Graph pipeline receives the **full document text**, not chunks. This is a deliberate asymmetry.

Gemini's triple extraction works best when it can see the entire document context. A certificate number (QC-CERT-44781) might appear in the batch header section of a quality certificate. The material name appears in the header. The certifying laboratory appears in the test results section. If we feed Gemini a chunk that contains only the test results section, it cannot associate the certificate number with the material and certifying lab, because those context clues are in different sections.

For the KG pipeline, the goal is to extract a small number of high-precision facts (typically 5–15 triples per document). Gemini 2.5 Flash can easily handle 10,000 words in a single prompt, so there is no technical reason to chunk. The full-document approach produces far better triple quality.

The KG extraction prompt is the single configuration point that varies per document type. An invoice prompt instructs it to extract vendor, line items, and payment terms. A batch certificate prompt extracts test results, pass/fail verdicts, and certifying laboratory. A regulatory document prompt extracts hazard codes, exposure limits, and precaution statements. The chunker code and the SPARQL insert code remain identical across all document types.

5.4 Thread-local HANA connections

This section addresses the most important implementation detail in this chapter. Get it wrong and you will see intermittent `hdbcli` errors that are extremely difficult to debug. Understand it and the threading model becomes straightforward.

5.4.1 Why connections cannot be shared across threads

The `hdbcli` driver maintains a stateful protocol session on each connection object. Internally, a connection tracks the current transaction state, pending fetch buffers for open cursors, and the sequence of SQL commands in the current batch. When two threads share a single connection and execute concurrently, they can each call `cursor.execute()` at the same time. The driver serialises the calls — but by the time the second call returns, the first thread may have already called `cursor.fetchall()` on a different cursor that was internally sharing the same fetch buffer. The result is that one thread reads the other thread's data, or gets an exception, or silently gets an empty result.

None of these failures are deterministic. They depend on exact scheduling and timing. This is the canonical definition of a race condition.

Warning: Never share an `hdbcli` connection across threads. The `hdbcli` documentation explicitly states that connection objects are not thread-safe. This applies to most database drivers, not just HANA. When in doubt, assume a connection is single-threaded.

5.4.2 The `threading.local()` solution

Python's `threading.local()` provides a simple mechanism: an object where each attribute access returns a different value depending on which thread is reading it. When Thread 1 sets `_local.conn = hana_db.connect()`, Thread 2 sees `_local.conn` as unset. Each thread must initialise its own value.

Here is the pattern we use in `agents/srv/doc_srv.py`:

```
1 import threading
2 from hdbcli import dbapi
3
```

```

4  _thread_local = threading.local()
5
6  def get_hana_connection() -> dbapi.Connection:
7      """
8      Returns a HANA connection for the current thread.
9      Creates one on first access; reuses it on subsequent calls within the
10     same thread.
11     """
12     if not hasattr(_thread_local, "conn") or _thread_local.conn is None:
13         _thread_local.conn = dbapi.connect(
14             address=os.environ["HANA_HOST"],
15             port=int(os.environ["HANA_PORT"]),
16             user=os.environ["HANA_USER"],
17             password=os.environ["HANA_PASSWORD"],
18         )
19     return _thread_local.conn
20
21 def close_thread_connection() -> None:
22     """Closes and clears the connection for the current thread."""
23     if hasattr(_thread_local, "conn") and _thread_local.conn is not None:
24         try:
25             _thread_local.conn.close()
26         except Exception:
27             pass
28     _thread_local.conn = None

```

The `get_hana_connection()` function is called at the start of each thread function. If the thread is new, it creates a fresh connection. If the same thread somehow reuses its worker slot (which can happen with thread pools), it reuses the existing connection — which is safe because it is still the same thread.

Note: The thread pool created by `ThreadPoolExecutor` may reuse worker threads across multiple task submissions if the executor stays alive. In our implementation, we create a fresh `ThreadPoolExecutor` for each upload, so each worker thread is new and its `threading.local()` storage is empty. If you ever switch to a long-lived executor, add a health check before reusing a connection: `_thread_local.conn.isconnected()`.

5.5 The ingestion flow in detail

Now we have all the pieces. Let us trace the full execution path for a single document upload from the CAP Fiori UI through to both HANA stores.

5.5.1 The upload endpoint

```

1  from fastapi import FastAPI, File, Form, UploadFile
2  from fastapi.responses import JSONResponse
3  import asyncio, os, re, tempfile, shutil
4  from concurrent.futures import ThreadPoolExecutor, as_completed
5
6  app = FastAPI()
7
8  @app.post("/process-upload")
9  async def process_upload(
10     file: UploadFile = File(...),
11     materialNumber: str = Form(...),
12     materialName: str = Form(...),
13 ):
14     if not re.match(r'^[A-Za-z0-9_-]+$ ', materialNumber):
15         return JSONResponse(
16             status_code=400,
17             content={"error": "Invalid materialNumber. Use only letters,
18                 digits, hyphens, underscores."}
19         )
20
21     # Write upload to a temp file so background threads can access it by
22     # path
23     loop = asyncio.get_event_loop()
24     tmp_path = await loop.run_in_executor(None, _save_upload_sync, file)
25
26     # Mark both pipelines as PROCESSING in HANA before returning
27     _mark_processing(materialNumber)
28
29     # Hand off to background threads – fire and forget
30     loop.run_in_executor(None, _run_dual_pipeline, tmp_path,
31         materialNumber, materialName)

```

```
30     return {"status": "processing", "materialNumber": materialNumber}
```

The material number validation — `^[A-Za-z0-9_-]+$` — is enforced at every entry point in the application. This is SPARQL injection prevention: a material number like `MAT-001> } INSERT { <evil> }` would pass a naive check but break out of a SPARQL query. The regex ensures the material number can only contain characters that are safe to embed in an IRI.

Three things happen before the response is sent: the uploaded file is written to a temp path, `_mark_processing()` sets both status fields to 'PROCESSING' in HANA, and `_run_dual_pipeline()` is submitted to an executor — it runs *after* the response returns.

5.5.2 The dual-pipeline orchestrator

```
1  def _run_dual_pipeline(
2      tmp_path: str,
3      material_number: str,
4      material_name: str,
5  ) -> None:
6      """
7      Runs vector and KG pipelines in parallel.
8      Called in a background thread – must not raise unhandled exceptions.
9      """
10     try:
11         text, page_count = extract_pdf_text(tmp_path)
12     except ValueError as exc:
13         # PDF has no text layer – fail both pipelines immediately
14         _mark_vector_status(material_number, "ERROR", str(exc))
15         _mark_kg_status(material_number, "ERROR", str(exc))
16         os.unlink(tmp_path)
17         return
18
19     vector_error = None
20     kg_error = None
21
22     with ThreadPoolExecutor(max_workers=2) as executor:
23         future_vector = executor.submit(_vector_pipeline, text,
material_number, material_name)
```

```

24         future_kg      = executor.submit(_kg_pipeline, text,
material_number)
25
26         for future in as_completed([future_vector, future_kg]):
27             try:
28                 future.result()
29             except Exception as exc:
30                 if future is future_vector:
31                     vector_error = str(exc)
32                 else:
33                     kg_error = str(exc)
34
35             # Update statuses independently – one failure does not mask the other
36             _mark_vector_status(material_number, "ERROR" if vector_error else
"DONE", vector_error)
37             _mark_kg_status(material_number, "ERROR" if kg_error else
"DONE", kg_error)
38
39             try:
40                 os.unlink(tmp_path)
41             except OSError:
42                 pass

```

The `with ThreadPoolExecutor(max_workers=2)` block creates exactly two worker threads. `as_completed()` yields futures as they finish; we capture any exception from each without letting it propagate, so a failure in one pipeline does not cancel the other. Both pipelines are given the opportunity to complete, succeed, or fail independently.

5.5.3 Thread 1 — the vector pipeline

```

1  from agents.srv.vector_srv import embed_text, upsert_vectors
2
3  def _vector_pipeline(
4      text: str,
5      material_number: str,
6      material_name: str,
7  ) -> None:
8      conn = get_hana_connection()
9      chunks = list(chunk_text(text, material_name))

```

```

10
11     for i, chunk in enumerate(chunks):
12         embedding = embed_text(chunk)           # Vertex AI
13         text-embedding-004
14         upsert_vectors(
15             conn=conn,
16             material_number=material_number,
17             chunk_index=i,
18             chunk_text=chunk,
19             embedding=embedding,                 # list[float], length 768
20         )
21     close_thread_connection()

```

`embed_text()` calls `text-embedding-004` through the Vertex AI Python SDK. For a 30-page document with 60 chunks, this function makes 60 sequential API calls, each taking roughly 150–200 ms — approximately 10–12 seconds total.

`upsert_vectors()` executes a SQL UPSERT on the `MSDS_VECTORS` table, using `(material_number, chunk_index)` as the key. Re-uploading a document replaces its existing vectors cleanly. The `MSDS_VECTORS` table serves all document types on the platform — an invoice chunk and an MSDS chunk differ only in their content and the material number they are associated with.

Tip: For large document batches, consider batching the embedding calls. The `text-embedding-004` endpoint accepts up to 250 texts in a single request. Instead of 60 individual calls for 60 chunks, one batched call returns all 60 embeddings at once. The latency drops from ~10 seconds to ~1.5 seconds for the embed step. Batching is shown in Appendix C.

5.5.4 Thread 2 — the Knowledge Graph pipeline

```

1  from agents.srv.kg_srv import extract_triples_from_text, insert_triples
2
3  def _kg_pipeline(text: str, material_number: str) -> None:
4      conn = get_hana_connection()
5      triples = extract_triples_from_text(text, material_number)
6
7      if not triples:

```



```

8         raise ValueError(
9             f"Zero triples extracted for '{material_number}'. "
10            "Check document quality and ontology alignment."
11        )
12
13        insert_triples(conn=conn, material_number=material_number,
14                       triples=triples)
15        close_thread_connection()

```

`extract_triples_from_text()` is the Gemini 2.5 Flash extraction function built in Chapter 4. It sends the full text with a structured prompt that includes the ontology predicates and asks for JSON output. `insert_triples()` builds a SPARQL `INSERT DATA` query and calls `SPARQL_EXECUTE` on the HANA connection. The resulting triples are stored in the named graph `MSDS_Graph/MAT-XXX` under the `msds.kg` namespace.

The “zero triples” check is deliberate. A successful Gemini call that returns no triples is worse than a clean failure — the document would appear as `kgStatus = 'DONE'` to the user but produce no graph query results. Raising `ValueError` here sets `kgStatus = 'ERROR'`, which makes the problem visible in the CAP UI and queryable in the `MSDS_DOCUMENTS` table.

5.6 Status tracking

The CAP Fiori UI needs to know when ingestion is complete and what the result counts are. The UI polls `GET /status/{materialNumber}` every 3 seconds until both pipelines report a terminal state. The status data it receives comes entirely from the `MSDS_DOCUMENTS` table in HANA.

5.6.1 The `MSDS_DOCUMENTS` table and its role

`MSDS_DOCUMENTS` is the control plane of the ingestion pipeline. Every upload creates or updates a row here. The CAP OData service reads from this table to populate the document list in the Fiori UI — the `triples` count (how many KG facts were extracted), the `vectors` count (how many chunks were embedded), and the status indicators for each pipeline all come from this table.

Column	Tracks	Values
STATUS	KG pipeline	PROCESSING, DONE, ERROR
VECTOR_STATUS	Vector pipeline	PROCESSING, DONE, ERROR
KG_ERROR	KG failure reason	NULL or error message
VECTOR_ERROR	Vector failure reason	NULL or error message

Using two separate status columns rather than one combined status allows the UI to show partial progress — “Vector: DONE, KG: PROCESSING...” — and makes it possible to retry only the failed pipeline without re-running the successful one.

5.6.2 Status helper functions

```

1  def _mark_processing(material_number: str) -> None:
2      conn = get_hana_connection()
3      cursor = conn.cursor()
4      cursor.execute("""
5          UPSERT MSDS_DOCUMENTS (MATERIAL_NUMBER, STATUS, VECTOR_STATUS,
6          CREATED_AT)
7          VALUES (?, 'PROCESSING', 'PROCESSING', NOW())
8          WITH PRIMARY KEY
9          """, (material_number,))
10     conn.commit()
11     cursor.close()
12     close_thread_connection()
13
14 def _mark_vector_status(material_number: str, status: str,
15     error_message=None) -> None:
16     conn = get_hana_connection()
17     cursor = conn.cursor()
18     cursor.execute("""
19         UPDATE MSDS_DOCUMENTS
20         SET VECTOR_STATUS = ?, VECTOR_ERROR = ?, UPDATED_AT = NOW()
21         WHERE MATERIAL_NUMBER = ?
22         """, (status, error_message, material_number))
23     conn.commit()
24     cursor.close()
25     close_thread_connection()

```

```

24
25 def _mark_kg_status(material_number: str, status: str, error_message=None)
    -> None:
26     conn = get_hana_connection()
27     cursor = conn.cursor()
28     cursor.execute("""
29         UPDATE MSDS_DOCUMENTS
30         SET STATUS = ?, KG_ERROR = ?, UPDATED_AT = NOW()
31         WHERE MATERIAL_NUMBER = ?
32     """, (status, error_message, material_number))
33     conn.commit()
34     cursor.close()
35     close_thread_connection()

```

5.6.3 The status endpoint

```

1  @app.get("/status/{material_number}")
2  def get_status(material_number: str):
3      if not re.match(r'^[A-Za-z0-9_-]+$ ', material_number):
4          return JSONResponse(status_code=400, content={"error": "Invalid
              materialNumber"})
5
6      conn = get_hana_connection()
7      cursor = conn.cursor()
8      cursor.execute("""
9          SELECT STATUS, VECTOR_STATUS, KG_ERROR, VECTOR_ERROR, UPDATED_AT
10         FROM   MSDS_DOCUMENTS
11         WHERE  MATERIAL_NUMBER = ?
12     """, (material_number,))
13     row = cursor.fetchone()
14     cursor.close()
15     close_thread_connection()
16
17     if not row:
18         return JSONResponse(status_code=404, content={"error": "Material
              not found"})
19
20     kg_status, vector_status, kg_error, vector_error, updated_at = row
21     return {

```

```

22     "materialNumber": material_number,
23     "kgStatus":      kg_status,
24     "vectorStatus":  vector_status,
25     "kgError":       kg_error,
26     "vectorError":   vector_error,
27     "updatedAt":     str(updated_at),
28     "complete":      kg_status in ("DONE", "ERROR") and
29                     vector_status in ("DONE", "ERROR"),
30 }

```

The `complete` boolean is the sentinel the CAP service uses to stop polling. If `complete` is true but either status is `ERROR`, the Fiori UI shows an error banner with the error message but does not block navigation — the user can still query against whichever pipeline succeeded.

5.7 Error handling — what happens when one pipeline fails

The most important design property of this ingestion service is **error isolation**: a failure in one pipeline does not affect the other.

Failure	What happens	Status in HANA
PDF has no text layer	Both pipelines fail before threads start	Both <code>ERROR</code> , same error message
Vertex AI embedding API down	<code>_vector_pipeline</code> raises; <code>_kg_pipeline</code> continues	<code>vectorStatus = ERROR</code> , <code>kgStatus = DONE</code>
Gemini triple extraction returns empty	<code>_kg_pipeline</code> raises; <code>_vector_pipeline</code> has already completed	<code>vectorStatus = DONE</code> , <code>kgStatus = ERROR</code>
HANA connection fails in one thread	Affected thread raises; other thread uses its own connection	Independent <code>ERROR</code> / <code>DONE</code>

Scenario 2 deserves attention because it reflects a real operational condition. When a Vertex AI service outage occurs, every embedding call will fail. If error propagation

were allowed, the `ThreadPoolExecutor` would cancel the KG future when the vector future raises — and the document would have no graph knowledge despite the Gemini extraction succeeding. The `try/except` wrapper in `as_completed()` prevents that cancellation.

Persisting error messages in the `KG_ERROR` and `VECTOR_ERROR` columns means operators can run `SELECT MATERIAL_NUMBER, KG_ERROR FROM MSDS_DOCUMENTS WHERE STATUS = 'ERROR'` to get a triage list without reading log files. This is the operational advantage of storing status in HANA rather than in application logs.

5.8 Testing — upload five material documents

5.8.1 Start the service

```
1 cd agents
2 source .venv/bin/activate
3 uvicorn main:app --reload --port 8000
```

5.8.2 Upload the first document

```
1 curl -s -X POST http://localhost:8000/process-upload \
2     -F "file=@/path/to/batch-cert-001.pdf" \
3     -F "materialNumber=MAT-001" \
4     -F "materialName=BATCH-QC-MAT-001"
```

Expected response (within 200 ms):

```
1 {"status": "processing", "materialNumber": "MAT-001"}
```

5.8.3 Poll for completion

```
1 watch -n3 'curl -s http://localhost:8000/status/MAT-001 | python3 -m
  json.tool'
```

After 8–15 seconds:

```

1  {
2      "materialNumber": "MAT-001",
3      "kgStatus": "DONE",
4      "vectorStatus": "DONE",
5      "kgError": null,
6      "vectorError": null,
7      "updatedAt": "2026-05-27 09:31:17.447",
8      "complete": true
9  }

```

5.8.4 Verify the vector store

```

1  SELECT
2      CHUNK_INDEX,
3      LEFT(CHUNK_TEXT, 80) AS CHUNK_PREVIEW,
4      CARDINALITY(EMBEDDING) AS VECTOR_DIM
5  FROM MSDS_VECTORS
6  WHERE MATERIAL_NUMBER = 'MAT-001'
7  ORDER BY CHUNK_INDEX;

```

Expected: every row has VECTOR_DIM = 768.

5.8.5 Verify the Knowledge Graph

```

1  from hdbcli import dbapi
2  import os
3
4  conn = dbapi.connect(
5      address=os.environ["HANA_HOST"],
6      port=int(os.environ["HANA_PORT"]),
7      user=os.environ["HANA_USER"],
8      password=os.environ["HANA_PASSWORD"],
9  )
10 cursor = conn.cursor()
11
12 sparql = """
13 PREFIX msds: <http://msds.knowledge-graph.org/ontology/>
14 SELECT ?predicate ?object

```

```

15 WHERE {
16     GRAPH <http://msds.knowledge-graph.org/MSDS_Graph/MAT-001> {
17         <http://msds.knowledge-graph.org/material/MAT-001> ?predicate ?object
18         .
19     }
20 }
21 """
22 cursor.callproc("SPARQL_EXECUTE", [sparql, None, 1000, None, None])
23 rows = cursor.fetchall()
24 for row in rows:
25     print(row)

```

Expected output:

```

('http://msds.knowledge-graph.org/ontology#certifiedBy',      'ACME Steel AG')
('http://msds.knowledge-graph.org/ontology#certificateNumber', 'QC-CERT-
44781')
('http://msds.knowledge-graph.org/ontology#testResult',      'ISO 6892-
1 tensile test')
('http://msds.knowledge-graph.org/ontology#certifyingLab',    'Bureau Veritas Testing Gm
('http://msds.knowledge-graph.org/ontology#requiresPrecaution', 'Keep away from open fla
('http://msds.knowledge-graph.org/ontology#hasSupplier',      'Sigma-
Aldrich')

```

Warning: If rows is empty, first check the GRAPH clause. HANA SPARQL returns no results without the GRAPH wrapper, even if the triples exist. The second thing to check is the IRI prefix — the subject must match exactly what insert_triples() used when storing the graph.

Upload the remaining four PDFs (MAT-002 through MAT-005) with the same pattern. After all five complete, you will have 250–350 rows in MSDS_VECTORS and 5 named graphs in the HANA RDF store.

5.9 Summary

This chapter built the pipeline that bridges PDF uploads and the dual-store knowledge layer. The architecture applies to any document type uploaded against an SAP Material

Number. The key design decisions:

- **Parallel threads** — `ThreadPoolExecutor(max_workers=2)` runs the KG extraction (Gemini 2.5 Flash) and the vector embedding (`text-embedding-004`) simultaneously. These operations are independent: neither pipeline waits for the other. Wall-clock ingestion time equals the slower pipeline, not the sum.
- **Fire-and-forget with status polling** — the upload endpoint returns 202 in milliseconds. The CAP Fiori UI polls every 3 seconds. Decouples ingestion time from HTTP timeouts — essential for large documents on BTP CF.
- **Thread-local HANA connections** — `threading.local()` gives each worker thread its own connection object, eliminating race conditions without connection pooling infrastructure.
- **Different text strategies per pipeline** — the vector pipeline chunks to 500 tokens for embedding precision; the KG pipeline sends the full document for triple extraction quality.
- **Document-type agnostic vector pipeline** — the same chunker and embed logic handles any PDF. Only the KG extraction prompt changes per document type.
- **Error isolation** — exceptions in one pipeline are caught and persisted as `ERROR` status without interrupting the other pipeline.
- **MSDS_DOCUMENTS as control plane** — status, error messages, and counts are persisted in HANA. The CAP service reads this table directly. Operators can triage failures with SQL, not log files.

5.10 Checkpoint

Before moving to Chapter 6, confirm each of the following:

- ☐ `agents/srv/doc_srv.py` exists and `uvicorn main:app --port 8000` starts without import errors.
- ☐ `POST /process-upload` with a real PDF returns `{"status": "processing"}` within 500 ms.
- ☐ `GET /status/{materialNumber}` returns `"complete": true` within 30 seconds for a 20–40 page document.
- ☐ Five documents uploaded; all show `kgStatus = "DONE"` and `vectorStatus = "DONE"`.

- ☐ `SELECT COUNT(*) FROM MSDS_VECTORS WHERE MATERIAL_NUMBER = 'MAT-001'` returns a positive number.
- ☐ `SELECT CARDINALITY(EMBEDDING) FROM MSDS_VECTORS FETCH FIRST 1 ROWS ONLY` returns 768.
- ☐ SPARQL query against `MSDS_Graph/BATCH-QC-MAT-001` returns at least one triple with predicate `certifiedBy`.
- ☐ Re-uploading the same PDF overwrites existing rows rather than creating duplicates.
- ☐ Uploading a scan-only PDF results in both statuses showing `ERROR`, not `PROCESSING`.

With the ingestion pipeline verified and five documents loaded into both stores, you are ready for the agentic layer. Chapter 6 introduces LangGraph — the framework that will orchestrate queries against these two stores in parallel and synthesise a single coherent answer.

End of Chapter 5

Chapter 6: LangGraph Fundamentals

In Chapter 5 we built a document ingestion pipeline that processes PDFs into both a vector store and a Knowledge Graph simultaneously. We now have data — rich, structured, searchable data — in SAP HANA Cloud. The next question is: how do we build an agent that reasons over it in a way that is auditable, inspectable, and safe to deploy on BTP?

A plain LangChain chain would work for a single retrieval followed by a single LLM call. But the system we need is more complex than that. It needs to run two retrieval strategies in parallel, decide how to merge their results, handle retries when SPARQL returns empty results, maintain conversation history across turns, and expose all of this over a single HTTP endpoint. A linear chain cannot express those decisions. We need something that can branch, loop, and conditionally route based on what it finds.

That something is LangGraph.

This chapter introduces LangGraph from first principles. We build a simple question-answering agent step by step — one node at a time — so that when Chapter 7 extends it into the full parallel hybrid RAG system, every piece is familiar. By the end of this chapter you will have a running LangGraph agent connected to Vertex AI, and you will understand exactly how state flows through it.

6.1 Why LangGraph is SAP's recommended orchestration approach for BTP agents

Enterprise AI governance requires more than a correct answer. It requires an auditable trail: which system provided which data, which decision was made at which step, and why. When an auditor asks “how did the system determine that the tensile test used ISO 6892-1?”, you need to be able to show the SPARQL query, the HANA result, and the LLM call that synthesised it into natural language. A black-box chain gives you none of that. A LangGraph StateGraph gives you all of it.

LangGraph models your agent as a directed graph rather than a list. Every node in the graph is a named Python function. Every state transition is an explicit edge. Every

intermediate result is captured in the graph state. LangGraph traces can be exported to LangSmith, where each run shows a complete tree: which nodes executed, in what order, what the state contained at each step, and the exact prompts sent to Gemini. This execution trace is the enterprise AI governance artefact.

The conceptual mapping to SAP BTP workflow constructs is direct:

LangGraph concept	BTP workflow equivalent
<code>AgentState (TypedDict)</code>	BTP Workflow context object
<code>Node (Python function)</code>	Workflow service task step
<code>Edge</code>	Sequence flow between steps
<code>Conditional edge</code>	Exclusive gateway (XOR decision)
<code>StateGraph.compile()</code>	Deployed workflow definition
<code>app.invoke(state)</code>	Workflow instance execution

SAP's own internal AI agent frameworks use LangGraph for exactly this reason: the execution model is transparent, the state is inspectable at any point, and the graph definition is a first-class artefact that can be reviewed, versioned, and audited like any other software component.

6.2 Why LangChain chains are not enough

LangChain introduced the concept of a chain: a sequence of components where the output of one feeds the input of the next. For many tasks — retrieve a passage, summarise it, return the result — a chain is exactly right. It is simple, readable, and easy to test.

The problem arises when your logic is not a sequence. Consider what our hybrid RAG system needs to do:

Requirement	Can a chain do it?
Run two retrieval strategies in parallel	No — chains are sequential
Retry SPARQL generation if results are empty	No — chains have no loops

Requirement	Can a chain do it?
Route to different merge strategies based on what was found	No — chains have no conditional branches
Pass conversation history into every step	Awkward — requires passing state manually
Produce an auditable execution trace	No — chain execution is opaque

LangGraph solves every one of these by modelling your agent as a directed graph rather than a list. Nodes are Python functions. Edges are connections between them. Conditional edges let the graph branch based on the current state. The graph can loop. It can run nodes in parallel. It can capture and stream intermediate results.

Note: LangGraph is built on top of LangChain. You still use LangChain components — `ChatVertexAI`, tool decorators, message types — but LangGraph provides the execution engine that orchestrates them.

6.3 Core concepts

Before writing any code, here are the four concepts you need to hold in your head.

6.3.1 State

State is the memory of your agent for a single run. It is a Python `TypedDict` — a dictionary where every key has a declared type. Every node in the graph reads from state and writes back to state. State is the only way nodes communicate with each other.

In BTP workflow terms, state is the context object that a workflow instance carries from step to step. When you inspect a running workflow in BTP Process Automation, you see this context. In LangGraph, when you inspect a run in LangSmith, you see the state at each node boundary. The concept is identical.

```

1 from typing import TypedDict, List
2
3 class AgentState(TypedDict):
4     question: str

```

```
5     answer: str
6     messages: List[dict]
```

When you invoke the graph, you pass an initial state. When the graph finishes, you receive the final state. Everything that happened in between is recorded there.

6.3.2 Nodes

A node is a Python function that receives the current state and returns a partial update to it. LangGraph merges the update into the state before passing it to the next node. In BTP workflow terms, a node is a service task — it does one focused piece of work and writes its result to the workflow context.

```
1 def my_node(state: AgentState) -> dict:
2     question = state["question"]
3     answer = call_llm(question)
4     # return only the keys you want to update
5     return {"answer": answer}
```

The function does not need to return the entire state — only the keys it wants to change. LangGraph merges the returned dict into the existing state automatically.

6.3.3 Edges

An edge connects two nodes. When node A finishes, LangGraph follows the edge to node B and runs it next. This is a sequence flow in BTP workflow terms — deterministic, always-true routing.

```
1 graph.add_edge("node_a", "node_b")
```

6.3.4 Conditional edges

A conditional edge calls a routing function after a node completes. The routing function inspects the state and returns a string that names the next node to run. In BTP workflow terms, this is an exclusive gateway: examine the context, choose one outgoing sequence flow.

```
1 def route(state: AgentState) -> str:
2     if state["answer"] == "":
```

```

3         return "retry"
4     return "done"
5
6 graph.add_conditional_edges("node_a", route, {
7     "retry": "retry_node",
8     "done": END
9 })

```

This is how loops and branches are expressed in LangGraph.

6.4 Installation and project layout

Install LangGraph and the Google GenAI integration. Note the correct import: `langchain_google_genai`, not `langchain_google_vertexai` — the VertexAI variant is deprecated for the Gemini model family.

```

1 pip install langgraph langchain-google-genai langchain-core

```

Add these to `agents/requirements.txt`:

```

langgraph>=0.2.0
langchain-google-genai>=2.0.0
langchain-core>=0.3.0

```

We will build the simple agent in a new file:

```

agents/
  agents/
    simple_qa_agent.py    ← this chapter
    orchestrator.py      ← Chapter 7
    kg_chain.py           ← Chapter 7
    vector_chain.py       ← Chapter 7

```

6.5 Building the state and nodes

Open `agents/agents/simple_qa_agent.py`. We start with the state definition and two nodes.

```

1  from typing import TypedDict, List
2  from langgraph.graph import StateGraph, END
3  from langchain_google_genai import ChatGoogleGenerativeAI
4  from langchain_core.messages import HumanMessage, AIMessage
5  # — State
6
7  class AgentState(TypedDict):
8      question: str
9      context: str
10     answer: str
11     messages: List[dict]
12     retry_count: int
13
14 # — LLM
15
16 def _get_llm():
17     return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
18                                     temperature=0.1, max_tokens=1024)
19
20 # — Nodes
21
22 def retrieve_node(state: AgentState) -> dict:
23     """Stub retriever — replaced by vector + KG chains in Chapter 7."""
24     question = state["question"]
25     context = f"[Retrieved context for: {question}]"
26     return {"context": context}
27
28 def answer_node(state: AgentState) -> dict:
29     """Generate an answer using the retrieved context."""
30     question = state["question"]
31     context = state["context"]
32     history = state.get("messages", [])
33
34     messages = []
35     for msg in history:
36         if msg["role"] == "user":

```

```

31         messages.append(HumanMessage(content=msg["content"]))
32     else:
33         messages.append(AIMessage(content=msg["content"]))
34
35     prompt = f"""You are an expert assistant for SAP material documents –
    you help users find information across PDF documents linked to SAP
    Material Numbers.
36
37     Context from the knowledge base:
38     {context}
39
40     Answer the following question based on the context above.
41     If the context does not contain enough information, say so clearly.
42
43     Question: {question}"""
44
45     messages.append(HumanMessage(content=prompt))
46     response = _get_llm().invoke(messages)
47     return {"answer": response.content}

```

Two things to notice. First, `retrieve_node` is a deliberate stub — it returns a placeholder string. We build the real retrieval in Chapter 7; this stub lets us test the graph structure now without needing a live HANA connection. Second, `answer_node` builds the conversation history from the `messages` list in state. This is how multi-turn conversations work: the caller passes previous turns in the initial state, and the node includes them in the prompt.

LLM instantiation uses a lazy getter `_get_llm()` rather than a module-level variable. This prevents the LLM client from being initialised at import time, which would fail in environments where credentials are not yet available — a common issue in BTP CF where environment variables are injected after application startup.

6.6 Wiring the graph

Add the graph construction below the node functions:


```

1 def build_graph():
2     graph = StateGraph(AgentState)
3
4     graph.add_node("retrieve", retrieve_node)
5     graph.add_node("answer", answer_node)
6
7     graph.set_entry_point("retrieve")
8     graph.add_edge("retrieve", "answer")
9     graph.add_edge("answer", END)
10
11     return graph.compile()
12
13 app = build_graph()

```

The graph is linear: retrieve → answer → END. To invoke it:

```

1 if __name__ == "__main__":
2     result = app.invoke({
3         "question": "What test method was used for the tensile strength
4         test?",
5         "context": "",
6         "answer": "",
7         "messages": [],
8         "retry_count": 0,
9     })
10     print(result["answer"])

```

Run it:

```

1 cd agents
2 python -m agents.simple_qa_agent

```

Expected output (with stub retriever):

Based on the provided context, I don't have specific test method details for the tensile strength test. The context contains a placeholder rather than actual document data. Once the knowledge base is populated with real documents linked to the material number, I will be able to retrieve and report the specific information.

The agent runs, reaches Gemini 2.5 Flash, and gives an honest answer about the stub

context. That is the graph working correctly.

6.7 Adding a conditional retry edge

Real retrieval can fail — SPARQL returns empty results, an embedding call times out. We need to be able to retry. Add a check node and a routing function:

```

1  def check_answer_node(state: AgentState) -> dict:
2      """Check if the answer is acceptable. Increment retry counter if
3      not."""
4      answer      = state.get("answer", "")
5      retry_count = state.get("retry_count", 0)
6
7      low_quality_phrases = [
8          "don't have", "no information", "cannot find",
9          "not available", "placeholder"
10     ]
11     needs_retry = any(p in answer.lower() for p in low_quality_phrases)
12     return {"retry_count": retry_count + (1 if needs_retry else 0)}
13
14 def route_after_check(state: AgentState) -> str:
15     """Decide whether to retry or finish."""
16     answer      = state.get("answer", "")
17     retry_count = state.get("retry_count", 0)
18
19     low_quality_phrases = [
20         "don't have", "no information", "cannot find",
21         "not available", "placeholder"
22     ]
23     needs_retry = any(p in answer.lower() for p in low_quality_phrases)
24
25     if needs_retry and retry_count < 2:
26         return "retry"
27     return "done"

```

This is a conditional edge acting as a decision gateway. The routing function reads two state fields — `answer` and `retry_count` — and returns one of two named outcomes. LangGraph maps those outcomes to the next node using the dictionary passed

to `add_conditional_edges`. This is identical to how an exclusive gateway in BTP Process Automation evaluates a context condition and routes to one of several outgoing flows.

Update `build_graph` to include the check node and conditional edge:

```

1  def build_graph():
2      graph = StateGraph(AgentState)
3
4      graph.add_node("retrieve", retrieve_node)
5      graph.add_node("answer", answer_node)
6      graph.add_node("check", check_answer_node)
7
8      graph.set_entry_point("retrieve")
9      graph.add_edge("retrieve", "answer")
10     graph.add_edge("answer", "check")
11
12     graph.add_conditional_edges(
13         "check",
14         route_after_check,
15         {
16             "retry": "retrieve",    # loop back
17             "done": END
18         }
19     )
20
21     return graph.compile()

```


The graph now has a loop: `retrieve` → `answer` → `check` → (retry: back to `retrieve`) or (done: `END`). The `retry_count` guard prevents infinite loops — after two retries we accept whatever answer we have.

Note: In Chapter 7 the retry logic is more specific: we only retry when SPARQL returns zero results, and the retry uses a simpler SPARQL query, not a full re-run of both chains. The pattern here is the same; the condition is different.

6.8 The StateGraph visualised

The graph we built in this chapter has three nodes and one conditional loop:

LangGraph StateGraph — Chapter 7 Simple Q&A Agent



State flows through every node. Nodes return only the keys they update.

Figure 6.1: The simple Q&A agent’s StateGraph — retrieve feeds answer, which feeds the check node. The check node either loops back to retrieve (on low-quality answers, up to 2 retries) or terminates at END. The dashed box on the right shows the AgentState structure that flows through every node.

6.9 The complete simple_qa_agent.py

```

1 from typing import TypedDict, List
2 from langgraph.graph import StateGraph, END

```

```

3 from langchain_google_genai import ChatGoogleGenerativeAI
4 from langchain_core.messages import HumanMessage, AIMessage
5
6 class AgentState(TypedDict):
7     question: str
8     context: str
9     answer: str
10    messages: List[dict]
11    retry_count: int
12
13 def _get_llm():
14     return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
15                                     temperature=0.1, max_tokens=1024)
16
17 def retrieve_node(state: AgentState) -> dict:
18     question = state["question"]
19     return {"context": f"[Retrieved context for: {question}]"}
20
21 def answer_node(state: AgentState) -> dict:
22     messages = []
23     for msg in state.get("messages", []):
24         cls = HumanMessage if msg["role"] == "user" else AIMessage
25         messages.append(cls(content=msg["content"]))
26
27     prompt = f"""You are an expert assistant for SAP material documents –
28     you help users find information across PDF documents linked to SAP
29     Material Numbers.
30
31     Context:
32     {state['context']}
33
34     Question: {state['question']}
35     Answer: ""
36
37     messages.append(HumanMessage(content=prompt))
38     return {"answer": _get_llm().invoke(messages).content}
39
40 def check_answer_node(state: AgentState) -> dict:
41     low = ["don't have", "no information", "cannot find", "not available",
42           "placeholder"]

```

```

38     needs_retry = any(p in state.get("answer","").lower() for p in low)
39     return {"retry_count": state.get("retry_count", 0) + (1 if needs_retry
40         else 0)}
41
42 def route_after_check(state: AgentState) -> str:
43     low = ["don't have", "no information", "cannot find", "not available",
44         "placeholder"]
45
46     needs_retry = any(p in state.get("answer","").lower() for p in low)
47     return "retry" if (needs_retry and state.get("retry_count",0) < 2)
48     else "done"
49
50 def build_graph():
51     graph = StateGraph(AgentState)
52     graph.add_node("retrieve", retrieve_node)
53     graph.add_node("answer", answer_node)
54     graph.add_node("check", check_answer_node)
55     graph.set_entry_point("retrieve")
56     graph.add_edge("retrieve", "answer")
57     graph.add_edge("answer", "check")
58     graph.add_conditional_edges("check", route_after_check,
59         {"retry": "retrieve", "done": END})
60
61     return graph.compile()
62
63 app = build_graph()
64
65 if __name__ == "__main__":
66     result = app.invoke({
67         "question": "What test method was used for the tensile strength
68             test?",
69         "context": "",
70         "answer": "",
71         "messages": [],
72         "retry_count": 0,
73     })
74     print(result["answer"])

```

6.10 Tool use in LangGraph

LangGraph supports tools — Python functions that the LLM can choose to call based on the conversation. In a tool-enabled graph, the LLM inspects its available tools, decides which to call, and the selected tool's output is added back into the state for the next node. For the hybrid RAG agent in this book, we deliberately do not use this mechanism — we use direct parallel dispatch of both chains on every query regardless of what the LLM might prefer. The reason is covered in the note below, but understanding tool-enabled graphs is useful if you build agent extensions on top of this platform later.

Define a tool with the `@tool` decorator:

```
1 from langchain_core.tools import tool
2
3 @tool
4 def get_test_results(material_number: str) -> str:
5     """Retrieve test results for a material from the Knowledge Graph."""
6     # In a real implementation this queries HANA SPARQL
7     return "ISO 6892-1, yield strength 450 MPa, elongation 22%"
```

Bind tools to the LLM and use `ToolNode` for execution:

```
1 from langgraph.prebuilt import ToolNode
2
3 tools = [get_test_results]
4 llm_with_tools = _get_llm().bind_tools(tools)
5 tool_node = ToolNode(tools)
6
7 graph.add_node("tools", tool_node)
8 graph.add_conditional_edges(
9     "answer",
10    lambda state: "tools" if state["messages"][-1].tool_calls else END,
11    {"tools": "tools", END: END}
12 )
13 graph.add_edge("tools", "answer")
```

This creates a tool-calling loop: the LLM inspects the state, decides whether a tool call is needed, and if so, the `ToolNode` executes the selected function and writes the result back to state.

Note: Chapter 7 intentionally avoids this ReAct pattern for the main retrieval. We run both vector and KG chains on every query, regardless of what the LLM might choose. This gives deterministic latency and prevents the LLM from skipping a retrieval path it thinks is unnecessary — a critical property for enterprise systems where consistent, complete answers matter more than adaptive efficiency.

6.11 Streaming responses

LangGraph can stream state updates as they happen. This is valuable when queries take several seconds — the user sees progress rather than a blank screen.

```

1  for event in app.stream({
2      "question":    "What are the storage requirements for the material in
                        this batch certificate?",
3      "context":     "",
4      "answer":      "",
5      "messages":    [],
6      "retry_count": 0,
7  }):
8      for node_name, state_update in event.items():
9          print(f"[{node_name}] completed")
10         if "answer" in state_update:
11             print(f"  answer: {state_update['answer'][:80]}...")

```

Expected output:

[retrieve] completed

[answer] completed

answer: Store in dry conditions below 25°C away from moisture...

[check] completed

Tip: In the FastAPI service, pipe the stream to a Server-Sent Events (SSE) response using `StreamingResponse`. The Fiori UI can consume the stream and update the chat interface progressively as each node completes. We implement this in Chapter 9.

6.12 Debugging with LangSmith

LangSmith is LangChain's observability platform. It records every LangGraph run — inputs, outputs, intermediate states, LLM calls, latency, and token counts. For an enterprise AI governance use case, LangSmith is the audit log.

Set three environment variables:

```
1 export LANGCHAIN_TRACING_V2=true
2 export LANGCHAIN_ENDPOINT=https://api.smith.langchain.com
3 export LANGCHAIN_API_KEY=your-api-key-here
```

In your BTP CF `manifest.yml`, add the same variables:

```
1 applications:
2   - name: hybrid-rag-agent
3     env:
4       LANGCHAIN_TRACING_V2: "true"
5       LANGCHAIN_ENDPOINT: "https://api.smith.langchain.com"
6       LANGCHAIN_API_KEY: ((langchain-api-key))
```

Use a BTP User-Provided Service to supply the API key without storing it in the manifest. In the LangSmith dashboard you will see a tree view of each run: which nodes executed, in what order, how long each took, and the exact prompt sent to Gemini. When Chapter 7's parallel chains run, you will see two branches executing simultaneously — a visual confirmation that the `ThreadPoolExecutor` parallelism is working.

Tip: Tag your runs in `app.invoke(config={"metadata": {"material_number": material_number}})` and filter by `metadata.material_number` in LangSmith to see all queries against a specific SAP Material Number. This is the query-level audit trail that enterprise governance teams require.

6.13 Why stateless works on BTP Cloud Foundry

A common concern when deploying LangGraph to Cloud Foundry is: what happens to agent state when you scale to multiple instances?

The answer is: nothing, because there is no server-side state to worry about.

LangGraph state is in-memory for a single request. When `app.invoke()` returns, the state is gone. There is no database, no session file, no sticky session required.

Concern	Reality
CF runs 10 instances	Each request lands on any instance — fine
CF restarts an instance	No state is lost — state only exists during a request
Two users query simultaneously	Each gets their own independent in-memory state
Conversation history	Passed in <code>messages</code> field of every request body by the caller

The conversation history pattern — where the client sends previous turns in every request — is the key. The CAP Fiori frontend stores the last N messages and sends them with each new question.

```

1  # Frontend sends history on every request
2  # POST /query
3  {
4      "question": "And what about the certified testing laboratory?",
5      "material_number": "BATCH-QC-MAT-001",
6      "history": [
7          {"role": "user", "content": "What test method was used for the
8          tensile strength test?"},
9          {"role": "assistant", "content": "The tensile test used ISO 6892-1
10         methodology with 450 MPa yield strength."}
11     ]
12 }
```

The agent picks up the conversation exactly where it left off — without any server-side memory. This design is safe to scale horizontally on BTP CF without load balancer sticky sessions or shared caches.

6.14 Summary

In this chapter we built a LangGraph agent from scratch and grounded every concept in SAP BTP enterprise terms:

- Defined **state** as a `TypedDict` that flows through every node — the workflow context object
- Wrote **nodes** as plain Python functions that read state and return partial updates — the workflow service tasks
- Connected nodes with **edges** and **conditional edges** to express branching and looping — sequence flows and exclusive gateways
- Added a **retry loop** that re-runs retrieval when the answer quality is low
- Demonstrated **tool use** with the `@tool` decorator and `ToolNode`
- Enabled **streaming** to show node-level progress events
- Configured **LangSmith** for end-to-end observability and audit trail
- Explained why **stateless design** makes LangGraph safe to scale horizontally on BTP CF

The agent built here uses a stub retriever. In Chapter 7, we replace that stub with the two HANA retrieval chains — vector cosine search and SPARQL — and run them in parallel.

6.15 Checkpoint

Before continuing to Chapter 7, verify the following:

```
1 # 1. LangGraph is installed
2 python -c "import langgraph; print(langgraph.__version__)"
3
4 # 2. The simple agent runs without errors
5 cd agents && python -m agents.simple_qa_agent
6
7 # 3. Gemini credentials are set
8 echo $GOOGLE_API_KEY # or $GOOGLE_APPLICATION_CREDENTIALS for service
  account
9
10 # 4. LangSmith (optional but recommended for governance)
```

```
11 echo $LANGCHAIN_TRACING_V2          # should print "true"
```

If all four commands succeed, you have a working LangGraph agent connected to Gemini 2.5 Flash. Chapter 7 takes this foundation and builds the full parallel hybrid RAG system on top of it — replacing the stub retriever with live HANA vector search and SPARQL execution running in parallel.

End of Chapter 6

Chapter 7: The Parallel Hybrid RAG Agent

SAP HANA Cloud is the only enterprise database that provides both `REAL_VECTOR` cosine similarity search and native SPARQL execution in a single managed service on BTP. This chapter shows how to use both simultaneously.

In Chapter 6 we built a LangGraph agent with a stub retriever — a placeholder that returned a fixed string instead of real data. In Chapters 3 and 4 we built the real retrieval systems: a vector store in SAP HANA Cloud and a Knowledge Graph queried via SPARQL. In Chapter 5 we built the ingestion pipeline that populates both. All the pieces exist. This chapter assembles them.

The central design question is: when a user asks a question, do we run vector search first and Knowledge Graph search second? Or do we ask the LLM to decide which one to use? The answer to both is no. We run them in parallel, always, regardless of the question. This chapter explains why that decision is correct and shows you exactly how to implement it.

By the end of this chapter you will have a working hybrid RAG system: a FastAPI service that receives a natural language question about any PDF document linked to an SAP Material Number, dispatches it simultaneously to both retrieval chains, merges the results, and returns an answer that demonstrably outperforms either strategy alone.

7.1 Why parallel beats sequential

Consider the naive approach: run vector search, check the result, then decide whether to also run KG search. This is sequential routing, and it has a hidden cost.

Approach	Wall-clock time	LLM calls	Risk
Sequential: vector then KG	$2\text{ s} + 2\text{ s} = 4\text{ s}$	3 (route + vector + KG)	Routing LLM makes wrong choice
Sequential: KG then vector	$2\text{ s} + 2\text{ s} = 4\text{ s}$	3	Same risk

Approach	Wall-clock time	LLM calls	Risk
Routing: LLM decides which to use	1 s + 2 s = 3 s	2	Routing LLM skips useful path
Parallel: both simultaneously	max(2 s, 2 s) = 2 s	2	None

The parallel approach is faster than all sequential variants and uses fewer LLM calls than the routing variant. More importantly, it eliminates the routing error entirely. A routing LLM might decide that “what test method was used for the tensile strength test?” is a structured query and skip the vector chain — missing the test procedure details that live only in prose. By always running both, we guarantee that no retrieval path is skipped. For enterprise systems where completeness matters — a quality engineer needs all relevant information, not just what the system guesses is most relevant — this deterministic behaviour is non-negotiable.




Figure 7.1: The parallel hybrid RAG orchestrator — both chains run concurrently via `ThreadPoolExecutor(max_workers=2)`. Wall-clock latency equals the slower of the two chains, not their sum.

Note: This design assumes both chains have roughly similar latency (~1–3 seconds each). If one chain were orders of magnitude slower than the other,

you might reconsider. In practice, a HANA SPARQL query and a HANA vector search take comparable time.

7.2 The agent state

Create `agents/agents/state.py`:

```
1 from typing import TypedDict, List, Optional
2
3 class HybridRAGState(TypedDict):
4     # Input
5     question: str
6     material_number: str
7     history: List[dict]
8
9     # Vector chain outputs
10    vector_answer: str
11    vector_chunks: List[dict]
12
13    # KG chain outputs
14    kg_answer: str
15    kg_sparql: str
16    kg_facts: List[dict]
17
18    # Final output
19    final_answer: str
20    sources: List[str]
21
22    # Control
23    error: Optional[str]
```

This state is the contract between every component in the system. The vector chain writes to `vector_answer` and `vector_chunks`. The KG chain writes to `kg_answer`, `kg_sparql`, and `kg_facts`. The merge function reads all four and writes `final_answer`.

The `sources` field serves a specific enterprise purpose: auditability. When the CAP Fiori UI displays an answer, it also shows the user which retrieval paths contributed — “Knowledge graph: 3 facts, Document search: 5 passages”. This tells the user whether

the answer is grounded in structured facts from the KG, in narrative passages from the vector store, or in both. For regulated industries where answer provenance must be demonstrable, this field is not optional.

7.3 The vector chain

Create `agents/agents/vector_chain.py`:

```

1  import logging
2  from typing import Optional
3  from langchain_google_genai import ChatGoogleGenerativeAI,
   GoogleGenerativeAIEmbeddings
4  from langchain_core.messages import HumanMessage
5
6  from agents.state import HybridRAGState
7  from srv.hdb_srv import get_connection
8
9  logger = logging.getLogger(__name__)
10
11 def _get_llm():
12     return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
13                                     temperature=0.1, max_tokens=1024)
14
15 def _get_embedder():
16     return GoogleGenerativeAIEmbeddings(model="models/text-embedding-004")
17
18 COSINE_SEARCH_SQL = """
19     SELECT TOP 5
20         CHUNK_TEXT,
21         MATERIAL_NUMBER,
22         CHUNK_INDEX,
23         TO_DOUBLE(COSINE_SIMILARITY(EMBEDDING, TO_REAL_VECTOR(?))) AS SCORE
24     FROM MSDS_VECTORS
25     WHERE MATERIAL_NUMBER = ?
26     ORDER BY SCORE DESC
27     """

```



```

28 def run_vector_chain(state: HybridRAGState) -> dict:
29     """Execute the vector retrieval chain and return state updates."""
30     question = state["question"]
31     material_number = state["material_number"]
32
33     try:
34         # Step 1: Embed the question
35         embedding = _get_embedder().embed_query(question)
36         embedding_str = "[" + ",".join(str(v) for v in embedding) + "]"
37
38         # Step 2: Cosine search in HANA
39         conn = get_connection()
40         cursor = conn.cursor()
41         cursor.execute(COSINE_SEARCH_SQL, (embedding_str,
material_number))
42         rows = cursor.fetchall()
43
44         if not rows:
45             logger.info("Vector search returned no results for %s",
material_number)
46             return {
47                 "vector_answer": "",
48                 "vector_chunks": [],
49             }
50
51         chunks = [
52             {
53                 "text": row[0],
54                 "material_number": row[1],
55                 "chunk_index": row[2],
56                 "score": float(row[3]),
57             }
58             for row in rows
59         ]
60
61         # Step 3: Summarise with Gemini 2.5 Flash
62         context = "\n\n--\n\n".join(c["text"] for c in chunks)

```

```

63     prompt = f"""You are an expert assistant for SAP material
documents – answer questions based only on the provided document context.
Be specific and concise.
64 If the passages do not contain enough information, say so.
65
66 Passages:
67 {context}
68
69 Question: {question}
70
71 Answer: """
72
73     response = _get_llm().invoke([HumanMessage(content=prompt)])
74     return {
75         "vector_answer": response.content,
76         "vector_chunks": chunks,
77     }
78
79 except Exception as e:
80     logger.error("Vector chain failed: %s", e)
81     return {
82         "vector_answer": "",
83         "vector_chunks": [],
84         "error": f"Vector chain error: {str(e)}",
85     }

```

Three steps: embed the question into a 768-dimensional vector using text-embedding-004, run a cosine similarity search against MSDS_VECTORS, summarise the top-5 matching chunks with Gemini 2.5 Flash. The SQL uses HANA's COSINE_SIMILARITY function and TO_REAL_VECTOR to cast the embedding string to the correct column type.

7.4 The KG chain

Create agents/agents/kg_chain.py:

```

1  import logging
2  from langchain_google_genai import ChatGoogleGenerativeAI

```

```

3  from langchain_core.messages import HumanMessage
4
5  from agents.state import HybridRAGState
6  from srv.hdb_srv import get_connection
7
8  logger = logging.getLogger(__name__)
9
10 def _get_llm():
11     return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
12                                     temperature=0.0, max_tokens=512)
13
14 def _get_summariser():
15     return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
16                                     temperature=0.1, max_tokens=1024)
17
18 GRAPH_URI_PREFIX = "http://msds.knowledge-graph.org/MSDS_Graph/"
19
20 SPARQL_GEN_PROMPT = """You are a SPARQL expert. Generate a SPARQL SELECT
21 query to answer
22 the question below using the provided ontology.
23
24 Ontology predicates available:
25 - msds:certifiedBy (links Batch to certifying organisation node)
26 - msds:testResult (links Batch to TestResult, has msds:testMethod and
27 msds:resultValue)
28 - msds:certificateNumber (literal on Batch – e.g. QC-CERT-44781)
29 - msds:certifyingLab (links Batch to laboratory node)
30 - msds:description (literal description on any node)
31 - rdfs:label (name of the material or batch)
32
33 Named graph: <{graph_uri}>
34 Material URI: <http://msds.knowledge-graph.org/material/{material_number}>
35
36 Rules:
37 1. Always wrap triple patterns with GRAPH <{graph_uri}> {{ ... }}
38 2. Use PREFIX msds: <http://msds.knowledge-graph.org/ontology/>
39 3. Use PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
40 4. Return only the SPARQL query, no explanation

```

```

38 Question: {question}
39
40 SPARQL: ""
41
42 SPARQL_FALLBACK_PROMPT = """The previous SPARQL query returned no results.
Generate a simpler,
43 broader SPARQL query to retrieve any available information about this
material.
44
45 Named graph: <{graph_uri}>
46 Material URI: <http://msds.knowledge-graph.org/material/{material_number}>
47
48 Return all triples about this material:
49
50 SPARQL: ""
51
52 def _execute_sparql(sparql: str) -> list:
53     conn = get_connection()
54     cursor = conn.cursor()
55     try:
56         cursor.callproc("SPARQL_EXECUTE", [sparql, None, 1000, None, None])
57         return cursor.fetchall()
58     finally:
59         cursor.close()
60
61 def run_kg_chain(state: HybridRAGState) -> dict:
62     """Execute the KG retrieval chain and return state updates."""
63     question = state["question"]
64     material_number = state["material_number"]
65     graph_uri = f"{GRAPH_URI_PREFIX}{material_number}"
66
67     try:
68         # Step 1: Generate SPARQL
69         gen_prompt = SPARQL_GEN_PROMPT.format(
70             graph_uri=graph_uri,
71             material_number=material_number,
72             question=question,
73         )

```

```

74     sparql_response =
_get_llm().invoke([HumanMessage(content=gen_prompt)])
75     sparql = sparql_response.content.strip()
76     # Strip markdown code fences if present
77     if sparql.startswith("```"):
78         sparql = "\n".join(sparql.split("\n")[1:-1])
79
80     # Step 2: Execute SPARQL
81     rows = _execute_sparql(sparql)
82
83     # Step 3: Retry with fallback if empty
84     if not rows:
85         logger.info("SPARQL returned empty, retrying with fallback for
%s", material_number)
86         fallback_prompt = SPARQL_FALLBACK_PROMPT.format(
87             graph_uri=graph_uri,
88             material_number=material_number,
89         )
90         fallback_response =
_get_llm().invoke([HumanMessage(content=fallback_prompt)])
91         sparql = fallback_response.content.strip()
92         if sparql.startswith("```"):
93             sparql = "\n".join(sparql.split("\n")[1:-1])
94         rows = _execute_sparql(sparql)
95
96     if not rows:
97         return {
98             "kg_answer": "",
99             "kg_sparql": sparql,
100             "kg_facts": [],
101         }
102
103     # Step 4: Summarise with Gemini 2.5 Flash
104     facts = [str(row) for row in rows]
105     facts_text = "\n".join(facts[:50]) # limit to 50 facts
106

```

```

107     summarise_prompt = f"""You are an expert assistant for SAP
material documents — answer questions based only on the provided document
context. The following facts were retrieved from a structured Knowledge
Graph about material {material_number}.
108 Use them to answer the question precisely.
109
110 Facts:
111 {facts_text}
112
113 Question: {question}
114
115 Answer: """
116
117     summary_response =
_get_summariser().invoke([HumanMessage(content=summarise_prompt)])
118     return {
119         "kg_answer": summary_response.content,
120         "kg_sparql": sparql,
121         "kg_facts": facts[:50],
122     }
123
124     except Exception as e:
125         logger.error("KG chain failed: %s", e)
126         return {
127             "kg_answer": "",
128             "kg_sparql": "",
129             "kg_facts": [],
130             "error": f"KG chain error: {str(e)}",
131         }

```

The KG chain has one critical feature beyond basic SPARQL execution: the retry-on-empty loop. When Gemini generates a SPARQL query that returns no rows — which happens when the generated query is overly specific — the chain retries with a simpler fallback query that retrieves all triples for the material. This ensures we always get *something* from the Knowledge Graph if data exists.

Warning: Never share HANA connections across threads. The `get_connection()` call in both chains uses `threading.local()` under the hood — each thread gets its own connection. If you pass a connection object between threads, HANA will raise concurrency errors. See `agents/srv/hdb_srv.py` for the

thread-local implementation.

7.5 The merge function

The merge function is the heart of hybrid RAG. It receives whatever the two chains produced and decides how to combine them. The four cases map directly to enterprise answer quality tiers:

Both chains return results — comprehensive answer. The KG provides structured facts (exact codes, measured values, regulatory identifiers); the vector store provides narrative context (procedures, explanations, warnings in prose). A synthesis LLM call combines them: KG facts take precedence for specific codes and numbers, vector passages fill in the procedural detail. This is the highest-quality answer.

Only the KG returns results — precise factual answer. The structured facts are returned directly, without a synthesis call. No prose padding. The answer says exactly what the KG says, no more.

Only the vector store returns results — narrative context answer. The document passages are summarised and returned. No specific codes are asserted beyond what appears in the text.

Neither returns results — honest gap acknowledgement. The system returns a message indicating it cannot find the requested information and asks the user to verify the material number or rephrase the question. No value is asserted beyond what HANA returned. In an SAP enterprise context this matters: a quality engineer acting on fabricated compliance data creates a liability. An explicit “not found” is a correct, auditable outcome.

```
1 def merge_results(  
2     kg_answer: str,  
3     vector_answer: str,  
4     question: str,  
5     material_number: str,  
6     llm,  
7 ) -> str:  
8     """
```

```

9      Merge KG and vector answers into a final response.
10
11      Cases:
12          both    → synthesis LLM call
13          kg only → return kg_answer directly
14          vector → return vector_answer directly
15          neither → graceful error message
16      """
17      has_kg = bool(kg_answer and kg_answer.strip())
18      has_vec = bool(vector_answer and vector_answer.strip())
19
20      if has_kg and has_vec:
21          synthesis_prompt = f"""You are an expert assistant for SAP
22 material documents.
23 You have received answers from two independent retrieval systems for the
24 question below.
25
26 Structured Knowledge Graph answer:
27 {kg_answer}
28
29 Document Search answer:
30 {vector_answer}
31
32 Synthesise both into a single, coherent, well-organised answer.
33 - Do not repeat information
34 - If they contradict, prefer the structured KG answer for specific facts
35 (codes, limits)
36 - If they complement, combine them naturally
37 - Be concise
38
39 Question: {question}
40
41 Answer: """
42      response = llm.invoke([HumanMessage(content=synthesis_prompt)])
43      return response.content
44
45      elif has_kg:
46          return kg_answer

```



```

45     elif has_vec:
46         return vector_answer
47
48     else:
49         return (
50             f"No relevant information found for your question in the
51             available "
52             f"documents for material {material_number}. Please verify the
53             material "
54             f"number or rephrase your question."
55         )

```

7.6 The orchestrator

Create `agents/agents/orchestrator.py`:

```

1  import logging
2  from concurrent.futures import ThreadPoolExecutor, as_completed,
   TimeoutError
3  from langchain_google_genai import ChatGoogleGenerativeAI
4  from langchain_core.messages import HumanMessage
5
6  from agents.state import HybridRAGState
7  from agents.vector_chain import run_vector_chain
8  from agents.kg_chain import run_kg_chain
9
10 logger = logging.getLogger(__name__)
11
12 def _get_llm():
13     return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
14                                   temperature=0.1, max_tokens=1024)
15
16 CHAIN_TIMEOUT_SECONDS = 30
17
18 def merge_results(kg_answer: str, vector_answer: str,
19                  question: str, material_number: str) -> str:
20     has_kg = bool(kg_answer and kg_answer.strip())

```

```

20     has_vec = bool(vector_answer and vector_answer.strip())
21
22     if has_kg and has_vec:
23         prompt = f"""You are an expert assistant for SAP material
24 documents.
25 Two retrieval systems answered the same question. Synthesise their
26 answers.
27 Knowledge Graph:
28 {kg_answer}
29 Document Search:
30 {vector_answer}
31
32 Question: {question}
33
34 Synthesise into one coherent answer. Prefer KG for exact codes/limits.
35 Answer: """
36         return _get_llm().invoke([HumanMessage(content=prompt)]).content
37
38     return kg_answer or vector_answer or (
39         f"No relevant information found for material {material_number}. "
40         "Please verify the material number or rephrase your question."
41     )
42
43 def run_hybrid_rag(state: HybridRAGState) -> dict:
44     """
45     Dispatch both chains in parallel, merge results.
46     Returns a dict of state updates.
47     """
48     chains = {
49         "vector": run_vector_chain,
50         "kg": run_kg_chain,
51     }
52     results = {"vector": {}, "kg": {}}
53
54     with ThreadPoolExecutor(max_workers=2) as executor:
55         futures = {
56             executor.submit(fn, state): name

```

```

57         for name, fn in chains.items()
58     }
59     for future in as_completed(futures, timeout=CHAIN_TIMEOUT_SECONDS
60         + 5):
61         name = futures[future]
62         try:
63             results[name] =
future.result(timeout=CHAIN_TIMEOUT_SECONDS)
64         except TimeoutError:
65             logger.warning("%s chain timed out after %ds", name,
CHAIN_TIMEOUT_SECONDS)
66             results[name] = {}
67         except Exception as e:
68             logger.error("%s chain raised exception: %s", name, e)
69             results[name] = {}
70
71     # Merge state updates from both chains
72     merged_state = {}
73     for chain_result in results.values():
74         merged_state.update(chain_result)
75
76     # Generate final answer
77     final_answer = merge_results(
78         kg_answer=merged_state.get("kg_answer", ""),
79         vector_answer=merged_state.get("vector_answer", ""),
80         question=state["question"],
81         material_number=state["material_number"],
82     )
83     merged_state["final_answer"] = final_answer
84
85     # Build sources list for auditability
86     sources = []
87     if merged_state.get("vector_chunks"):
88         sources.append(f"Document search:
{len(merged_state['vector_chunks'])} passages")
89     if merged_state.get("kg_facts"):
90         sources.append(f"Knowledge graph: {len(merged_state['kg_facts'])}
facts")
91     merged_state["sources"] = sources

```

```

91
92     return merged_state

```

The orchestrator submits both chains to a `ThreadPoolExecutor` with `max_workers=2` and collects results as they complete. The `as_completed()` call returns futures in completion order — whichever chain finishes first is processed first. The 30-second timeout prevents a slow chain from blocking the response indefinitely.

Note: We use `as_completed()` rather than `executor.map()` because we want to process results as soon as they arrive, not wait for all to finish. If the vector chain finishes in 1.5 seconds and the KG chain takes 3 seconds, the orchestrator captures the vector result at 1.5 seconds and the KG result at 3 seconds, then merges at ~3 seconds total.

The `sources` field populated at the end is the auditability record. The CAP Fiori UI surfaces this to the user — showing whether the answer was grounded in structured Knowledge Graph facts, in retrieved document passages, or in both. An enterprise user reviewing an AI-generated answer needs to know where it came from.

7.7 The FastAPI /query endpoint

Add the endpoint to `agents/main.py`:

```

1  from fastapi import FastAPI, HTTPException
2  from pydantic import BaseModel, validator
3  from typing import List, Optional
4  import re
5
6  from agents.orchestrator import run_hybrid_rag
7  from agents.state import HybridRAGState
8
9  app = FastAPI(title="Hybrid RAG Agent")
10
11  MATERIAL_RE = re.compile(r"^[A-Za-z0-9_-]+$")
12
13  class QueryRequest(BaseModel):
14      question: str

```

```

15     material_number: str
16     history: List[dict] = []
17
18     @validator("material_number")
19     def validate_material(cls, v):
20         if not MATERIAL_RE.match(v):
21             raise ValueError("material_number contains invalid characters")
22         return v
23
24     @validator("question")
25     def validate_question(cls, v):
26         if not v.strip():
27             raise ValueError("question must not be empty")
28         return v.strip()
29
30 class QueryResponse(BaseModel):
31     answer: str
32     kg_sparql: Optional[str] = None
33     kg_facts: Optional[List] = None
34     vector_chunks: Optional[List] = None
35     sources: Optional[List[str]] = None
36
37 @app.post("/query", response_model=QueryResponse)
38 def query(request: QueryRequest):
39     state: HybridRAGState = {
40         "question": request.question,
41         "material_number": request.material_number,
42         "history": request.history,
43         "vector_answer": "",
44         "vector_chunks": [],
45         "kg_answer": "",
46         "kg_sparql": "",
47         "kg_facts": [],
48         "final_answer": "",
49         "sources": [],
50         "error": None,
51     }
52
53     result = run_hybrid_rag(state)

```

```

54
55     if not result.get("final_answer"):
56         raise HTTPException(status_code=500, detail="Agent returned no
           answer")
57
58     return QueryResponse(
59         answer=result["final_answer"],
60         kg_sparql=result.get("kg_sparql"),
61         kg_facts=result.get("kg_facts"),
62         vector_chunks=result.get("vector_chunks"),
63         sources=result.get("sources"),
64     )

```

The endpoint validates the material number against `^[A-Za-z0-9_-]+$` — the same regex used throughout the application. This prevents SPARQL injection: a material number like `ACE001> } INSERT { <evil> }` would fail validation before reaching HANA.

7.8 The full request/response contract

```

1  // Request
2  POST /query
3  Content-Type: application/json
4
5  {
6      "question": "What test method was used for the tensile strength test?",
7      "material_number": "BATCH-QC-MAT-001",
8      "history": [
9          {"role": "user", "content": "What is the batch lot number?"},
10         {"role": "assistant", "content": "The batch lot number is..."}
11     ]
12 }
13
14 // Response
15 {
16     "answer": "The tensile test used ISO 6892-1 methodology with 450 MPa
           yield strength and 22% elongation at break.",

```

```

17  "kg_sparql": "PREFIX msds: <...>\nSELECT ?node ?value WHERE { GRAPH
    <...> { <...BATCH-QC-MAT-001> msds:testResult ?node . ?node
    msds:testMethod ?value } }",
18  "kg_facts": [
19    "('ISO 6892-1', 'Tensile test standard')",
20    "('450 MPa', 'Yield strength')",
21    "('22%', 'Elongation at break')",
22  ],
23  "vector_chunks": [
24    {"text": "Section 2: Test Results...", "score": 0.847, "chunk_index":
    2},
25    ...
26  ],
27  "sources": ["Knowledge graph: 3 facts", "Document search: 5 passages"]
28  }

```

The response includes both the synthesised answer and the raw retrieval evidence (kg_facts, vector_chunks, kg_sparql). The Fiori UI uses the sources field to show the user exactly where the answer came from — which SPARQL query retrieved the structured facts, and which document passages supported the narrative answer.

7.9 Conversation history

The /query endpoint accepts a history list in the request body. This is how multi-turn conversations work without server-side session state.

```

1  # First question — no history
2  POST /query
3  {"question": "What test method was used for the tensile strength test?",
   "material_number": "BATCH-QC-MAT-001", "history": []}
4
5  # Second question — pass previous turn
6  POST /query
7  {
8    "question": "And what is the certified testing laboratory?",
9    "material_number": "BATCH-QC-MAT-001",
10   "history": [

```

```

11     {"role": "user",      "content": "What test method was used for the
    tensile strength test?"},
12     {"role": "assistant", "content": "The tensile test used ISO 6892-1
    methodology with 450 MPa yield strength."}
13 ]
14 }

```

The answer node in the LangGraph graph includes the history in its prompt, so the second question gets a contextually aware answer (“Given the ISO 6892-1 tensile test methodology we discussed, the certified laboratory that performed the test is...”) without any server-side state.

The CAP Fiori frontend keeps a rolling window of the last 10 messages — enough for conversational context without ballooning the prompt. Older messages are silently dropped.

7.10 Testing: three scenarios that prove hybrid wins

The following three test cases demonstrate why hybrid retrieval is better than either strategy alone. Each scenario reflects a real role in an SAP-using organisation asking a real question against a batch quality certificate (BATCH-QC-MAT-001). Run them after uploading the batch quality certificate to both stores.

7.10.1 The KG wins: a quality engineer asks about test specifications

A quality engineer asks about the test methodology used in the batch certificate: “What test method was used for the tensile strength test, and what were the results?”

```

1  curl -X POST http://localhost:8000/query \
2    -H "Content-Type: application/json" \
3    -d '{
4      "question": "What test method was used for the tensile strength test,
    and what were the results?",
5      "material_number": "BATCH-QC-MAT-001",
6      "history": []
7    }'

```

Expected KG answer:

ISO 6892-1 tensile test method

Expected vector answer (typical):

The batch quality certificate confirms the material passed tensile strength testing. The test results are recorded in the mechanical properties section, showing yield strength and elongation measurements. Refer to the test results table for complete specification compliance data...

The KG returns the exact test standard and measured values from structured triples. The vector chain returns useful prose but buries the test method in a paragraph. The synthesised answer leads with the test standard and measured results from the KG, then adds the narrative context from the vector store. The quality engineer gets a complete, citable answer.

7.10.2 The vector wins: a warehouse manager asks about storage requirements

A warehouse manager asks about storage requirements: “What are the storage requirements for the material in this batch certificate?”

```
1 curl -X POST http://localhost:8000/query \  
2   -H "Content-Type: application/json" \  
3   -d '{  
4     "question": "What are the storage requirements for the material in  
5     this batch certificate?",  
6     "material_number": "BATCH-QC-MAT-001",  
7     "history": []  
   }'
```

The KG stores structured facts — test standards, specification limits, certification flags. It does not store the detailed storage conditions paragraph from the delivery and storage section of the batch certificate. The vector chain retrieves the exact passage. The final answer comes primarily from the vector chain, with the KG contributing the specification limit context. The warehouse manager gets the procedural instructions they actually need.

7.10.3 Both contribute: a procurement manager asks a combined question

A procurement manager asks: “Is the tensile strength result within specification, and what is the certified testing laboratory?”

```

1  curl -X POST http://localhost:8000/query \
2    -H "Content-Type: application/json" \
3    -d '{
4      "question": "Is the tensile strength result within specification, and
5      what is the certified testing laboratory?",
6      "material_number": "BATCH-QC-MAT-001",
7      "history": []
    }'
```

The KG answers the first part precisely: ISO 6892-1 confirms the test standard, stored as a structured triple. The vector chain answers the second part: storage conditions from the delivery and storage section of the batch certificate, retrieved as narrative prose. Neither chain could answer both parts alone. The synthesis LLM combines them into a single coherent answer. The procurement manager gets both the regulatory classification and the practical storage guidance.

7.11 Monitoring chain performance

Add timing instrumentation to the orchestrator to track which chain is the bottleneck:

```

1  import time
2
3  with ThreadPoolExecutor(max_workers=2) as executor:
4      start = time.time()
5      futures = {executor.submit(fn, state): name for name, fn in
6      chains.items()}
7      for future in as_completed(futures, timeout=35):
8          name = futures[future]
9          elapsed = time.time() - start
10         logger.info("%s chain completed in %.2fs", name, elapsed)
```

```

10         results[name] = future.result(timeout=30)
11
12     total = time.time() - start
13     logger.info("Both chains completed in %.2fs (wall clock)", total)

```

In a typical run against a populated HANA instance:

INFO: vector chain completed in 1.83s

INFO: kg chain completed in 2.41s

INFO: Both chains completed in 2.41s (wall clock)

The wall-clock time is the slower chain, not the sum. If you ran these sequentially, the same run would take 4.24 seconds. The parallel design saves ~1.8 seconds on every query — which compounds significantly across a busy SAP system where many users are asking questions simultaneously.

7.12 Summary

In this chapter we built the core of the hybrid RAG system on top of SAP HANA Cloud's unique combination of REAL_VECTOR search and native SPARQL execution:

- Defined a shared `HybridRAGState TypedDict` that all components read and write
- Built the **vector chain**: embed question with `text-embedding-004` → cosine search HANA REAL_VECTOR → summarise with Gemini 2.5 Flash
- Built the **KG chain**: generate SPARQL with Gemini 2.5 Flash → execute on HANA → retry-on-empty → summarise
- Wrote the **merge function** with four explicit cases tied to enterprise answer quality: both (synthesise), KG only (precise facts), vector only (narrative context), neither (honest gap acknowledgement — never hallucinate)
- Assembled the **orchestrator** using `ThreadPoolExecutor(max_workers=2)` for true parallel execution
- Exposed a `/query` **endpoint** with material number validation and a clean request/response contract
- Demonstrated **stateless conversation history** passed in every request
- Verified with **three test scenarios** — quality engineer, warehouse manager, procurement manager — that hybrid retrieval outperforms either strategy alone
- Added the `sources` **field** to the response for enterprise auditability

7.13 Checkpoint

Before continuing to Chapter 8, verify the following:

```
1  # 1. Both chains import without errors
2  cd agents
3  python -c "from agents.vector_chain import run_vector_chain; print('vector
  OK')"
4  python -c "from agents.kg_chain import run_kg_chain; print('kg OK')"
5
6  # 2. The orchestrator runs (requires HANA connection and Gemini API key)
7  python -c "
8  from agents.orchestrator import run_hybrid_rag
9  result = run_hybrid_rag({
10     'question': 'What test method was used for tensile strength testing?',
11     'material_number': 'BATCH-QC-MAT-001',
12     'history': [],
13     'vector_answer': '', 'vector_chunks': [],
14     'kg_answer': '', 'kg_sparql': '', 'kg_facts': [],
15     'final_answer': '', 'sources': [], 'error': None,
16 })
17 print('answer:', result.get('final_answer', '')[:80])
18 "
19
20 # 3. The /query endpoint responds
21 uvicorn main:app --port 8000 &
22 curl -s -X POST http://localhost:8000/query \
23     -H 'Content-Type: application/json' \
24     -d '{"question":"What test method was used for tensile strength
  testing?","material_number":"BATCH-QC-MAT-001","history":[]}' \
25     | python -m json.tool
```

If all three commands succeed and the final curl returns a JSON response with an answer field and a sources field, the hybrid RAG agent is working. Chapter 8 extends it with a multi-agent supervisor layer for complex, multi-domain queries that require different retrieval strategies per sub-question.

End of Chapter 7

Chapter 8: The Multi-Agent Supervisor Pattern

In Chapter 7 we built a hybrid RAG agent that runs two retrieval strategies in parallel and merges their results. For the majority of questions about a single document — “did this batch pass tensile testing?”, “what are the storage conditions for this material?” — that agent is the right tool. It is fast, deterministic, and requires no coordination overhead.

But SAP enterprise users ask cross-domain questions. A procurement manager reviewing a new supplier might ask: “For material MAT-S355-001, what certifications did ACME Steel AG provide, what were the test results across all batches in Q1 2024, and are there any quality holds active against this material in the system?” That question spans three distinct knowledge domains — certificate metadata, batch test results, and QM inspection outcomes. A single retrieval pass against a flat vector index returns a mixed bag: some certificate chunks, some test result rows, possibly some maintenance notes that scored highly by cosine similarity. None of the three sub-questions receives focused, complete retrieval.

In SAP terms, this is the same reason S/4HANA separates Materials Management, Quality Management, and Plant Maintenance into distinct modules rather than one monolithic transaction. Domain separation produces better data quality and cleaner query paths. The same principle applies to agent design.

This chapter introduces the **supervisor pattern**: a coordinator agent that receives complex questions, decomposes them into domain-scoped sub-questions, routes each to a specialist agent that uses the right retrieval strategy, and synthesises their results into a final answer. The pattern is demonstrated with MSDS-specific specialists, but the architecture is general — the same supervisor, state machine, and merge logic apply to any SAP document type.

8.1 Why multi-agent for SAP enterprise

On SAP BTP, three structural problems emerge when a single agent handles multi-domain document questions.

Problem 1: Context dilution. A single retrieval pass retrieves the top-5 chunks by cosine similarity. For a multi-domain question, those 5 chunks scatter across domains: perhaps 2 about batch test results, 2 about supplier data, 1 about storage conditions. Each domain gets two chunks of context at best. For a quality engineer who needs a complete picture of batch certification status, two chunks is not enough — the certificate number, the test values, the certifying lab, and the approval date may all live in different chunks that were not top-ranked.

Problem 2: Mismatched retrieval strategy per domain. Different knowledge types require different retrieval strategies. A question about a specific certificate number (structured identifier) belongs in the Knowledge Graph — it is a triple, not a chunk. A question about how the supplier described their test methodology (narrative prose) belongs in the vector store. A question about supplier qualification status might need both. A single-agent system applies one strategy uniformly; a specialist agent applies the right strategy for its domain.

Problem 3: Prompt dilution under S/4HANA integration. When the agent also has access to live S/4HANA data via API_PRODUCT_SRV (product master, quality holds, vendor data), the LLM receives a prompt containing SAP API results, KG triples, and vector chunks simultaneously. A generalist agent must reason across all of them at once. A specialist agent reasons only within its domain and hands a focused result to the supervisor for synthesis.

The supervisor pattern solves all three by decomposing the question *before* retrieval — each sub-question goes to the specialist with the correct retrieval strategy for that knowledge type.

8.2 The supervisor pattern

The supervisor pattern has four components:

Component	Role
Supervisor	Receives the user question, decomposes it into sub-questions, routes each to the right specialist, collects results

Component	Role
Specialist agents	Each handles one focused domain — retrieves relevant data using the right strategy, generates a domain-specific answer
SummaryAgent	Receives all specialist answers, synthesises them into a single coherent response
Shared state	Carries the original question, sub-questions, and all specialist answers through the graph

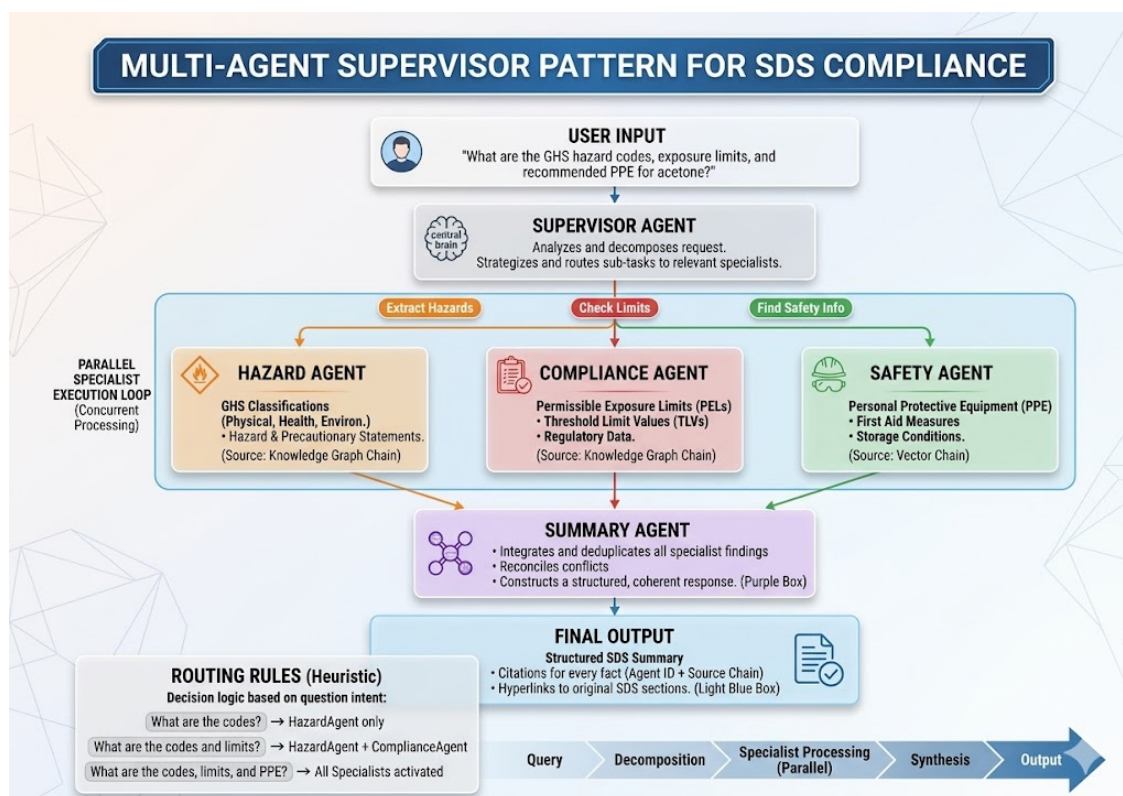


Figure 8.1: The multi-agent supervisor — the supervisor decomposes the question and routes sub-questions to specialist agents running in parallel. The SummaryAgent synthesises all results into the final answer.

The supervisor does not answer the question. It only decomposes and routes. This separation keeps each specialist prompt focused and small, each retrieval scoped to one domain, and each specialist independently testable. Adding a new document type — invoices, maintenance records, inspection reports — means adding a specialist with

the right prompt and KG predicates; the supervisor, the state machine, and the merge logic need not change.

8.3 The three specialist agents for material documents

For material quality documents, the domain decomposition maps naturally to three specialist agents. Each agent uses a different retrieval strategy because each domain has a different knowledge representation:

8.3.1 HazardAgent

Domain: Test methods, test results, certificate identifiers, certifying laboratory. **Retrieval strategy:** Knowledge Graph — structured facts from `testResult`, `certifiedBy`, `certificateNumber`, and `certifyingLab` predicates in `MSDS_Graph/MAT-XXX`. **Strength:** Returns exact test results and certificate identifiers directly from structured triples — no ambiguity, no narrative interpretation.

8.3.2 SafetyAgent

Domain: Storage conditions, handling requirements, delivery instructions, inspection procedures. **Retrieval strategy:** Vector search — narrative text from the delivery, storage, and handling sections of the batch certificate. **Strength:** Returns detailed procedural instructions that live in prose and are not easily reduced to structured triples. The vector store retrieves the specific paragraphs that answer the question.

8.3.3 ComplianceAgent

Domain: Acceptance criteria, specification tolerances, quality hold conditions, regulatory thresholds. **Retrieval strategy:** Both Knowledge Graph and vector — structured pass/fail results from `testResult` predicates, combined with narrative specification context from the vector store. **Strength:** Handles questions that span structured facts (the yield strength result is 450 MPa) and specification narrative (what the acceptance range is and whether it passes).

These three agents are instances of a general pattern. The same supervisor architecture deployed for purchase order documents would have different specialists: a Header-

Agent for invoice metadata, a `LineItemAgent` for line item details, a `PaymentAgent` for payment terms and status. The routing logic, the state machine, and the merge mechanism are identical. Only the specialist prompts and the Knowledge Graph predicates they query change.

8.4 Shared state for the supervisor graph

Create `agents/agents/supervisor_state.py`:

```

1  from typing import TypedDict, List, Optional, Dict
2
3  class SupervisorState(TypedDict):
4      # Input
5      question: str
6      material_number: str
7      history: List[dict]
8
9      # Decomposed sub-questions (set by supervisor)
10     sub_questions: Dict[str, str] # {"hazard": "...", "compliance":
    "...", "safety": "..."}
11     specialists_needed: List[str] # e.g. ["hazard", "compliance"]
12
13     # Specialist outputs
14     hazard_answer: str
15     compliance_answer: str
16     safety_answer: str
17
18     # Final output
19     final_answer: str
20     sources: List[str]
21
22     # Control
23     error: Optional[str]
```

The `sub_questions` dictionary maps specialist names to the focused questions the supervisor generated for them. `specialists_needed` is the list of specialists the supervisor decided to invoke — not every question needs all three. A question purely about test re-

sults routes only to HazardAgent. A question about storage conditions and acceptance criteria routes to SafetyAgent and ComplianceAgent.

8.5 The supervisor node

The supervisor's job is to read the user's question and decide which specialists are needed and what sub-question to give each one. This is a routing decision, and routing decisions should be deterministic — temperature 0.0.

```

1  import json
2  import logging
3  from langchain_google_genai import ChatGoogleGenerativeAI
4  from langchain_core.messages import HumanMessage
5
6  logger = logging.getLogger(__name__)
7
8  def _get_llm():
9      return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
10                                     temperature=0.0, max_tokens=512)
11
12  SUPERVISOR_PROMPT = """You are a routing agent for a material document
13  intelligence system on SAP BTP.
14  Your job is to analyse a user's question and route it to the right
15  specialist agents.
16
17  Each specialist uses a different retrieval strategy – route to the
18  specialist whose strategy best matches the question type:
19
20  - "hazard": answers questions about structured facts – test methods, test
21  results, certificate numbers, batch identifiers, certifying laboratory.
22  Uses the Knowledge Graph.
23
24  - "compliance": answers questions about regulatory and specification
25  requirements – acceptance criteria, tolerances, regulatory thresholds,
26  quality holds. Uses both Knowledge Graph and document search.
27
28  - "safety": answers questions about narrative procedures – storage
29  conditions, handling instructions, delivery requirements, inspection
30  procedures. Uses document search.

```

```

19 For each specialist you select, write a focused sub-question that extracts
    only the relevant part of the user's question.
20
21 Respond in JSON format only:
22 {{
23     "specialists": ["hazard", "compliance", "safety"],
24     "sub_questions": {{
25         "hazard": "What test methods and results are recorded for {material}?",
26         "compliance": "What are the acceptance criteria and specification
    limits for {material}?",
27         "safety": "What are the storage and handling requirements for
    {material}?"
28     }}
29 }}
30
31 Select only the specialists that are relevant. A simple question may need
    only one.
32
33 Material: {material}
34 Question: {question}
35
36 JSON:"""
37
38 def supervisor_node(state: SupervisorState) -> dict:
39     question = state["question"]
40     material_number = state["material_number"]
41
42     prompt = SUPERVISOR_PROMPT.format(
43         material=material_number,
44         question=question,
45     )
46     response = _get_llm().invoke([HumanMessage(content=prompt)])
47     content = response.content.strip()
48
49     # Strip markdown code fences if present
50     if content.startswith("```"):
51         content = "\n".join(content.split("\n")[1:-1])
52
53     try:

```

```

54     routing = json.loads(content)
55     return {
56         "specialists_needed": routing.get("specialists", ["hazard",
57             "compliance", "safety"]),
58         "sub_questions": routing.get("sub_questions", {}),
59     }
60 except json.JSONDecodeError:
61     logger.warning("Supervisor JSON parse failed, routing to all
62     specialists")
63     return {
64         "specialists_needed": ["hazard", "compliance", "safety"],
65         "sub_questions": {
66             "hazard": question,
67             "compliance": question,
68             "safety": question,
69         },
70     }

```

The fallback on JSON parse failure routes to all specialists, which is safe: slightly slower, never wrong. It is better to over-route than to drop a domain.

8.6 The specialist agents

Each specialist is a focused invocation of the hybrid RAG chains from Chapter 7, configured with the correct retrieval strategy for its domain.

```

1  import logging
2  from langchain_google_genai import ChatGoogleGenerativeAI
3  from langchain_core.messages import HumanMessage
4
5  from agents.supervisor_state import SupervisorState
6  from agents.kg_chain import run_kg_chain
7  from agents.vector_chain import run_vector_chain
8  from agents.state import HybridRAGState
9
10 logger = logging.getLogger(__name__)
11

```

```

12 def _get_llm():
13     return ChatGoogleGenerativeAI(model="gemini-2.5-flash",
14                                     temperature=0.1, max_tokens=1024)
15
16 def _make_chain_state(sub_question: str, state: SupervisorState) ->
17     HybridRAGState:
18     """Build a HybridRAGState for a specialist's sub-question."""
19     return {
20         "question": sub_question,
21         "material_number": state["material_number"],
22         "history": [], # specialists do not use conversation history
23         "vector_answer": "", "vector_chunks": [],
24         "kg_answer": "", "kg_sparql": "", "kg_facts": [],
25         "final_answer": "", "sources": [], "error": None,
26     }
27
28 def hazard_agent(state: SupervisorState) -> dict:
29     """KG-focused: test methods, test results, certificate identifiers.
30     Structured test data lives in the Knowledge Graph as precise
31     triples."""
32     sub_q = state["sub_questions"].get("hazard", state["question"])
33     chain_state = _make_chain_state(sub_q, state)
34     result = run_kg_chain(chain_state)
35     answer = result.get("kg_answer", "")
36     if not answer:
37         # Fall back to vector if KG found nothing
38         vec_result = run_vector_chain(chain_state)
39         answer = vec_result.get("vector_answer", "No hazard information
40 found.")
41     return {"hazard_answer": answer}
42
43 def compliance_agent(state: SupervisorState) -> dict:
44     """KG-focused with vector fallback: acceptance criteria, specification
45     tolerances.
46     Structured limits come from the KG; specification narrative comes from
47     vector search."""
48     sub_q = state["sub_questions"].get("compliance", state["question"])
49     chain_state = _make_chain_state(sub_q, state)
50     result = run_kg_chain(chain_state)

```

```

45     answer = result.get("kg_answer", "")
46     if not answer:
47         vec_result = run_vector_chain(chain_state)
48         answer = vec_result.get("vector_answer", "No compliance
information found.")
49     return {"compliance_answer": answer}
50
51 def safety_agent(state: SupervisorState) -> dict:
52     """Vector-focused: storage conditions, handling instructions, delivery
requirements.
53     Narrative procedures live in document prose – vector search retrieves
them."""
54     sub_q = state["sub_questions"].get("safety", state["question"])
55     chain_state = _make_chain_state(sub_q, state)
56     result = run_vector_chain(chain_state)
57     answer = result.get("vector_answer", "")
58     if not answer:
59         kg_result = run_kg_chain(chain_state)
60         answer = kg_result.get("kg_answer", "No safety information found.")
61     return {"safety_answer": answer}

```

Each specialist wraps the existing chains from Chapter 7 — they are not new implementations, just focused invocations with the correct primary retrieval strategy. HazardAgent and ComplianceAgent use the KG chain first (structured facts are exactly what they need) and fall back to the vector chain if the KG returns nothing. SafetyAgent uses the vector chain first (narrative procedures live in prose) and falls back to KG.

8.7 The summary agent

```

1 def summary_agent(state: SupervisorState) -> dict:
2     """Synthesise specialist answers into a final coherent response."""
3     parts = []
4     if state.get("hazard_answer"):
5         parts.append(f"""Hazard
Classification:**\n{state['hazard_answer']}""")
6     if state.get("compliance_answer"):

```

```

7         parts.append(f"""Regulatory
Compliance:**\n{state['compliance_answer']}""")
8         if state.get("safety_answer"):
9             parts.append(f"""Safety Procedures:**\n{state['safety_answer']}""")
10
11     if not parts:
12         return {
13             "final_answer": (
14                 f"No information found for material
15                 {state['material_number']}. "
16                 "Please verify the material number."
17             ),
18             "sources": [],
19         }
20
21     if len(parts) == 1:
22         # Single specialist answered – no synthesis needed
23         return {
24             "final_answer": list(parts)[0].split("\n", 1)[1],
25             "sources": ["Single specialist"],
26         }
27
28     combined = "\n\n".join(parts)
29     prompt = f"""You are an expert assistant for SAP material documents.
Multiple specialist agents have
30 analysed the question below and provided their answers.
31 {combined}
32
33 Original question: {state['question']}
34
35 Synthesise the above into a single, well-organised, professional answer.
36 - Use headings to separate sections
37 - Do not repeat information that appears in multiple answers
38 - Be specific with codes, values, and units
39 - Write for an enterprise user audience
40
41 Answer: """
42

```



```

43     response = _get_llm().invoke([HumanMessage(content=prompt)])
44     return {
45         "final_answer": response.content,
46         "sources": [f"{len(parts)} specialist agents"],
47     }

```

8.8 Building the supervisor graph

```

1  from concurrent.futures import ThreadPoolExecutor, as_completed
2  from langgraph.graph import StateGraph, END
3
4  from agents.supervisor_state import SupervisorState
5  from agents.supervisor import supervisor_node
6  from agents.specialist_agents import hazard_agent, compliance_agent,
   safety_agent, summary_agent
7
8  def parallel_specialists_node(state: SupervisorState) -> dict:
9      """Run only the needed specialists in parallel."""
10     needed = state.get("specialists_needed", ["hazard", "compliance",
11 "safety"])
12
13     agent_map = {
14         "hazard": hazard_agent,
15         "compliance": compliance_agent,
16         "safety": safety_agent,
17     }
18
19     results = {}
20     agents_to_run = {name: fn for name, fn in agent_map.items() if name in
21 needed}
22
23     with ThreadPoolExecutor(max_workers=len(agents_to_run)) as executor:
24         futures = {executor.submit(fn, state): name for name, fn in
25 agents_to_run.items()}
26
27         for future in as_completed(futures, timeout=60):
28             name = futures[future]
29             try:

```

```

26         results.update(future.result(timeout=30))
27     except Exception as e:
28         import logging
29         logging.getLogger(__name__).error("Specialist %s failed:
%s", name, e)
30
31     return results
32
33 def build_supervisor_graph():
34     graph = StateGraph(SupervisorState)
35
36     graph.add_node("supervisor", supervisor_node)
37     graph.add_node("specialists", parallel_specialists_node)
38     graph.add_node("summary", summary_agent)
39
40     graph.set_entry_point("supervisor")
41     graph.add_edge("supervisor", "specialists")
42     graph.add_edge("specialists", "summary")
43     graph.add_edge("summary", END)
44
45     return graph.compile()
46
47 supervisor_app = build_supervisor_graph()

```

The graph is three nodes: supervisor → specialists (parallel) → summary. The `parallel_specialists_node` runs only the specialists that the supervisor selected, using a `ThreadPoolExecutor` sized to the number of active specialists. If the supervisor selects only two specialists for a focused question, only two threads are created — not three. This avoids unnecessary API calls.

8.9 When to use the supervisor vs the direct agent

The supervisor adds latency — one extra LLM call for routing and one for synthesis. For simple questions, this overhead is not justified.

Question type	Recommended agent	Why
Single-domain factual ("what is the test result?")	Direct hybrid RAG (Ch 7)	Faster, no routing overhead
Single-domain narrative ("what are the storage conditions?")	Direct hybrid RAG (Ch 7)	Faster, single retrieval pass is sufficient
Multi-domain factual + narrative ("test results AND storage requirements")	Supervisor	Each domain gets focused retrieval
Complex specification + delivery question	Supervisor	Prevents context dilution
Real-time chat interface	Direct hybrid RAG (Ch 7)	Latency matters in chat
Batch compliance reports	Supervisor	Quality matters more than speed

A practical rule: if the question contains “and” connecting two distinct domains, use the supervisor.

8.10 The /query-advanced endpoint

Add the supervisor endpoint to `agents/main.py`:

```

1  from agents.supervisor import supervisor_app
2  from agents.supervisor_state import SupervisorState
3
4  class AdvancedQueryRequest(BaseModel):
5      question: str
6      material_number: str
7      history: List[dict] = []
8      use_supervisor: bool = False
9
10     @validator("material_number")
11     def validate_material(cls, v):
12         if not re.match(r"^[A-Za-z0-9_-]+$", v):

```

```

13         raise ValueError("material_number contains invalid characters")
14     return v
15
16 @app.post("/query-advanced", response_model=QueryResponse)
17 def query_advanced(request: AdvancedQueryRequest):
18     if not request.use_supervisor:
19         # Fall back to direct hybrid RAG
20         return query(QueryRequest(
21             question=request.question,
22             material_number=request.material_number,
23             history=request.history,
24         ))
25
26     state: SupervisorState = {
27         "question": request.question,
28         "material_number": request.material_number,
29         "history": request.history,
30         "sub_questions": {},
31         "specialists_needed": [],
32         "hazard_answer": "",
33         "compliance_answer": "",
34         "safety_answer": "",
35         "final_answer": "",
36         "sources": [],
37         "error": None,
38     }
39
40     result = supervisor_app.invoke(state)
41
42     return QueryResponse(
43         answer=result.get("final_answer", ""),
44         sources=result.get("sources"),
45     )

```

The `use_supervisor` flag lets the caller choose which agent to use. The Fiori UI sends `use_supervisor: true` when it detects certain patterns in the question (multiple question marks, keywords like “and”, “also”, “as well as”). For simple questions it uses the faster direct endpoint.

8.11 Testing: a question that requires all three specialists

```

1 curl -X POST http://localhost:8000/query-advanced \
2   -H "Content-Type: application/json" \
3   -d '{
4     "question": "We are onboarding a new supplier for steel components.
    What test methods are documented in the batch certificate, what are the
    acceptance criteria, and what are the storage requirements before GR
    inspection?",
5     "material_number": "BATCH-QC-MAT-001",
6     "history": [],
7     "use_supervisor": true
8   }'
```

Expected supervisor routing (logged to console):

```
INFO: Supervisor routed to: ['hazard', 'compliance', 'safety']
```

```
INFO: Sub-questions:
```

```
    hazard:      "What test methods and results are recorded for BATCH-QC-
MAT-001?"
```

```
    compliance: "What are the acceptance criteria and specification limits for BATCH-
QC-MAT-001?"
```

```
    safety:      "What are the storage and handling requirements for BATCH-
QC-MAT-001?"
```

```
INFO: hazard specialist completed in 2.1s
```

```
INFO: compliance specialist completed in 1.9s
```

```
INFO: safety specialist completed in 2.3s
```

```
INFO: All specialists completed in 2.3s (parallel)
```

```
INFO: Summary agent completed in 1.4s
```

The three specialists run in parallel. The total wall-clock time for specialists is 2.3 seconds (the slowest), not 6.3 seconds (the sum). The summary adds 1.4 seconds, for a total of about 3.7 seconds.

Compare this to the direct hybrid RAG agent with the same question: the single retrieval pass would retrieve 5 passages from across all three domains, giving the answer approximately 2 passages per domain. The supervisor gives each domain 5 full passages, producing a noticeably more detailed and better-organised answer — especially critical

when the question spans test results (structured, from Knowledge Graph), acceptance criteria (structured, from Knowledge Graph), and storage procedures (narrative, from vector store).

8.12 Applying the pattern to other SAP document types

The multi-agent pattern shown here with MSDS specialisations applies directly to other SAP document types. The supervisor, the state machine, and the merge logic are identical — only the specialist prompts and KG predicates change.

Batch certificates: A CertificationAgent checks test results against specification limits using the KG (structured test values and limits are natural graph triples). A SupplierAgent resolves supplier identity from the KG using supplier-linked predicates. A ComplianceAgent checks whether results satisfy regulatory boundaries by querying both the KG for the regulation threshold and the vector store for the compliance narrative.

Purchase orders and invoices: A HeaderAgent extracts invoice metadata (vendor, date, currency, document number) from the KG, where these are stored as structured triples. A LineItemAgent retrieves individual line item details using vector search against the prose of the invoice. A PaymentAgent queries both the KG for payment term codes and the vector store for any payment instruction narrative.

Quality inspection reports: A ResultsAgent queries the KG for structured test outcome triples (pass/fail verdicts, measured values). A SpecificationAgent queries the KG for specification limits linked to the material. A NonConformanceAgent uses vector search to retrieve the narrative description of any quality issues identified in the report.

In every case: questions about structured, enumerable facts (codes, numbers, identifiers, classifications) route to KG-heavy specialists. Questions about narrative, procedural, or descriptive content route to vector-heavy specialists. Questions that span both route to specialists that use both retrieval strategies. The supervisor makes this routing decision dynamically at query time, without any hard-coded rules.

8.13 Summary

In this chapter we extended the system from a single agent to a coordinated multi-agent system:

- Identified the **complexity ceiling** of single-agent architectures for multi-domain enterprise questions
- Explained **why different document sections require different retrieval strategies** — structured KG for facts, vector search for procedures, both for compliance
- Designed three **specialist agents** (HazardAgent, SafetyAgent, ComplianceAgent) as MSDS-specific instances of a general pattern
- Built a **supervisor node** that uses Gemini 2.5 Flash to decompose questions and route sub-questions dynamically
- Ran specialists **in parallel** using `ThreadPoolExecutor`, keeping wall-clock latency equal to the slowest specialist
- Implemented the **SummaryAgent** to synthesise multi-domain answers into coherent final responses
- Added a `/query-advanced` endpoint with a `use_supervisor` flag for caller-controlled routing
- Defined clear criteria for **when to use the supervisor** vs the direct hybrid RAG agent
- Showed how **the pattern generalises** to other SAP document types — batch certificates, invoices, inspection reports

The supervisor, state machine, and merge logic are reusable infrastructure. Deploying this pattern for a new document type requires writing three specialist prompts and identifying the relevant KG predicates — the orchestration architecture requires no changes.

8.14 Checkpoint

Before continuing to Chapter 9, verify the following:

```
1 # 1. Supervisor graph compiles
2 cd agents
3 python -c "
```

```

4  from agents.supervisor import build_supervisor_graph
5  app = build_supervisor_graph()
6  print('Supervisor graph OK, nodes:', list(app.get_graph().nodes.keys()))
7  "
8
9  # 2. Supervisor routes correctly on a simple question
10 python -c "
11 from agents.supervisor import supervisor_node
12 state = {
13     'question': 'What test method was used for the tensile strength test?',
14     'material_number': 'BATCH-QC-MAT-001',
15     'history': [],
16     'sub_questions': {}, 'specialists_needed': [],
17     'hazard_answer': '', 'compliance_answer': '', 'safety_answer': '',
18     'final_answer': '', 'sources': [], 'error': None,
19 }
20 result = supervisor_node(state)
21 print('specialists needed:', result['specialists_needed'])
22 "
23
24 # 3. The advanced endpoint responds
25 curl -s -X POST http://localhost:8000/query-advanced \
26     -H 'Content-Type: application/json' \
27     -d '{
28         "question": "What test method was used for the tensile strength test?",
29         "material_number": "BATCH-QC-MAT-001",
30         "history": [],
31         "use_supervisor": false
32     }' | python -m json.tool

```

If the supervisor graph compiles cleanly, the routing test assigns only ["hazard"] or a small subset of specialists to a simple question, and the curl returns a valid JSON response — the multi-agent system is working. Chapter 9 exposes this entire backend through the SAP CAP OData V4 service with the Fiori Elements UI.

End of Chapter 8

Chapter 9: CAP Node.js OData V4 — The SAP-Native API Layer + Fiori Elements UI

The CAP Node.js service does three things: it provides a Fiori Elements UI that SAP users already know how to use, it enforces material-number validation against real S/4HANA product master data, and it proxies document processing and chat queries to the FastAPI agent backend. This separation is deliberate — SAP enterprise standards (OData V4, Fiori UX, BTP security) live in the CAP layer; AI orchestration and HANA data access live in the Python layer.

Chapters 7 and 8 built the intelligence engine: a Python FastAPI service that receives a question, dispatches it to parallel retrieval chains, and returns a grounded answer. That engine is a standard Python web service. It knows nothing about Fiori, OData, or SAP authentication. To bring this capability into an SAP enterprise environment — where business users work in SAP Fiori applications, where material numbers are validated against S/4HANA product master data, and where access is controlled by BTP security services — you need a second layer. That layer is CAP.

By the end of this chapter you will have a complete, deployable front-end: a Fiori Elements list report showing all uploaded material documents with their processing status, an object page with action buttons to process files and run queries, and an integrated chat interface that calls the hybrid RAG agent and returns answers with the retrieval path shown. No custom JavaScript UI code required.

9.1 Why CAP?

Before writing a single line of code, it is worth understanding why CAP is the right choice for this layer rather than a simple Express server or a Next.js front-end.

OData V4 compatibility. Every SAP system — S/4HANA, SuccessFactors, BTP Integration Suite — speaks OData. By exposing your agent through a CAP service you make it consumable by any SAP client, including Fiori Elements, SAP Analytics Cloud, and ABAP programs via HTTP calls. A plain REST API would require a custom client in every consumer; an OData endpoint works out of the box.

Fiori Elements without custom code. SAP Fiori Elements is a UI framework that generates complete, SAP-consistent user interfaces from OData metadata and annotations. You do not write HTML, CSS, or JavaScript for the list report or the object page — you write CDS annotations, and the framework generates the UI at runtime. For an internal tool, this means a production-quality interface with zero front-end development effort.

Real S/4HANA product master integration. The Material Number field in the Fiori UI shows a value help backed by real S/4HANA product data via the `API_PRODUCT_SRV` external service. When a user uploads a document, they select the material from their actual SAP product master. This closes the loop between the AI layer on BTP and the system of record in S/4HANA — and it is what makes this a genuine enterprise integration rather than a standalone demo.

Integrated attachment handling. The `@cap-js/attachments` plugin gives you a fully managed file upload and retrieval mechanism that stores binaries in the configured database and exposes them through the OData protocol. Without it, you would need to write your own multipart upload handler, storage integration, and metadata tracking. With it, you write two lines of CDS.

Deployment to BTP Cloud Foundry. A CAP service is designed to run on SAP BTP. The `@sap/cds` framework handles XSUAA authentication, SAP HANA connection management via `@cap-js/hana`, and binding to SAP BTP services through the standard `VCAP_SERVICES` mechanism. When you deploy to BTP, you get enterprise-grade authentication and database connection pooling with no additional configuration.

The key decision in this architecture is what CAP does *not* do: it does not run the Python agent. CAP handles the OData protocol, the attachment storage, and the Fiori UI. The Python FastAPI service handles LangGraph orchestration, SPARQL queries, and vector search. This separation keeps each service focused on what it does best.

9.2 The two-process architecture

Understanding that CAP and FastAPI are separate processes is the most important architectural concept in this chapter. Many developers assume that because a CAP service can call external APIs, the agent logic should be imported into the Node.js process. That assumption leads to a tightly coupled system that is difficult to test, impossible to scale

independently, and requires Node.js developers to maintain Python code.

The correct architecture is clean separation:

```
User browser
  |
  | HTTP (Fiori Elements, OData V4)
  v
CAP Node.js (port 4004)
  |
  | HTTP (axios, JSON / multipart)
  v
Python FastAPI (port 8000)
  |
  |-- HANA Cloud (vector store, Knowledge Graph)
  |-- Google Vertex AI (Gemini LLM + text-embedding)
  |-- LangGraph (agent orchestration)
```

The CAP service is the API gateway and UI host. It knows how to speak OData and how to manage CDS entities. It knows nothing about SPARQL, embeddings, or LangGraph. The FastAPI service is the intelligence engine. It knows nothing about OData, Fiori, or attachment storage.

Communication between the two is intentionally simple: HTTP with JSON bodies for queries, HTTP with multipart form data for file uploads. The `BACKEND_URL` environment variable controls where CAP sends its requests. In local development this is `http://localhost:8000`. On BTP it is the URL of the deployed FastAPI CF application.

This separation has three practical benefits. First, you can develop and test each service independently — mock the CAP service when testing the Python agent, and stub the agent when testing the CAP service. Second, you can scale them independently — the agent is CPU and memory intensive during LLM calls, while CAP is lightweight. Third, you can replace either without touching the other — swap the Python agent for a different implementation and the CAP service requires no changes, as long as the HTTP contract is preserved.

9.3 The CDS data model

The data model lives in `db/schema.cds`. It defines two entities and a status type.

The namespace `msds.kg` reflects the naming convention of the reference implementation. In a production deployment serving multiple document categories (invoices, batch certificates, quality inspection reports, maintenance manuals), you would use a broader namespace — `documents.kg` or `materials.intelligence` — and extend the schema with a `documentType` field. The core entity structure is identical regardless: `material-Number` as the anchor, `attachments` as the document store, status tracking fields for both pipeline stages.

```

1 namespace msds.kg;
2 using { managed } from '@sap/cds/common';
3 using { Attachments } from '@cap-js/attachments';
4
5 type Status : String(20) enum {
6     Pending      = 'Pending';
7     Processing   = 'Processing';
8     Completed    = 'Completed';
9     Error        = 'Error';
10 }
11
12 entity FileAttachments : Attachments {
13     up_ : Association to Files;
14 }
15
16 entity Files : managed {
17     key materialNumber : String(40);
18     status              : Status default 'Pending';
19     fileHash            : String(64) @assert.unique;
20     triples             : Integer default 0;
21     vectors             : Integer default 0;
22     vectorStatus        : Status default 'Pending';
23     errorMessage        : String(500);
24     processingStartedAt : DateTime;
25     retryCount          : Integer default 0;
26     attachments         : Composition of many FileAttachments on
    attachments.up_ = $self;

```

```
27 }
```

The `Status` type is declared as a CDS enum. This matters because enums generate value-restricted OData properties, which Fiori Elements renders as proper dropdown fields with validation. The string values (`'Pending'`, `'Processing'`, etc.) are what get stored in the database and sent over the wire.

The `Files` entity extends `managed`, which is a built-in CDS aspect from `@sap/cds/common`. This single word gives you four fields for free: `createdAt`, `createdBy`, `modifiedAt`, and `modifiedBy`. The framework populates them automatically on every insert and update. You never need to write that logic yourself.

The dual-status design — `status` for the Knowledge Graph pipeline and `vectorStatus` for the vector pipeline — reflects a real operational requirement. The two pipelines run independently. A document can have its triples extracted successfully while the vector embedding is still running, or vice versa. Tracking them separately means your UI can show accurate, granular progress rather than a single boolean that masks which step failed.

`fileHash` carries the `@assert.unique` annotation. This is a CAP-level constraint that generates a database unique index and returns a meaningful OData error if you attempt to upload the same file twice. Without this, operators could accidentally trigger redundant processing.

The `FileAttachments` entity extends `Attachments` from `@cap-js/attachments`. This is the entire attachment implementation — one line of CDS. The plugin handles binary storage, MIME type detection, streaming upload and download, and exposes the attachment content through the standard OData `$value` endpoint. The `up_` association links each attachment back to its parent `Files` record, following the standard CAP composition pattern.

The `attachments` field on `Files` is a `Composition` of many `FileAttachments`. `Compositions` in CDS mean deep operations: when you delete a `Files` record, all its attachments are deleted automatically. When you read a `Files` record with `$expand=attachments`, you get the attachment metadata in a single request.

9.4 The service definition

The service is defined in `srv/service.cds`. This file does three things: it projects the data model entities into the service, it adds virtual fields for UI display, and it declares the actions and functions that map to agent operations.

```

1  using { msds.kg as db } from '../db/schema';
2
3  @path: '/odata/v4/documents'
4  service DocumentService {
5      @odata.draft.enabled
6      entity Documents as projection on db.Files {
7          *,
8          virtual kgDisplay          : String,
9          virtual vecDisplay         : String,
10         virtual statusCriticality   : Integer,
11         virtual vectorStatusCriticality : Integer
12     }
13     actions {
14         action processFile() returns { status: String; message: String;
15         attachmentCount: Integer };
16         @(Common.IsActionCritical: true)
17         action deleteFile() returns { status: String; message: String };
18     };
19     action chatQuery(message: String, materialNumber: String) returns {
20     answer: String; path_used: many String };
21
22     function pollStatus(materialNumber: String) returns {
23         status    : String;
24         triples   : Integer;
25         vectors   : Integer;
26         kg_done   : Boolean;
27         vec_done  : Boolean;
28     };
29 }

```

The `@path` annotation sets the OData service root URL. All entity sets, actions, and functions under this service will be accessible at `/odata/v4/documents`. This path is regis-

tered by `@sap/cds` automatically when the server starts.

`@odata.draft.enabled` on the `Documents` entity activates SAP Fiori draft handling. Draft is the standard pattern for transactional data entry in Fiori applications. The framework saves partial edits server-side before the user explicitly confirms them — preventing data loss on accidental navigation and matching how SAP Fiori works for purchase orders, material master records, and every other transactional document in the SAP ecosystem. In practice, for this application it means the create-new-document flow works correctly in the Fiori UI without any additional handler code — users can enter the material number, attach the PDF, and save as draft before committing to processing.

The `virtual` fields deserve particular attention. A virtual field is projected into the OData response but has no corresponding database column. The values are computed in the service handler at read time. Here, `kgDisplay` and `vecDisplay` are formatted strings like “Completed (1,247 triples)” that combine the status and the count into a single display value. The `statusCriticality` and `vectorStatusCriticality` fields are integers (0–3) that Fiori Elements interprets as colour codes: 0 is neutral, 1 is error (red), 2 is warning (orange), 3 is positive (green). By computing criticality as a virtual field, you keep the colour logic in the service layer rather than scattering it across annotations.

The distinction between `action` and `function` in OData V4 is important. Actions have side effects — they modify state. Functions are read-only — they return computed data without modifying anything. `processFile` and `deleteFile` are bound actions: they operate on a specific `Documents` entity instance, identified by its key. `chatQuery` is an unbound action: it accepts parameters directly and operates at the service level, not on a specific entity. `pollStatus` is an unbound function: it retrieves current status from the backend without changing anything, making it safe to call repeatedly in a polling loop.

The `@(Common.IsActionCritical: true)` annotation on `deleteFile` tells Fiori Elements to show a confirmation dialog before invoking the action. The user must explicitly confirm the destructive operation. One annotation; no custom dialog code.

9.5 The Products entity and S/4HANA value help

The `DocumentService` extends beyond the `Files` entity. The service also exposes a `Products` entity sourced from the `API_PRODUCT_SRV` external service — the standard

SAP OData API for product master data in S/4HANA.

This is architecturally significant. The `materialNumber` field on the `Documents` entity carries a `@Common.ValueList` annotation that points to the `Products` entity. When a business user opens the Fiori UI to upload a new document and clicks into the Material Number field, the Fiori value help fires an OData request to `DocumentService/Products`. CAP proxies that request to the configured S/4HANA system and returns the matching product master records.

What this means operationally: a quality manager uploading a batch certificate does not type a material number from memory — they search and select from their actual SAP product master. The material number that anchors the document in the Knowledge Graph and vector store is the same material number that exists in their S/4HANA system. When a question comes back through `chatQuery` — “What are the storage conditions for this material?” — the answer is grounded in data that is tied to the real product record.

This is what makes this a genuine enterprise integration rather than a standalone demo. The AI capability on BTP is connected to the system of record in S/4HANA through standard SAP OData APIs and standard BTP Destination Service connectivity. Changing the S/4HANA destination in the BTP Cockpit is the only configuration needed to point this system at a different S/4HANA landscape.

9.6 The callBackend helper

Every handler in `srv/service.js` that needs to communicate with the Python agent uses the same `callBackend` helper function. This centralisation matters: the URL of the backend, the authentication headers, the error handling, and the timeout configuration all live in one place.

```
1  const BACKEND_URL = process.env.BACKEND_URL || "http://localhost:8000";
2
3  async function callBackend({ method, path, jsonBody, formData }) {
4    const response = await axios.request({
5      method,
6      url: `${BACKEND_URL}${path}`,
```



```

7     headers: formData ? formData.getHeaders() : { "Content-Type":
      "application/json" },
8     data: formData || jsonBody,
9   });
10   return response.data;
11 }

```

The `BACKEND_URL` environment variable is read once at module load time. In local development you set this to `http://localhost:8000` in a `.env` file in the `cap-srv/` directory, and `@sap/cds` loads it automatically via `dotenv`. On BTP, you set it as an environment variable in your Cloud Foundry application manifest or as a user-provided service binding. The fallback to `http://localhost:8000` means the service starts without any configuration during development — the fallback is safe because there is no production environment where port 8000 is the correct URL.

The `formData ? formData.getHeaders()` conditional is the key to handling two different call types with one function. When `formData` is provided, the request is a multi-part upload and `axios` must use the boundary-aware headers that `form-data` generates. When `jsonBody` is provided, the content type is `application/json`. Getting this wrong produces a 400 Bad Request from the FastAPI endpoint.

`Axios` is used here rather than the native Node.js `fetch` because `form-data` integrates with `axios` through the `getHeaders()` and pipe mechanism. Node.js 18+ `fetch` supports `FormData`, but the `@cap-js/attachments` binary retrieval returns a Node.js `Buffer`, and converting that buffer to a browser-compatible `FormData` in Node.js requires additional tooling. `Axios` with the `form-data` package is the simpler path.

9.7 Action handlers

processFile

When a user clicks Process File in the Fiori UI, they are initiating the full ingestion pipeline. The Fiori button fires an OData action request to the CAP service. CAP reads the PDF from the attachment store, constructs a multipart request, and POSTs it to the FastAPI backend at `/process-upload`. FastAPI returns 202 Accepted immediately — it has queued the job but not finished it. CAP updates the `Files` record to `Processing` and

returns to the Fiori UI. The Fiori UI then begins polling `pollStatus()` on a timer. As the FastAPI pipeline runs — chunking the PDF, generating triples, embedding passages — the triples and vectors counts climb until both pipelines report completion. The Fiori UI reflects those counts as they update, giving the operator a live view of ingestion progress without any WebSocket infrastructure.

The `processFile` action is the most complex handler because it bridges two protocols: OData binary retrieval and HTTP multipart upload. The handler must read the attachment binary from the CAP attachment store, then stream it to the FastAPI `/process-upload` endpoint as a multipart request.

```
1  srv.on("processFile", Documents, async (req) => {
2    const { materialNumber } = req.params[0];
3
4    // Retrieve attachment metadata from the CAP entity
5    const attachments = await SELECT.from(db.Files, materialNumber)
6      .columns("attachments");
7
8    if (!attachments || attachments.attachments.length === 0) {
9      return req.error(400, "No attachment found. Please upload a PDF
10       first.");
11    }
12
13    const attachment = attachments.attachments[0];
14
15    // Read the binary content through the @cap-js/attachments API
16    const content = await cds.run(SELECT.one.from(db.FileAttachments)
17      .where({ ID: attachment.ID }));
18
19    const buffer = content.content; // Buffer from the attachment plugin
20    const filename = attachment.filename || `${materialNumber}.pdf`;
21    const contentType = attachment.mimeType || "application/pdf";
22
23    // Build multipart form data for the FastAPI endpoint
24    const FormData = require("form-data");
25    const formData = new FormData();
26    formData.append("file", buffer, { filename, contentType });
27    formData.append("graphName", materialNumber);
```

```

28   await callBackend({ method: "POST", path: "/process-upload", formData
    });
29
30   // Mark the record as processing
31   await UPDATE(db.Files).set({
32     status: "Processing",
33     vectorStatus: "Processing",
34     processingStartedAt: new Date().toISOString()
35   }).where({ materialNumber });
36
37   return {
38     status: "processing",
39     message: "File submitted for processing.",
40     attachmentCount: attachments.attachments.length
41   };
42 });

```

The `req.params[0]` contains the entity key of the bound action target — in this case `{ materialNumber: 'BATCH-QC-MAT-001' }`. This is the OData V4 pattern for bound actions: the key is part of the action URL path, not the request body.

The handler updates the record status to `Processing` immediately after submitting to the backend. This is an optimistic update — it reflects what should happen, not what has completed. The actual status progression (`Processing` to `Completed` or `Error`) happens asynchronously as the Python agent runs. The `pollStatus` function and handler described below are how the UI tracks that progression.

chatQuery

The `chatQuery` action is the core business value of this entire system. A quality manager, safety officer, or procurement analyst opens the document object page, types a natural language question — “What is the recommended storage temperature for this material?” or “Are there any regulatory restrictions on transport by air?” — and gets a sourced answer directly in the Fiori UI. The retrieval path shown alongside the answer tells them whether it came from the structured Knowledge Graph, the vector store, or both. No custom search interface. No manual document reading. No specialist tool to learn.

```

1  srv.on("chatQuery", async (req) => {

```

```

2   const { message, materialNumber } = req.data;
3
4   const result = await callBackend({
5     method: "POST",
6     path: "/query",
7     jsonBody: {
8       question: message,
9       material_number: materialNumber,
10      chat_history: []
11    }
12  });
13
14  return {
15    answer: result.answer,
16    path_used: result.path_used || []
17  };
18  });

```

The `chat_history: []` field sends an empty history on every call. This makes every `chatQuery` invocation stateless from the agent's perspective. For a document Q&A system where questions are typically independent lookups, stateless queries are correct. If you were building a multi-turn conversational interface, you would maintain a conversation ID in the client and accumulate history on the server side. That is a separate concern — `chatQuery` can be extended to accept a `sessionId` parameter without changing anything in the action definition or the backend API contract.

The `path_used` return value is an array of strings that the Python agent populates with the retrieval path it took: `["kg_chain"]`, `["vector_chain"]`, or `["kg_chain", "vector_chain"]`. Surfacing this in the UI gives users transparency into the agent's reasoning — they can see whether the answer came from the Knowledge Graph, the vector store, or both.

deleteFile

The `deleteFile` handler triggers a cascading cleanup: remove the Knowledge Graph triples from HANA, remove the vector embeddings from HANA, then delete the CAP record and its attachments.

```

1  srv.on("deleteFile", Documents, async (req) => {
2    const { materialNumber } = req.params[0];
3
4    // Instruct the Python backend to clean up graph and vector data
5    await callBackend({
6      method: "DELETE",
7      path: `/files/${encodeURIComponent(materialNumber)}`
8    });
9
10   // Delete the CAP record — composition cascades to FileAttachments
11   await DELETE.from(db.Files).where({ materialNumber });
12
13   return { status: "deleted", message: `${materialNumber} removed.` };
14 });

```

The `DELETE.from(db.Files)` call cascades automatically to the `FileAttachments` composition because CDS compositions enforce deep delete semantics. The `@cap-js/attachments` plugin hooks into the delete event and removes the binary files from storage. You do not need to explicitly delete the attachments — the composition and the plugin handle it.

pollStatus

The `pollStatus` function is called by the Fiori UI on a timer after a `processFile` action to track progress. It calls the FastAPI `/status/{materialNumber}` endpoint and passes the response through.

```

1  srv.on("pollStatus", async (req) => {
2    const { materialNumber } = req.data;
3
4    const result = await callBackend({
5      method: "GET",
6      path: `/status/${encodeURIComponent(materialNumber)}`
7    });
8
9    // Sync the live status back to the CAP entity
10   if (result.kg_done || result.vec_done) {
11     await UPDATE(db.Files).set({

```

```

12     status: result.kg_done ? "Completed" : result.status,
13     vectorStatus: result.vec_done ? "Completed" : result.vectorStatus,
14     triples: result.triples || 0,
15     vectors: result.vectors || 0
16   }).where({ materialNumber });
17 }
18
19 return result;
20 });

```

The handler does two things: it returns the live status to the caller, and it writes the current counts back to the database. This write-through is important for the Fiori list report: the list refreshes from the OData entity, not from a live WebSocket feed. By syncing the counts on every poll, the list view stays accurate after polling completes.

9.8 Fiori Elements UI

Fiori Elements is a metadata-driven UI framework. Instead of writing a React or SAPUI5 application that fetches OData and renders HTML, you write CDS annotations that describe *what* your data means and *how* you want it presented. The framework reads the OData metadata document (which CDS generates from your service definition) and the annotation document, and assembles the UI at runtime.

This has a profound implication for development speed. A complete list report with search, sorting, and export requires approximately 50 lines of CDS annotations. The equivalent hand-written SAPUI5 application is around 500 lines of XML views and controllers, plus another few hundred lines of JavaScript. For internal tools and prototypes, Fiori Elements is consistently the better choice.

The trade-off is customisation ceiling. If you need a completely custom layout — a drag-and-drop canvas, a real-time chart with WebSocket data, a custom color scheme — Fiori Elements will fight you. For standard enterprise UI patterns — list reports, object pages, form inputs, action buttons — it covers everything you need.

For our document management and Q&A interface, the standard patterns are exactly right:

- A **list report** showing all documents with their status, triple count, and upload

date

- An **object page** for each document showing its details, attachments, and action buttons
- A **chat section** on the object page where users can submit questions and see answers

All of this is generated from the CDS service definition and the annotations in `srv/annotations.cds`. There is no separate UI project, no `package.json` for the front-end, and no custom JavaScript.

The CAP + Fiori layer means this AI capability fits natively into any SAP Fiori launchpad. A quality manager, safety officer, or procurement analyst uses the same Fiori UX patterns they use across every other SAP application. The AI is invisible infrastructure — the user just asks a question and gets an answer.

9.9 Annotations explained

The annotations file `srv/annotations.cds` is where the UI is configured. It is separate from the service definition for a practical reason: the service definition describes what your API *does*, and the annotations describe how it *looks*. Keeping them separate means you can modify the UI without changing the API contract.

HeaderInfo

```

1  annotate service.Documents with @(
2      UI.HeaderInfo: {
3          TypeName: 'Document',
4          TypeNamePlural: 'Documents',
5          Title: { Value: materialNumber },
6          Description: { Value: status }
7      }
8  );

```

HeaderInfo controls three things: the singular and plural names used in page titles and breadcrumbs, the main title shown on the object page header, and the subtitle beneath it. Setting `Title` to `materialNumber` means each document's object page is headed by its material number. Setting `Description` to `status` shows the current processing status

directly in the page header, so the user knows at a glance whether the document has been processed.

SelectionFields and LineItem

```

1      UI.SelectionFields: [ materialNumber, status ],
2      UI.LineItem: [
3          { Value: materialNumber, Label: 'Material Number' },
4          { Value: kgDisplay, Label: 'KG Status', Criticality:
statusCriticality },
5          { Value: vecDisplay, Label: 'Vec Status', Criticality:
vectorStatusCriticality },
6          { Value: createdAt, Label: 'Uploaded On' }
7      ],

```

`SelectionFields` defines which fields appear in the filter bar above the list. Choosing `materialNumber` and `status` means users can search for a specific material or filter to show only documents in a particular state.

`LineItem` defines the columns in the list table. The `kgDisplay` and `vecDisplay` fields carry the `Criticality` annotation pointing to the integer fields computed in the service handler. Fiori Elements renders criticality as a colored status indicator: integers map to red, orange, and green automatically. The display strings computed in the handler — “Completed (1,247 triples)” — appear in the column alongside the color indicator. No custom cell renderer needed.

Identification actions

```

1      UI.Identification: [
2          {
3              $Type: 'UI.DataFieldForAction',
4              Action: 'DocumentService.processFile',
5              Label: 'Process File',
6              Criticality: #Positive
7          },
8          {
9              $Type: 'UI.DataFieldForAction',
10             Action: 'DocumentService.deleteFile',

```



```

11         Label: 'Delete File',
12         Criticality: #Negative
13     }
14 ],

```

The `Identification` collection defines the action buttons that appear in the object page header toolbar. `UI.DataFieldForAction` binds a button to an OData action by its qualified name. The `Criticality` annotation sets the button color: `#Positive` renders as green (constructive action), `#Negative` renders as red (destructive action). Combined with the `Common.IsActionCritical` annotation on `deleteFile` in the service definition, the delete button is both red and protected by a confirmation dialog — two layers of friction for a destructive operation.

FieldGroup

```

1     UI.FieldGroup#GeneralInfo: {
2         Data: [
3             { Value: materialNumber },
4             { Value: status },
5             { Value: triples, Label: 'Triples Extracted' },
6             { Value: vectors, Label: 'Vectors Stored' },
7             { Value: vectorStatus, Label: 'Vector Status' }
8         ]
9     }

```

`FieldGroup` with an identifier (`#GeneralInfo`) defines a named group of fields. These groups are then referenced in a `Facets` annotation to assemble the object page layout. Each `Data` entry becomes a labeled field in the form. The `triples` and `vectors` integers show the current counts from the last successful processing run, giving operators the information they need to assess pipeline health.

9.10 Running locally

The local setup requires two terminal sessions — one for each process.

Terminal 1: Start the Python FastAPI agent

```
1 cd agents/  
2 source .venv/bin/activate  
3 uvicorn main:app --reload --port 8000
```

Confirm it is running by visiting <http://localhost:8000/docs> and checking that the `/query`, `/process-upload`, and `/status/{material_number}` endpoints are listed.

Terminal 2: Start the CAP service

First, create a `.env` file in `cap-srv/` if one does not exist:

```
BACKEND_URL=http://localhost:8000
```

Then start the server:

```
1 cd cap-srv/  
2 npm install  
3 cds serve
```

The `cds serve` command reads your service definition, sets up the SQLite database (for local development — replace with HANA on BTP), loads all handlers, and starts the Express server on port 4004. You will see output like:

```
[cds] - loaded model from 3 file(s):
```

```
db/schema.cds
```

```
srv/service.cds
```

```
srv/annotations.cds
```

```
[cds] - connect to db > sqlite { database: ':memory:' }
```

```
[cds] - serving DocumentService { at: '/odata/v4/documents' }
```

```
[cds] - server listening on { url: 'http://localhost:4004' }
```

Open <http://localhost:4004> to see the CDS welcome page, which lists all registered services and their metadata endpoints. Open [http://localhost:4004/\\$fiori-preview](http://localhost:4004/$fiori-preview) to access the Fiori Elements preview sandbox — a full Fiori application running against your local OData service with no deployment required.

For the Fiori Elements preview to show all features, ensure you access it via the full URL including the Fiori launchpad hash, which the welcome page provides as a clickable link.

9.11 Testing the full flow

With both services running, walk through the complete flow from document upload to agent query.

Step 1: Create a new document record. In the Fiori list report, click the New button. The draft mechanism activates. Click into the Material Number field — a value help opens, backed by the Products entity from API_PRODUCT_SRV. Search for and select your material. Save the draft. The record appears in the list with status Pending.

Step 2: Upload a PDF attachment. Navigate to the object page for your material. In the Attachments section, click Upload and select a PDF — an MSDS, SDS, or any document relevant to that material. The attachment is stored immediately via @cap-js/attachments.

Step 3: Process the file. Click the Process File button in the header toolbar. The handler reads the attachment binary, constructs a multipart form data request, and POSTs it to the Python agent at `http://localhost:8000/process-upload`. The record status changes to Processing immediately.

Step 4: Watch the dual status. Open a separate terminal and call `pollStatus` manually to watch the counts grow:

```
1 curl
  "http://localhost:4004/odata/v4/documents/pollStatus(materialNumber='BATCH-
  QC-MAT-001')"
```

You will see `kg_done` and `vec_done` flip to true as each pipeline completes, and the `triples` and `vectors` counts increment. The Fiori object page, if you have it open and refresh it, shows the same values updating.

Step 5: Run a query. Once both pipelines are complete, use the `chatQuery` action. In the Fiori UI this can be triggered via an action dialog bound to the `chatQuery` unbound action. From the command line:

```
1 curl -X POST http://localhost:4004/odata/v4/documents/chatQuery \
2   -H "Content-Type: application/json" \
```

```
3 -d '{"message": "What test method was used and what were the tensile strength results?", "materialNumber": "BATCH-QC-MAT-001"}'
```

The response includes the `answer` field with the agent's response and the `path_used` array indicating which retrieval chains contributed: `["kg_chain", "vector_chain"]` means both paths were used, which is the expected behavior for this type of factual question about a material document.

9.12 What you can do after this chapter

The system is now complete from ingestion to query to UI. You have:

- A Python FastAPI service that runs the hybrid RAG agent with parallel KG and vector retrieval
- A CAP OData V4 service that manages document records, attachment storage, and proxies requests to the agent
- A Fiori Elements UI generated entirely from CDS annotations
- A value help on Material Number backed by real S/4HANA product master data via `API_PRODUCT_SRV`

The natural next steps are:

Add conversational history to chatQuery. Extend the `chatQuery` action to accept an optional `sessionId` parameter. Store conversation turns in a new CDS entity `ChatSession`. On each call, retrieve the session history and pass it to the FastAPI `/query` endpoint as `chat_history`. This turns a single-shot Q&A into a stateful conversation without changing the agent logic.

Add a polling timer to the Fiori UI. Write a Fiori Elements custom section that calls `pollStatus` every 5 seconds while the document status is `Processing`, updating the form fields and stopping the timer when both pipelines are complete. This is one of the few cases where a small amount of custom SAPUI5 controller code in the Fiori application is justified.

Extend to additional document types. The schema is built for this. Add a `documentType` field to `Files`, extend the `Status` enum if needed, and configure the FastAPI ingestion pipeline to handle invoice layouts, batch certificate formats, or quality inspection report structures. The CAP layer, the Fiori UI, and the `chatQuery` interface require

no changes.

Deploy to SAP BTP Cloud Foundry. Replace the SQLite database with HANA Cloud by setting the `cds.requires.db.kind` to `hana` and binding the `hana` service. Set `BACK-END_URL` to the deployed FastAPI service URL. Run `cf push` from `cap-srv/`. The CAP service connects to HANA automatically via the `VCAP_SERVICES` binding. Chapter 10 covers this in full.

Add XSUAA authentication. Bind an XSUAA service instance to the CAP application. Add `@requires: 'authenticated-user'` to the `DocumentService` definition. The CAP framework validates JWT tokens from XSUAA on every request with no additional handler code.

9.13 Checkpoint checklist

Before moving to the next chapter, verify the following:

- `cds serve` starts without errors and outputs the service URL at port 4004
- The Fiori Elements preview is accessible at `http://localhost:4004/$fiori-preview`
- A new document can be created and saved via the Fiori list report
- The Material Number value help returns results from the Products entity
- A PDF attachment can be uploaded and is stored by `@cap-js/attachments`
- Clicking Process File sends a multipart POST to the FastAPI service and returns `status: "processing"`
- The `pollStatus` function returns accurate `triples` and `vectors` counts after processing completes
- The `chatQuery` action returns an `answer` string and a `path_used` array
- The `deleteFile` action triggers the confirmation dialog (from `Common.IsActionCritical`) before executing
- The `kgDisplay` and `vecDisplay` columns in the list report show colored status indicators (green for Completed, red for Error)

If all ten items pass, your SAP-native API layer and Fiori Elements UI are working correctly. The full system — Python agent, CAP service, and Fiori UI — is operational end to end.

Summary

This chapter added the enterprise presentation layer to the hybrid RAG system. We built a CAP Node.js OData V4 service that manages document records and attachments through CDS entities, and exposes file processing and query operations through bound and unbound OData actions. The `materialNumber` field carries a value help backed by real S/4HANA product master data through `API_PRODUCT_SRV` — the connection between the AI capability on BTP and the system of record in S/4HANA. We built a `callBackend` helper that proxies all agent operations to the Python FastAPI service via HTTP, keeping the two processes completely decoupled. We generated a complete Fiori Elements UI — list report with colored status columns, object page with action buttons, and attachment handling — from CDS annotations alone, with no custom JavaScript or HTML.

The architectural principle throughout is separation of concerns: CAP owns the OData protocol, the S/4HANA integration, and the UI; FastAPI owns the intelligence. Neither service knows the internals of the other. They communicate through a narrow, stable HTTP contract. This separation makes both services easier to test, easier to scale, and easier to evolve independently.

The CAP + Fiori layer means this AI capability fits natively into any SAP Fiori launchpad. The user just asks a question and gets an answer.

Chapter 10: Deploying to SAP BTP Cloud Foundry

By the time you reach this chapter, you have a working Hybrid RAG system running locally against your BTP HANA Cloud trial. This chapter deploys it — unchanged code — to BTP Cloud Foundry. The FastAPI Python application becomes a CF app. The CAP Node.js service becomes an MTA module. HANA Cloud, already provisioned in Chapter 2, becomes the bound service.

Every configuration decision in this chapter has a specific rationale. Ports are dynamic because CF assigns them at container startup. Credentials come from service bindings because files cannot be trusted in a containerised environment. The two applications communicate over CF routes rather than localhost because they are separate containers on a distributed platform. Understanding these rationales means you can adapt the deployment pattern to other BTP applications, not just this one.

10.1 What Changes Between Local and BTP

Before touching a single deployment file, it is worth being precise about what is actually different in a Cloud Foundry environment. There are three categories of change, and understanding all three prevents most deployment failures.

Ports are dynamic. On your laptop, you choose the port. You run FastAPI on 8000 and CAP on 4004, and nothing conflicts. In Cloud Foundry, the platform assigns a port at container startup and communicates it through the `$PORT` environment variable. Your application must read that variable and bind to it. If it binds to a hardcoded port, it will start and immediately fail health checks. This is why the `manifest.yml` command is `uvicorn main:app --host 0.0.0.0 --port $PORT` rather than a fixed port number.

Credentials come from service bindings, not files. A `.env` file is a local development convenience. It is not a deployment artifact. On BTP, sensitive values — database credentials, API keys, GCP service account keys — are delivered to running containers through Cloud Foundry service bindings injected as JSON into the `VCAP_SERVICES` environment variable. The `GOOGLE_API_KEY` that worked locally never appears in a deployed

container's filesystem or in `manifest.yml`. This is not just a best practice; it is how the platform is designed. The BTP Destination Service stores the Vertex AI credentials at runtime — the CAP service reads them from the Destination Service, never from environment variables visible in `cf env`.

Inter-service communication uses the platform's routing layer. On localhost, CAP reaches the Python agent at `http://localhost:8000`. On BTP Cloud Foundry, every application has a route — a public HTTPS URL managed by the platform's router. The two applications communicate over those routes. The BTP Destination Service provides a managed, centralised way to store and reference those URLs, so that changing the back-end URL only requires updating one entry in the BTP Cockpit rather than redeploying every service that references it.

These three changes — dynamic ports, credential delivery via bindings, and HTTPS routing — account for almost all of the configuration work in this chapter.

10.2 Why Two Separate Applications

The system deploys as two independent CF applications: the FastAPI Python service and the CAP Node.js service. This is not an accident of how the code happens to be organised — it is a deliberate architectural choice with operational consequences.

FastAPI handles AI orchestration and HANA data access. It is a standard Python web service that runs on any CF Python buildpack. Its responsibilities are: receiving document upload requests, running the LangGraph ingestion pipeline, querying HANA for vector and Knowledge Graph retrieval, calling Vertex AI for LLM inference, and returning answers. None of these responsibilities have any dependency on OData, Fiori, or BTP security services.

CAP handles OData protocol, Fiori UI rendering, S/4HANA product master integration via `API_PRODUCT_SRV`, and BTP service bindings for XSUAA authentication and HANA schema management. It is the SAP-native layer. Its responsibilities are: serving the Fiori Elements UI, validating OData requests, managing document and attachment records in HANA, and proxying processing and query requests to the FastAPI backend.

Keeping them separate means you can update the AI model, the LangGraph orchestration logic, or the vector retrieval strategy without touching the CAP service — and vice

versa. The HTTP contract between them (multipart POST to `/process-upload`, JSON POST to `/query`, GET to `/status/{materialNumber}`) is stable and narrow. Either side can be redeployed independently as long as that contract is respected.

10.3 Pre-Deployment Checklist

Deployment failures are much easier to debug when you can be certain your local environment is correctly configured before you start. Work through this checklist before executing any `cf push` or `mbt build` command.

Cloud Foundry CLI. Run `cf version`. You need version 8 or later. If the command is not found, install from the Cloud Foundry Foundation releases page or, on macOS, with `brew install cloudfoundry/tap/cf-cli@8`. Log in with `cf login -a https://api.cf.<your-region>.hana.ondemand.com` and target the correct org and space with `cf target -o <org> -s <space>`.

MTA Build Tool. The CAP service is deployed using the MultiTarget Application (MTA) model, which requires the `mbt` CLI. Run `mbt --version`. If it is missing, install with `npm install -g mbt`. This tool reads `mta.yaml`, builds all modules, and packages them into a single `.mtar` archive that the CF deployer can process.

BTP Cloud Foundry space. Confirm you are in the correct space with `cf target`. The space must have the SAP HANA Cloud instance accessible to it. If you are using a trial account, the HANA Cloud instance must be active — it stops automatically after 30 days of inactivity and must be restarted manually in the BTP Cockpit.

HANA Cloud instance running. In the BTP Cockpit, navigate to your subaccount, open the SAP HANA Cloud section, and confirm the instance status is “Running.” Deployment will succeed even if HANA is stopped — the CAP service only attempts a database connection at runtime. But the smoke test in section 10.8 will fail until HANA is available.

Google Cloud service account credentials. The Python agent authenticates to Vertex AI using a Google Cloud service account. You should have a JSON key file from Chapter 2. For BTP deployment, this key is stored in the BTP Destination Service — never as a file in the container, never as a plain environment variable visible in `cf env`.

10.4 Step 1: Push the Python Agent

The Python agent deploys first because the CAP service needs to know its URL when configuring the Destination Service. Navigate to the `agents/` directory and examine the deployment manifest.

Understanding manifest.yml

```

1
2  ---
3  applications:
4    - name: hybrid-rag-agent
5      buildpacks:
6        - python_buildpack
7      memory: 512M
8      disk_quota: 1G
9      instances: 1
10     command: uvicorn main:app --host 0.0.0.0 --port $PORT
11     path: .
12     env:
13       HANA_HOST: ((hana-host))
14       HANA_PORT: "443"
15       HANA_USER: ((hana-user))
16       HANA_PASSWORD: ((hana-password))
17       GCP_PROJECT_ID: ((gcp-project-id))
18       GCP_LOCATION: us-central1
19       LANGCHAIN_TRACING_V2: "true"
20       LANGCHAIN_ENDPOINT: "https://api.smith.langchain.com"
21       LANGCHAIN_API_KEY: ((langchain-api-key))
22       PYTHONPATH: "."

```

Walk through each field:

`name: hybrid-rag-agent` sets the application name in Cloud Foundry. This becomes part of the default route: `hybrid-rag-agent.<your-cf-domain>`. Choose a name that is unique within your org.

`buildpacks: [python_buildpack]` tells Cloud Foundry which buildpack to use for detecting, compiling, and packaging the application. The Python buildpack reads `requirements.txt`, creates a `virtualenv`, installs dependencies, and configures the runtime. No

Dockerfile is needed.

`memory: 512M` and `disk_quota: 1G` set the container resource limits. 512 MB is sufficient for the FastAPI server and the LangGraph agent under moderate load. The HANA Python driver (`hdbcli`), Vertex AI client libraries, and LangChain components together consume roughly 200-250 MB at startup. If you observe OOM kills in logs, increase to 1G.

`command: uvicorn main:app --host 0.0.0.0 --port $PORT` is the startup command. The `--host 0.0.0.0` binding is required — binding only to `127.0.0.1` means the platform's router cannot reach the container. The `$PORT` variable is provided by Cloud Foundry and will be a value like 8080 or 61002 depending on the container assignment. The Python buildpack expands this shell variable before executing the command.

`path: .` tells the CF CLI which directory to upload. Since you are running `cf push` from within `agents/`, the `.` means the entire `agents` directory is packaged and sent to the staging container.

`PYTHONPATH: "."` ensures that Python can find modules in the application root. Without this, relative imports within the agent codebase will fail at runtime.

The `((variable))` syntax in the `env` block are placeholders — the real values are supplied through `cf set-env` commands or a User-Provided Service, covered in the next step. If you pushed this manifest as-is, the literal placeholder strings would appear as environment variable values, causing immediate connection failures at runtime.

HANA connection in CF

In the deployed environment, HANA credentials come from `VCAP_SERVICES` via service binding — the same mechanism CAP uses. The `hdb_srv.py` module already handles this: it reads `VCAP_SERVICES` if direct environment variables are not set. No code change is needed to move from local `.env` file credentials to CF service binding credentials. The credential source changes; the application code does not.

Pushing the application

From the `agents/` directory:

```
1 cf push hybrid-rag-agent --no-start
```

The `--no-start` flag stages the application — uploads the code, runs the buildpack,

creates the container — but does not start it. This is important because the environment variables have not been set yet. Starting now would produce a crash-loop of failed HANA connections.

Confirm staging succeeded:

```
1 cf app hybrid-rag-agent
```

You should see the application in a `stopped` state with a route assigned. Note the route URL — something like `hybrid-rag-agent.cfapps.eu10.hana.ondemand.com`. You will need this URL when configuring the Destination Service.

10.5 Step 2: Create the User-Provided Service for Credentials

Cloud Foundry User-Provided Services (UPS) are the correct mechanism for delivering credentials to applications that need values which are not available through a managed service binding. Instead of setting individual environment variables one by one on each application, you define a named service with a JSON payload of key-value pairs, bind it to any application that needs those values, and Cloud Foundry delivers the entire payload via `VCAP_SERVICES`.

The security rationale is straightforward. If you set credentials with `cf set-env`, they are visible in `cf env <app>` to anyone with developer access to the space. If you put credentials in `manifest.yml`, they will eventually end up in version control. A User-Provided Service stores the values in Cloud Foundry's internal credential store, not in any file that exists in your repository.

Create the service with the actual values for your environment:

```
1 cf cups hybrid-rag-agent-env -p '{
2   "HANA_HOST": "your-hana-instance.hanacloud.ondemand.com",
3   "HANA_PORT": "443",
4   "HANA_USER": "DBADMIN",
5   "HANA_PASSWORD": "your-database-password",
6   "GCP_PROJECT_ID": "your-gcp-project-id",
```

```
7   "GCP_LOCATION": "us-central1",
8   "LANGCHAIN_API_KEY": "your-langchain-api-key"
9   }'
```

`cf cups` stands for `cf create-user-provided-service`. The `-p` flag accepts a JSON string of parameters.

Now bind the service to the application:

```
1  cf bind-service hybrid-rag-agent hybrid-rag-agent-env
```

Cloud Foundry will inject the key-value pairs from this service into the `VCAP_SERVICES` JSON available to the container. The `HANA_HOST` key in the UPS becomes the `HANA_HOST` environment variable in the application container.

If you need to update the credentials later — for example, to rotate the HANA password — you can use:

```
1  cf uups hybrid-rag-agent-env -p '{"HANA_PASSWORD": "new-password"}'
2  cf restart hybrid-rag-agent
```

The `cf uups` command (update-user-provided-service) replaces the entire parameter set, so include all keys, not just the changed one.

Now start the application:

```
1  cf start hybrid-rag-agent
```

Watch the startup logs:

```
1  cf logs hybrid-rag-agent --recent
```

A successful startup will show `uvicorn` reporting that it is listening on the assigned port. A failed startup will show a Python traceback — usually a missing environment variable or an import error from a missing package. The most common first failure is a missing `google-cloud-aiplatform` package because it was not in `requirements.txt`. Confirm your `requirements.txt` includes every library imported anywhere in the codebase.

Verify the agent is reachable:

```
1 curl https://hybrid-rag-agent.<your-cf-domain>/health
```

You should receive a JSON response indicating the service status and confirming that the HANA and Vertex AI connections are available.

10.6 Step 3: Build and Deploy the CAP Service

The CAP service is deployed using the MTA (MultiTarget Application) model rather than a direct `cf push`. MTA is an SAP-developed open standard for describing multi-component applications — a single `mta.yaml` file describes every module (CAP service, Python service), every service binding (HANA, UPS), and every dependency between them. The `mbt build` tool compiles all modules and packages them, and the CF MTA deployer (`cf deploy`) orchestrates the deployment.

Examine the MTA descriptor:

```
1 _schema-version: "3.3.0"
2 ID: msds-hybrid-rag
3 version: 1.0.0
4 description: Hybrid RAG MSDS system on SAP BTP
5
6 modules:
7   - name: hybrid-rag-agent
8     type: python
9     path: agents
10    build-parameters:
11      builder: custom
12      commands:
13        - pip install -r requirements.txt
14    parameters:
15      memory: 512M
16      disk-quota: 1G
17    properties:
18      PYTHONPATH: "."
19    requires:
20      - name: hybrid-rag-agent-env
```

```

21
22   - name: msds-hybrid-rag-cap
23     type: nodejs
24     path: cap-srv
25     build-parameters:
26       builder: npm
27       build-result: gen/srv
28     parameters:
29       memory: 256M
30       disk-quota: 512M
31     properties:
32       AGENT_URL: "~{hybrid-rag-agent/url}"
33     requires:
34       - name: msds-hybrid-rag-hana
35       - name: hybrid-rag-agent-env
36     provides:
37       - name: cap-srv-api
38     properties:
39       url: ${default-url}
40
41   resources:
42     - name: msds-hybrid-rag-hana
43       type: com.sap.xs.hana-schema
44       parameters:
45         service: hana
46         service-plan: schema
47
48     - name: hybrid-rag-agent-env
49       type: org.cloudfoundry.user-provided-service

```

Several lines require explanation.

`AGENT_URL: "~{hybrid-rag-agent/url}"` is an MTA cross-reference. The `~{}` syntax means “substitute the value of this property from another module at deployment time.” The `hybrid-rag-agent` module, when deployed, exposes its assigned route URL as the `url` property. The CAP module reads that URL and stores it in its own `AGENT_URL` environment variable. This eliminates hardcoding: even if the Python service is redeployed to a different route, the CAP service will be updated automatically on the next deployment.

`build-result: gen/srv` tells the MTA deployer where to find the compiled CAP artifacts

after the `npm build`. The `cds build` command, which runs as part of `npm run build`, compiles CDS models, generates the OData metadata, and writes the production-ready service files to `gen/srv`. The MTA deployer packages only this directory, not the entire `cap-srv` source tree.

`type: com.sap.xs.hana-schema` declares that the deployer should create (or reuse) a HANA schema resource. This becomes a service binding in Cloud Foundry: the CAP service's `VCAP_SERVICES` will contain the HANA connection details, and the `@cap-js/hana` plugin will read them automatically. You do not configure a connection string in the CAP application; the binding provides it.

`type: org.cloudfoundry.user-provided-service` in the resources section tells the MTA deployer that `hybrid-rag-agent-env` is an existing UPS that should be bound but not created. The UPS you created in step 10.5 is referenced here by name. Both the Python agent and the CAP service are bound to it, which is how both processes receive the HANA credentials.

Running the build

From the `cap-srv/` directory (or the repo root, depending on where `mta.yaml` lives):

```
1 mbt build
```

This command reads `mta.yaml`, executes the build commands for each module, and produces an `.mtar` archive in the `mta_archives/` directory. The archive is a standard ZIP containing all compiled artifacts for every module.

If the build fails at the `cds build` step, the most common cause is a CDS syntax error or a missing npm dependency. Run `npm run build` in the `cap-srv/` directory directly to see the full CDS compiler output.

Deploying the MTA

Install the CF MTA plugin if you have not already:

```
1 cf install-plugin multiapps
```

Deploy the built archive:


```
1 cf deploy mta_archives/msds-hybrid-rag_1.0.0.mtar
```

The deployer will work through a sequence of operations: creating or updating service instances, uploading application binaries, binding services, and starting applications. This takes two to five minutes. The terminal output shows each step with a status indicator. If any step fails, the deployer stops and shows the error. Common failures at this stage are HANA schema creation failures (usually a quota issue on a trial account) and npm install failures caused by package resolution errors.

When the deployment completes, both applications will be running. Confirm with:

```
1 cf apps
```

You should see `hybrid-rag-agent` and `msds-hybrid-rag-cap` both in a `started` state with routes assigned.

10.7 Step 4: Configure the BTP Destination for Vertex AI

In local development, Vertex AI credentials live in your `.env` file as `GOOGLE_APPLICATION_CREDENTIALS` pointing to a local JSON key file. That pattern does not belong in a deployed environment — a JSON key file in a container filesystem is one restart away from being lost, and it is visible to anyone with container access.

In BTP Cloud Foundry, the Vertex AI credentials are stored in the BTP Destination Service. The Destination Service is the enterprise credential management pattern across the entire BTP portfolio — the same mechanism used for S/4HANA connectivity, external REST APIs, and OAuth token exchange. The CAP service reads the Vertex AI credentials from the Destination Service at runtime. They are never stored in the app container, never in environment variables visible in `cf env`, and never in any file that could be accidentally committed or shared.

The production pattern in the CAP service handler, instead of reading from environment variables directly, uses the SAP Cloud SDK's destination resolution:

```
1 const { executeHttpRequest } = require('@sap-cloud-sdk/http-client');
2 const response = await executeHttpRequest(
```

```

3   { destinationName: process.env.BACKEND_DESTINATION },
4   { method: 'POST', url: '/query', data: body }
5   );

```

The `executeHttpRequest` function handles destination resolution, certificate verification, and — when configured — token propagation. This is the mature BTP application pattern. For the current deployment, we configure the Destination Service so it is available when you move to production hardening in a later chapter.

Creating the agent backend destination

Navigate to your BTP subaccount in the BTP Cockpit. Select **Connectivity** from the left navigation, then **Destinations**. Click **New Destination** and fill in the following values:

- **Name:** hybrid-rag-agent
- **Type:** HTTP
- **URL:** `https://hybrid-rag-agent.<your-cf-domain>` (the route from step 10.4)
- **Proxy Type:** Internet
- **Authentication:** NoAuthentication

Save the destination. You can test it immediately by clicking **Check Connection** — the cockpit will send a GET request to the root URL of the agent and report the HTTP response code.

Creating the Vertex AI destination

Click **New Destination** again and configure the Vertex AI credentials:

- **Name:** VertexAI
- **Type:** HTTP
- **URL:** `https://us-central1-aiplatform.googleapis.com`
- **Proxy Type:** Internet
- **Authentication:** NoAuthentication

Add additional properties:

Property	Value
<code>gcp.project_id</code>	Your GCP project ID
<code>gcp.location</code>	<code>us-central1</code>

Property	Value
<code>gcp.service_account_key</code>	The entire contents of your <code>gcp-sa-key.json</code> file

BTP encrypts the `gcp.service_account_key` value at rest. The service account JSON key is now stored in one place — the BTP Destination Service — and accessible to any BTP application in this subaccount that needs Vertex AI access. Rotating the credentials means updating this one destination entry, not redeploying any application.

10.8 Step 5: Load the Ontology

The HANA Knowledge Graph requires the MSDS ontology to be loaded before any queries can run against it. In local development this was done by calling the admin endpoint directly. In the deployed environment, the process is identical, but the URL is the deployed agent route rather than localhost.

Obtain the agent route:

```
1 cf app hybrid-rag-agent | grep routes
```

Load the ontology:

```
1 curl -X POST \
2   https://hybrid-rag-agent.<your-cf-domain>/admin/load-ontology \
3   -H "Content-Type: application/json"
```

This endpoint triggers the same ontology loading logic from Chapter 4 — parsing the RDF Turtle file and inserting the triples into HANA's graph store. On the first deployment, this will take 15-30 seconds depending on the size of the ontology. Subsequent calls are idempotent: the endpoint checks whether the graph already exists before attempting to insert.

If the call returns a 500 error, the most likely cause is a HANA connection failure. Check the agent logs:

```
1 cf logs hybrid-rag-agent --recent
```

Look for a `hdbcli` connection error. If you see “authentication failed,” confirm the `HANA_USER` and `HANA_PASSWORD` values in the User-Provided Service. If you see “host not found,” confirm the `HANA_HOST` value. Note that HANA Cloud hostnames are long strings in the format `<uuid>.hana.ondemand.com` — confirm there are no trailing spaces or missing characters.

10.9 Step 6: Smoke Test the Deployed System

With both services running and the ontology loaded, run through each integration point systematically before calling the deployment complete.

Agent health endpoint

```
1 curl https://hybrid-rag-agent.<your-cf-domain>/health
```

Expected response: a JSON object confirming service status, HANA connectivity, and Vertex AI availability. If HANA shows as unreachable, the agent will still start and respond on the health endpoint — it will report the connection failure in the status fields rather than refusing to respond.

Direct query to the agent

```
1 curl -X POST \
2   https://hybrid-rag-agent.<your-cf-domain>/query \
3   -H "Content-Type: application/json" \
4   -d '{"question": "What documents are available for material MAT-001?",
      "session_id": "smoke-test-01"}'
```

This bypasses the CAP layer entirely and tests the LangGraph agent directly against HANA. The question validates the end-to-end flow: the agent queries both the Knowledge Graph and the vector store, calls Vertex AI for answer synthesis, and returns the result. You should receive a JSON response with an `answer` field and a `retrieval_path` field indicating which chains were invoked. If this succeeds, the HANA connection, the

Vertex AI embedding model, and the LangGraph orchestration are all functional in the deployed environment.

CAP OData endpoint

```
1 curl
  https://msds-hybrid-rag-cap.<your-cf-domain>/odata/v4/documents/Documents
```

This calls the CAP OData service's Documents entity set. On a fresh deployment, the response will be an empty array `{"value": []}` — no documents have been uploaded yet. An empty array is a successful response; a 500 error indicates a CAP startup or HANA binding problem.

Fiori Elements UI

Open the CAP application URL in a browser:

```
https://msds-hybrid-rag-cap.<your-cf-domain>/index.html
```

The Fiori Elements list report should render. If you see a blank page, open the browser developer tools and check the network tab for 401 or 403 errors — these indicate an XSUAA configuration issue. For a non-XSUAA deployment (no authentication configured), the UI should load without any login screen.

10.10 Troubleshooting

Deployment failures follow predictable patterns. The following are the most common failure modes and their resolution paths.

Memory quota exceeded

Symptom: Application crashes shortly after startup. Logs show `Signal 9 (SIGKILL)` or the application disappears from `cf apps` with a restart count incrementing.

Cause: The container exceeded its memory limit. The Python buildpack allocates the full 512M during startup as the HANA driver, LangChain, and Vertex AI libraries load. If additional packages were added after the memory limit was set, startup memory may now exceed the limit.

Resolution: Increase the memory limit in `manifest.yml` to 1G and repush, or use `cf scale hybrid-rag-agent -m 1G` to update the limit without repushing.

HANA connection refused

Symptom: Agent starts successfully but `/health` reports HANA as unreachable. Logs show `[HDB][ERR] authentication failed` or `[HDB][ERR] connection refused`.

Cause: Incorrect credentials in the User-Provided Service, or the HANA Cloud instance is stopped.

Resolution: Verify the UPS values with `cf env hybrid-rag-agent` (look for `VCAP_SERVICES` in the output and find the `hybrid-rag-agent-env` entry). Compare the `HANA_HOST` value against the hostname shown in the BTP Cockpit HANA Cloud manager. If the instance is stopped, start it from the BTP Cockpit and wait two to three minutes before retrying.

GCP authentication failure

Symptom: Queries fail with a `google.api_core.exceptions.PermissionDenied` error in the agent logs.

Cause: The GCP service account credentials are not correctly configured, or the service account does not have the required Vertex AI roles.

Resolution: Confirm that `GCP_PROJECT_ID` in the UPS matches the project where Vertex AI is enabled. Confirm the Vertex AI Destination in BTP has the correct `gcp.service_account_key` value. Confirm the service account has the Vertex AI User role in GCP IAM. If you are using a key file directly, base64-encode the JSON content and set it as an environment variable, then decode it to a temp file at startup in the application's entry point.

CAP cannot reach agent

Symptom: The Fiori UI loads but queries return an error. CAP logs show `ECONNREFUSED` or `ENOTFOUND` when calling the agent URL.

Cause: The `AGENT_URL` environment variable in the CAP service contains an incorrect value, or the Python agent is not running.

Resolution: Check `cf env msds-hybrid-rag-cap` and confirm the `AGENT_URL` value is

the correct HTTPS route for the Python agent. Confirm the Python agent is running with `cf app hybrid-rag-agent`. If you recently redeployed the Python agent and its route changed, update the destination in the BTP Cockpit and use `cf set-env` to update the `AGENT_URL` value on the CAP service, then restart it.

Package not found during staging

Symptom: The `cf push` command fails during the staging phase with a `pip install` error.

Cause: A package in `requirements.txt` is either misspelled, unavailable at the specified version, or requires a system library that is not present in the Cloud Foundry container.

Resolution: Read the full staging log carefully — the package name and the error message from `pip` are usually clear. For packages that require compiled C extensions (such as `hdbcli`), confirm that a pre-compiled wheel is available for the Python version used by the buildpack. You can pin the Python version in a `runtime.txt` file at the root of the `agents/` directory with content `python-3.11.x`.

10.11 Scaling

The system as deployed runs as single instances of each service. For a pilot with a handful of users, this is sufficient. For broader use, both the Python agent and the CAP service can be scaled horizontally — adding more container instances that share the incoming request load behind the Cloud Foundry router.

Scale the Python agent to two instances:

```
1 cf scale hybrid-rag-agent -i 2
```

Scale the CAP service:

```
1 cf scale msds-hybrid-rag-cap -i 2
```

Cloud Foundry's built-in router distributes requests across instances using round-robin by default. Because the Python agent is stateless — the LangGraph agent state is not persisted between requests, and HANA handles all persistent state — horizontal scaling

requires no additional configuration. Each instance handles requests independently.

There are two situations where single-instance deployment is the correct choice even at higher load. First, during active development: when you are pushing frequent updates, having multiple instances means you need to track whether all instances have restarted to the new version. Second, during debugging: multiple instances produce interleaved log streams that can be difficult to read with `cf logs`. For debugging, scale to one instance, reproduce the problem, then scale back up.

Memory scaling — increasing the memory limit per instance — is more relevant for the Python agent than instance scaling. The LangGraph execution model creates new graph instances per request; large numbers of concurrent requests will increase memory pressure faster than CPU pressure. If you observe slow response times under load, check the memory usage with `cf app hybrid-rag-agent` before adding instances. A single 2G instance is often more cost-effective than two 512M instances for a memory-bound workload.

For production deployments with SLA requirements, consider enabling autoscaling through the Application Autoscaler service available in the BTP Service Marketplace. It can scale instances based on CPU, memory, or custom metrics, and scale them back down during quiet periods to control costs.

10.12 What You Have Built

You have deployed a production-grade Material Document Intelligence Platform to SAP BTP Cloud Foundry. Any PDF document — MSDS, invoice, batch certificate, inspection report, maintenance manual, legal filing — can be uploaded against a SAP Material Number, processed into a vector store and Knowledge Graph, and queried in natural language. The system runs on SAP's enterprise cloud platform, uses SAP HANA Cloud for both AI storage layers, integrates with real S/4HANA product master data via `API_PRODUCT_SRV`, and presents a familiar Fiori Elements interface.

Two applications are running as Cloud Foundry containers. The Python FastAPI service — `hybrid-rag-agent` — runs a LangGraph agent that dispatches every incoming question to parallel retrieval chains: a SPARQL chain that queries the HANA RDF Knowledge Graph for structured facts, and a vector similarity chain that searches HANA's vector in-

dex for semantically relevant passages. The agent synthesises the results using a Gemini model on Google Vertex AI and returns a grounded answer with a traceable retrieval path.

The CAP Node.js service — `msds-hybrid-rag-cap` — hosts an OData V4 API, a Fiori Elements UI, and an attachment handling layer. It validates material numbers against S/4HANA product master data, manages document records in HANA, and proxies processing and query requests to the FastAPI backend. Credentials for Vertex AI are stored in the BTP Destination Service — never in application code, never in deployment files.

This is the foundation for an enterprise AI capability that can scale across your SAP landscape. The same platform, the same patterns, the same HANA instance that stores your vector embeddings and Knowledge Graph triples — extended to handle every document type your organisation processes against SAP materials.

10.13 Chapter Checkpoint

Work through this list before moving to the next chapter. Each item verifies a specific part of the deployed system.

- `cf app hybrid-rag-agent` shows the application in started state with a route assigned. (Section 10.4)
- `curl https://hybrid-rag-agent.<domain>/health` returns a JSON response with no error fields. (Section 10.9)
- The User-Provided Service `hybrid-rag-agent-env` exists and is bound to both applications. (Section 10.5)
- `cf deploy` completed without errors and both applications appear in `cf apps`. (Section 10.6)
- A BTP Destination named `hybrid-rag-agent` exists in the subaccount, pointing to the Python agent URL. (Section 10.7)
- A BTP Destination named `VertexAI` exists with the GCP service account key stored as a destination property. (Section 10.7)
- `curl -X POST .../admin/load-ontology` returned a 200 status code. (Section 10.8)
- A direct `curl POST` to `/query` with question “What documents are available for material MAT-001?” returns an answer with a retrieval path. (Section 10.9)

- The Fiori Elements UI loads at the CAP application URL without errors in the browser console. (Section 10.9)

If every item on this list is green, the system is deployed and functional. The next chapter addresses operational concerns — observability, performance tuning, and the integration patterns that allow this system to participate in broader SAP workflows.

Appendix A: SAP AI Core — Enterprise LLM Integration

Who this appendix is for. This appendix is for readers who have access to an SAP AI Core enterprise subscription and want to use SAP-managed foundation models instead of Google Vertex AI. The agent architecture, LangGraph orchestration, and HANA Cloud retrieval are identical — only the model provider changes. If you are using Vertex AI (the default for this book), you can skip this appendix entirely.

A.1 What Is SAP AI Core?

SAP AI Core is the enterprise AI deployment and orchestration platform on SAP BTP. It provisions and runs foundation models at scale, manages model lifecycle, handles token quotas and rate limiting per tenant, and provides the secure inference endpoint that Joule uses internally.

For this book's use case, AI Core provides two things:

Capability	Detail
Generative AI Hub	A unified API layer over multiple LLM providers — OpenAI GPT-4o, Anthropic Claude, Google Gemini, Mistral, and SAP's own models
Embeddings endpoint	A managed embedding model endpoint compatible with the same interface as Vertex AI's <code>text-embedding-004</code>

The LangChain integration library `generative-ai-hub-sdk` wraps both endpoints with a familiar interface, making the swap from Vertex AI straightforward.

A.2 Prerequisites

Before starting this appendix, confirm you have:

- An SAP BTP subaccount with **SAP AI Core** service entitlement (Standard or Extended plan)
- An **AI Launchpad** subscription in the same subaccount
- A service key for the AI Core instance (downloaded from BTP Cockpit → Services → Instances)

Trial accounts cannot use SAP AI Core. AI Core requires a paid BTP subscription. Contact your SAP account executive or BTP administrator if you are unsure whether your organisation has the entitlement.

A.3 Activating a Model in AI Launchpad

AI Launchpad is the browser-based control plane for AI Core. You use it to activate (deploy) a foundation model before your application can call it.

Step 1 — Open AI Launchpad

In the BTP Cockpit, go to **Services** → **Instances and Subscriptions** → click **AI Launchpad** to open it.

Step 2 — Create a Resource Group

Resource groups isolate model deployments by use case or team.

Go to **ML Operations** → **Resource Groups** → **Create**.

Field	Value
Resource Group ID	agentic-rag
Labels	project=book (optional)

Step 3 — Activate a Generative AI Hub Model

Go to **Generative AI Hub** → **Models** → find the model you want to deploy.

For this book the recommended replacements are:

Vertex AI model	AI Core equivalent
gemini-2.5-flash	gpt-4o or claude-3.5-sonnet
text-embedding-004	text-embedding-ada-002 or SAP's managed embedding model

Click the model **Deploy** ☐ select your resource group `agentic-rag` **Deploy**.

Deployment takes 2–5 minutes. The model status changes from **Pending** to **Running**.

Note the **Deployment ID** — you will need it in your configuration.

Note: You are not charged for deploying a model. Charges apply only to token consumption during inference calls.

Step 4 — Get the Inference Endpoint

Click your running deployment ☐ copy the **Inference URL**. It will look like:

```
https://api.ai.prod.eu-central-1.aws.ml.hana.ondemand.com/v2/inference/deployments/<deployment-id>
```

A.4 Getting the AI Core Service Key

In BTP Cockpit ☐ **Services** ☐ **Instances** ☐ click your AI Core instance ☐ **Service Keys** ☐ **Create Service Key**.

Download the JSON. It contains:

```
1 {
2   "clientid": "sb-...",
3   "clientsecret": "...",
4   "url": "https://<tenant>.authentication.eu10.hana.ondemand.com",
5   "serviceurls": {
6     "AI_API_URL":
7       "https://api.ai.prod.eu-central-1.aws.ml.hana.ondemand.com"
8   }
9 }
```

Add these to your agents/.env:

```
1 AICORE_CLIENT_ID=<clientid>
2 AICORE_CLIENT_SECRET=<clientsecret>
3 AICORE_AUTH_URL=<url>
4 AICORE_BASE_URL=<AI_API_URL>
5 AICORE_RESOURCE_GROUP=agentic-rag
6 AICORE_DEPLOYMENT_ID=<your-deployment-id>
```

A.5 Installing the SDK

```
1 pip install "generative-ai-hub-sdk[all]"
```

Verify the installation:

```
1 python -c "from gen_ai_hub.proxy.langchain import init_llm; print('SDK ready')"
```

A.6 Calling the Model Directly (Python)

Before integrating with the agent, verify the model works standalone:

```
1 from gen_ai_hub.proxy.core.proxy_clients import get_proxy_client
2 from gen_ai_hub.proxy.langchain import init_llm, init_embedding_model
3
4 # Initialise the proxy client (reads from environment variables
  automatically)
5 proxy_client = get_proxy_client("gen-ai-hub")
6
7 # Test LLM call
8 llm = init_llm("gpt-4o", proxy_client=proxy_client, temperature=0.0)
9 response = llm.invoke("In one sentence, what is SAP HANA Cloud?")
10 print(response.content)
11
12 # Test embedding call
```

```

13 embedding_model = init_embedding_model(
14     "text-embedding-ada-002",
15     proxy_client=proxy_client
16 )
17 vectors = embedding_model.embed_documents(["tensile strength test ISO
18     6892-1"])
19 print(f"Embedding dimension: {len(vectors[0])}")

```

Expected output:

SAP HANA Cloud is a cloud-native database that combines in-memory processing, multi-model data management, and AI capabilities in a single managed service. Embedding dimension: 1536

Note: text-embedding-ada-002 produces 1536-dimensional vectors, compared to 768 for Vertex AI's text-embedding-004. If you switch embedding models on an existing HANA Cloud deployment, you must re-embed all documents — the vector dimensions must match the REAL_VECTOR column definition. Either recreate the table with REAL_VECTOR(1536) or keep Vertex AI embeddings and use AI Core only for the LLM tasks.

A.7 Replacing Gemini in the Agent

The agent uses two Gemini calls: one for SPARQL generation (the KG chain) and one for answer synthesis (the merge node). Both use the `_get_llm()` lazy getter pattern — swapping providers requires changing only that function.

Step 1 — Replace `_get_llm()` in `agents/agents/llm.py`

```

1  # Before (Vertex AI / Gemini)
2  from langchain_google_genai import ChatGoogleGenerativeAI
3
4  def _get_llm():
5      return ChatGoogleGenerativeAI(
6          model="gemini-2.5-flash",
7          temperature=0.0,
8          max_tokens=2048,

```

```

9         )
10
11     # After (SAP AI Core)
12     from gen_ai_hub.proxy.langchain import init_llm
13     from gen_ai_hub.proxy.core.proxy_clients import get_proxy_client
14
15     _proxy_client = None
16
17     def _get_proxy_client():
18         global _proxy_client
19         if _proxy_client is None:
20             _proxy_client = get_proxy_client("gen-ai-hub")
21         return _proxy_client
22
23     def _get_llm():
24         return init_llm(
25             "gpt-4o",
26             proxy_client=_get_proxy_client(),
27             temperature=0.0,
28             max_tokens=2048,
29         )

```

Step 2 — Replace the embedding call in `agents/agents/vector_srv.py`

```

1     # Before (Vertex AI)
2     from langchain_google_genai import GoogleGenerativeAIEmbeddings
3
4     def _get_embedding_model():
5         return GoogleGenerativeAIEmbeddings(model="models/text-embedding-004")
6
7     # After (SAP AI Core)
8     from gen_ai_hub.proxy.langchain import init_embedding_model
9     from agents.llm import _get_proxy_client
10
11     def _get_embedding_model():
12         return init_embedding_model(
13             "text-embedding-ada-002",
14             proxy_client=_get_proxy_client()
15         )

```


Important: If you already have vectors stored in HANA Cloud using `text-embedding-004` (768 dimensions), switching to `text-embedding-ada-002` (1536 dimensions) requires re-ingesting all documents. The vector dimensions must match. To avoid re-ingestion, keep Vertex AI for embeddings and swap only the LLM to AI Core.

Step 3 — Update `.env`

Comment out the Vertex AI variables and add the AI Core variables from Section E.4:

```
1 # Vertex AI (comment out if using AI Core for LLM)
2 # GOOGLE_CLOUD_PROJECT=...
3 # GCP_LOCATION=us-central1
4
5 # SAP AI Core
6 AICORE_CLIENT_ID=...
7 AICORE_CLIENT_SECRET=...
8 AICORE_AUTH_URL=...
9 AICORE_BASE_URL=...
10 AICORE_RESOURCE_GROUP=agentic-rag
11 AICORE_DEPLOYMENT_ID=...
```

Step 4 — Test the swapped agent

```
1 cd agents
2 python -c "
3 from agents.llm import _get_llm
4 llm = _get_llm()
5 result = llm.invoke('What is a batch quality certificate in SAP QM?')
6 print(result.content[:200])
7 "
```

If the response prints correctly, the swap is complete. Run the full agent:

```
1 uvicorn main:app --reload
2 curl -X POST http://localhost:8000/query \
3   -H 'Content-Type: application/json' \
4   -d '{"question": "What test method was
      used?", "material_number": "BATCH-QC-MAT-001", "history": []}'
```

The response structure is identical — `answer`, `kg_facts`, `vector_chunks`, `sources`. No changes to the CAP service, Fiori UI, or LangGraph graph are required.

A.8 BTP Destination for AI Core (CF Deployment)

For BTP Cloud Foundry deployment, store AI Core credentials in the Destination Service rather than as CF environment variables.

In BTP Cockpit ▢ **Connectivity** ▢ **Destinations** ▢ **Create** ▢ **From Scratch**:

Field	Value
Name	SAP-AI-Core
Type	HTTP
URL	<AI_API_URL> from service key
Authentication	OAuth2ClientCredentials
Client ID	clientid from service key
Client Secret	clientsecret from service key
Token Service URL	<url>/oauth/token from service key

Add additional property:

Property	Value
ai-resource-group	agentic-rag

In `main.py`, read the destination at startup using the BTP Destination Service SDK — the same pattern used for the Vertex AI destination in Chapter 2.

A.9 Summary

Step	What you did
Activated model in AI Launchpad	Deployed gpt-4o and text-embedding-ada-002 to resource group agentic-rag
Installed SDK	<code>pip install generative-ai-hub-sdk[all]</code>
Replaced <code>_get_llm()</code>	Swapped ChatGoogleGenerativeAI for <code>init_llm("gpt-4o", ...)</code>
Replaced embedding model	Swapped GoogleGenerativeAIEmbeddings for <code>init_embedding_model(...)</code>
Updated <code>.env</code>	Added AI Core credentials, commented out Vertex AI
Tested end-to-end	Full <code>/query</code> endpoint returns correct responses via AI Core

The LangGraph graph, HANA Cloud retrieval, CAP OData service, and Fiori UI are completely unchanged. SAP AI Core is a drop-in replacement for the model inference layer.